

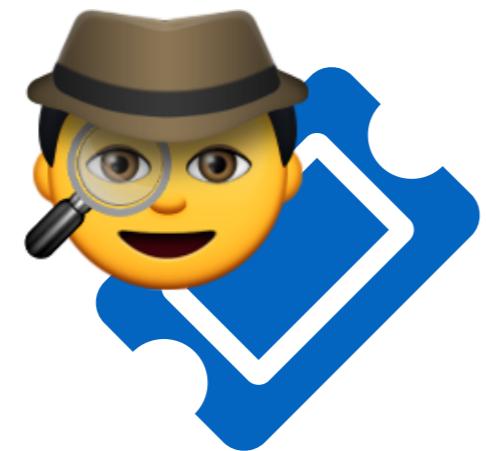
Mitigating Contention in Distributed Systems

Generals Exam

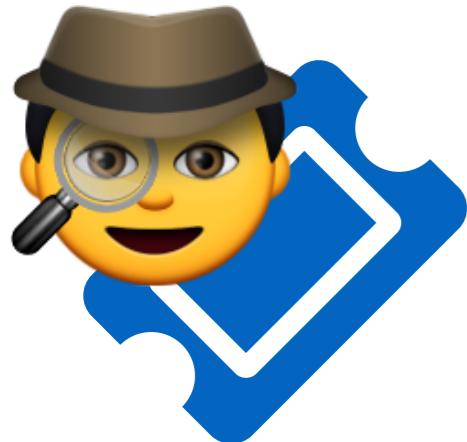
Brandon Holt

Committee

Luis Ceze, Mark Oskin, Dan Ports, Jevin West



TICKETSLEUTH



Customer Ratings

Current Version

All Versions

Average Rating:



12,725 Ratings

Click to rate:



Customer Reviews

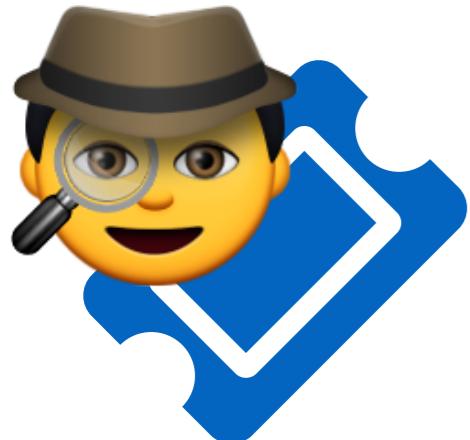
Write a Review

App Support

TicketSleuth is #winning!! ★★★★★

by bholt1111 – Nov 13, 2015





F Star Wars' Presales Crash

www.forbes.com/sites/hayleycuccinello/2015/10/20/star-wars-presales-crash-ticketing-sites-sets-record-for-fandango/

Forbes / Media & Entertainment

The Little Black Book of Billionaire Secrets

OCT 20, 2015 @ 03:55 PM 138,858 VIEWS

'Star Wars' Presales Crash Ticketing Sites, Set Record for Fandango

Hayley C. Cuccinello, CONTRIBUTOR
I write about television and theater

[FOLLOW ON FORBES \(11\)](#)

Opinions expressed by Forbes Contributors are their own.

[FULL BIO ▾](#)

COMMENT NOW

SHARE >

f t e in

SHARE >

A red brushstroke highlights the author's name, Hayley C. Cuccinello.

A red brushstroke highlights the 'COMMENT NOW' button.

A red brushstroke highlights the social sharing icons (Facebook, Twitter, Email, LinkedIn).

A red brushstroke highlights the 'SHARE >' button.

A red brushstroke highlights the 'FULL BIO ▾' button.



What causes contention?



Power Laws

Focus on few records.



Interactivity

Heavy write load → synchronize or conflict



Realtime Events

Focus in time.



Introduction: Contention



- ✓ Sources of contention.
- 3 broad approaches to mitigating contention.

Background: Trading off consistency

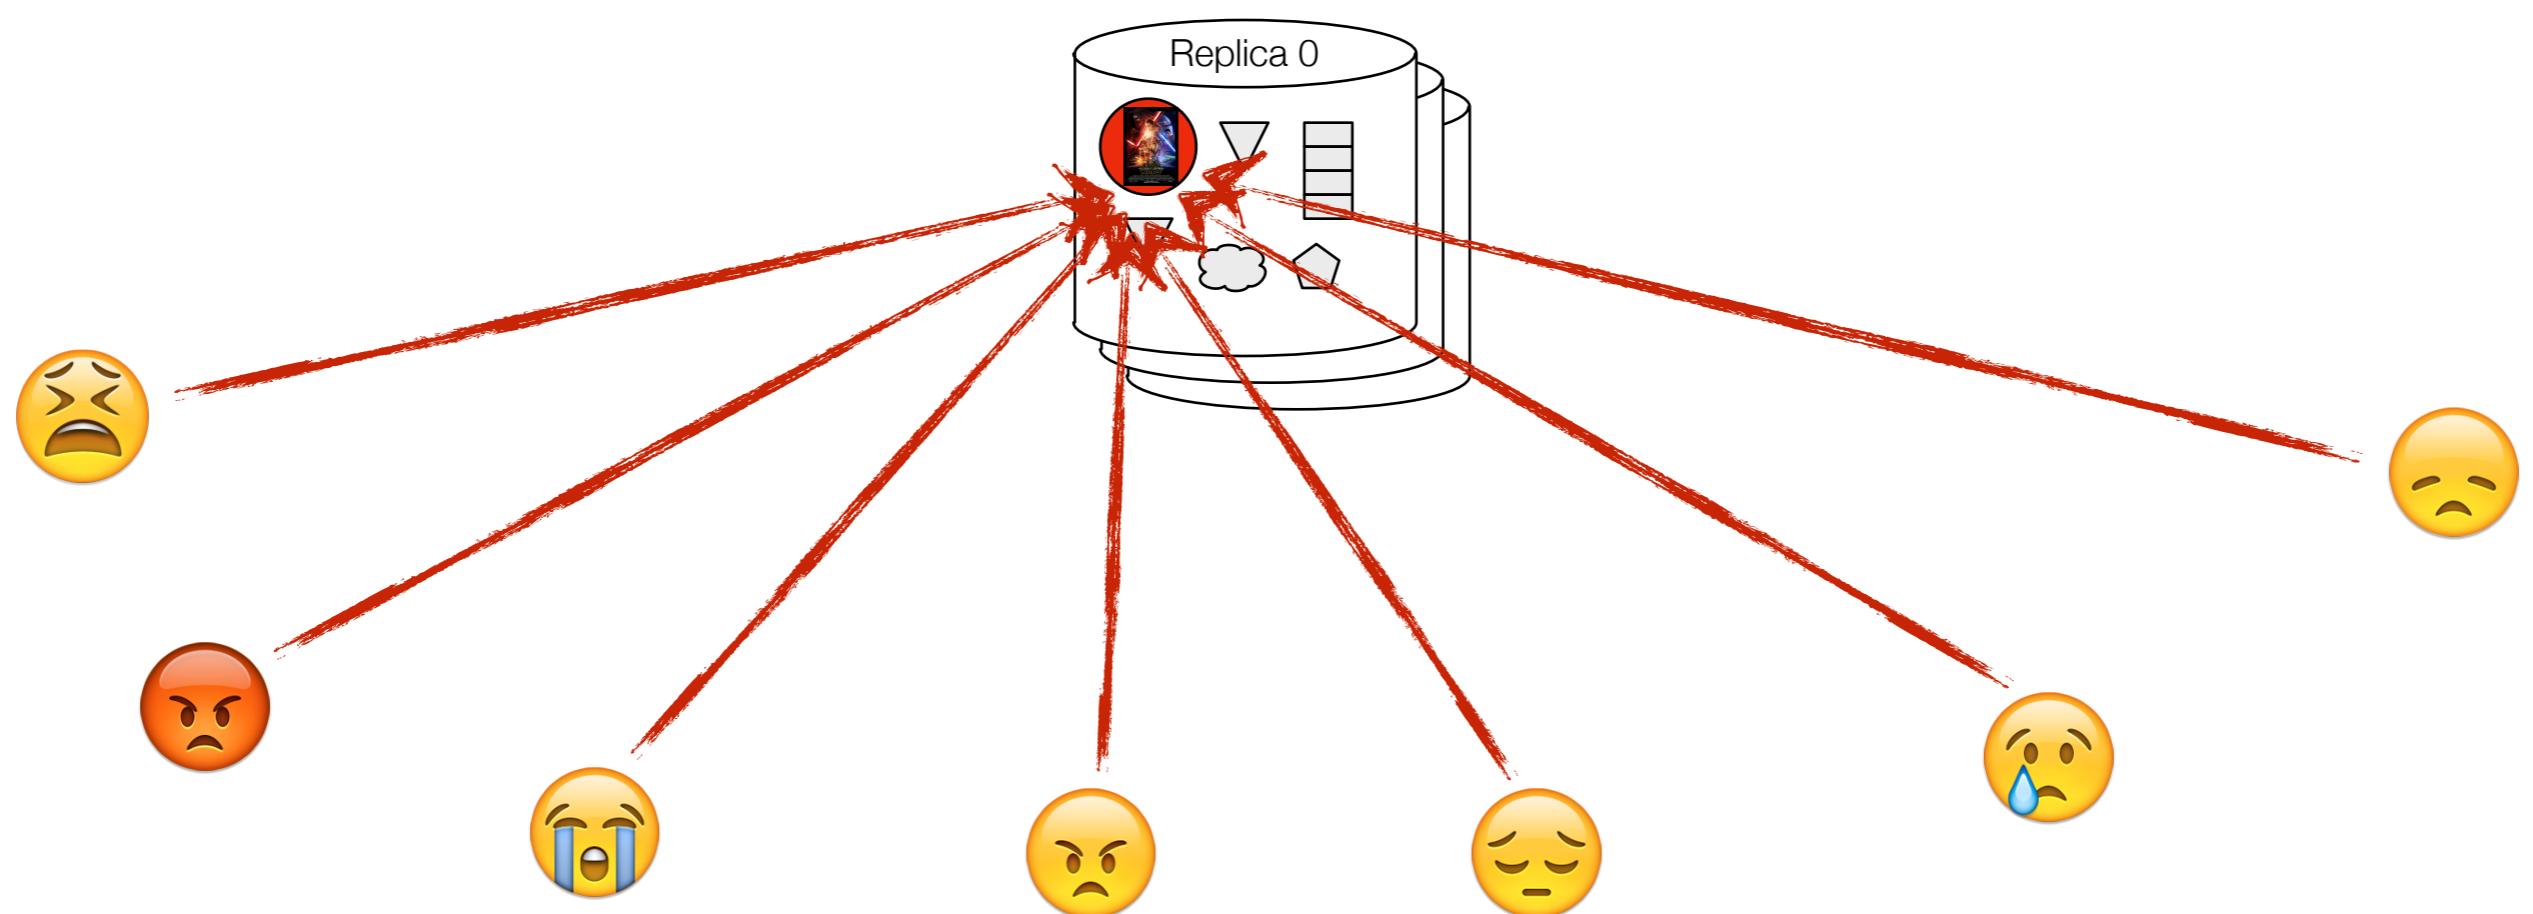
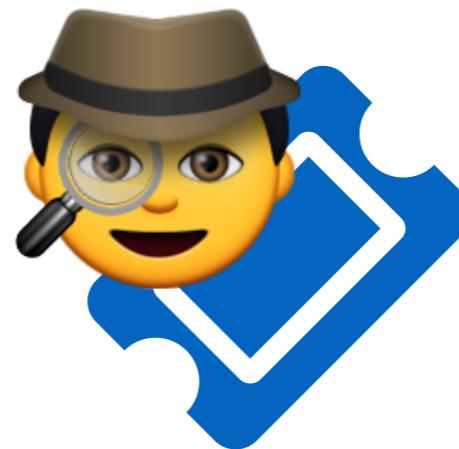


- Consistency for architects.
- Ordering and visibility constraints.
- Dealing with uncertainty: **converged state vs. staleness**
- Programming with tradeoffs.

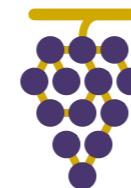


Proposal: Disciplined Inconsistency

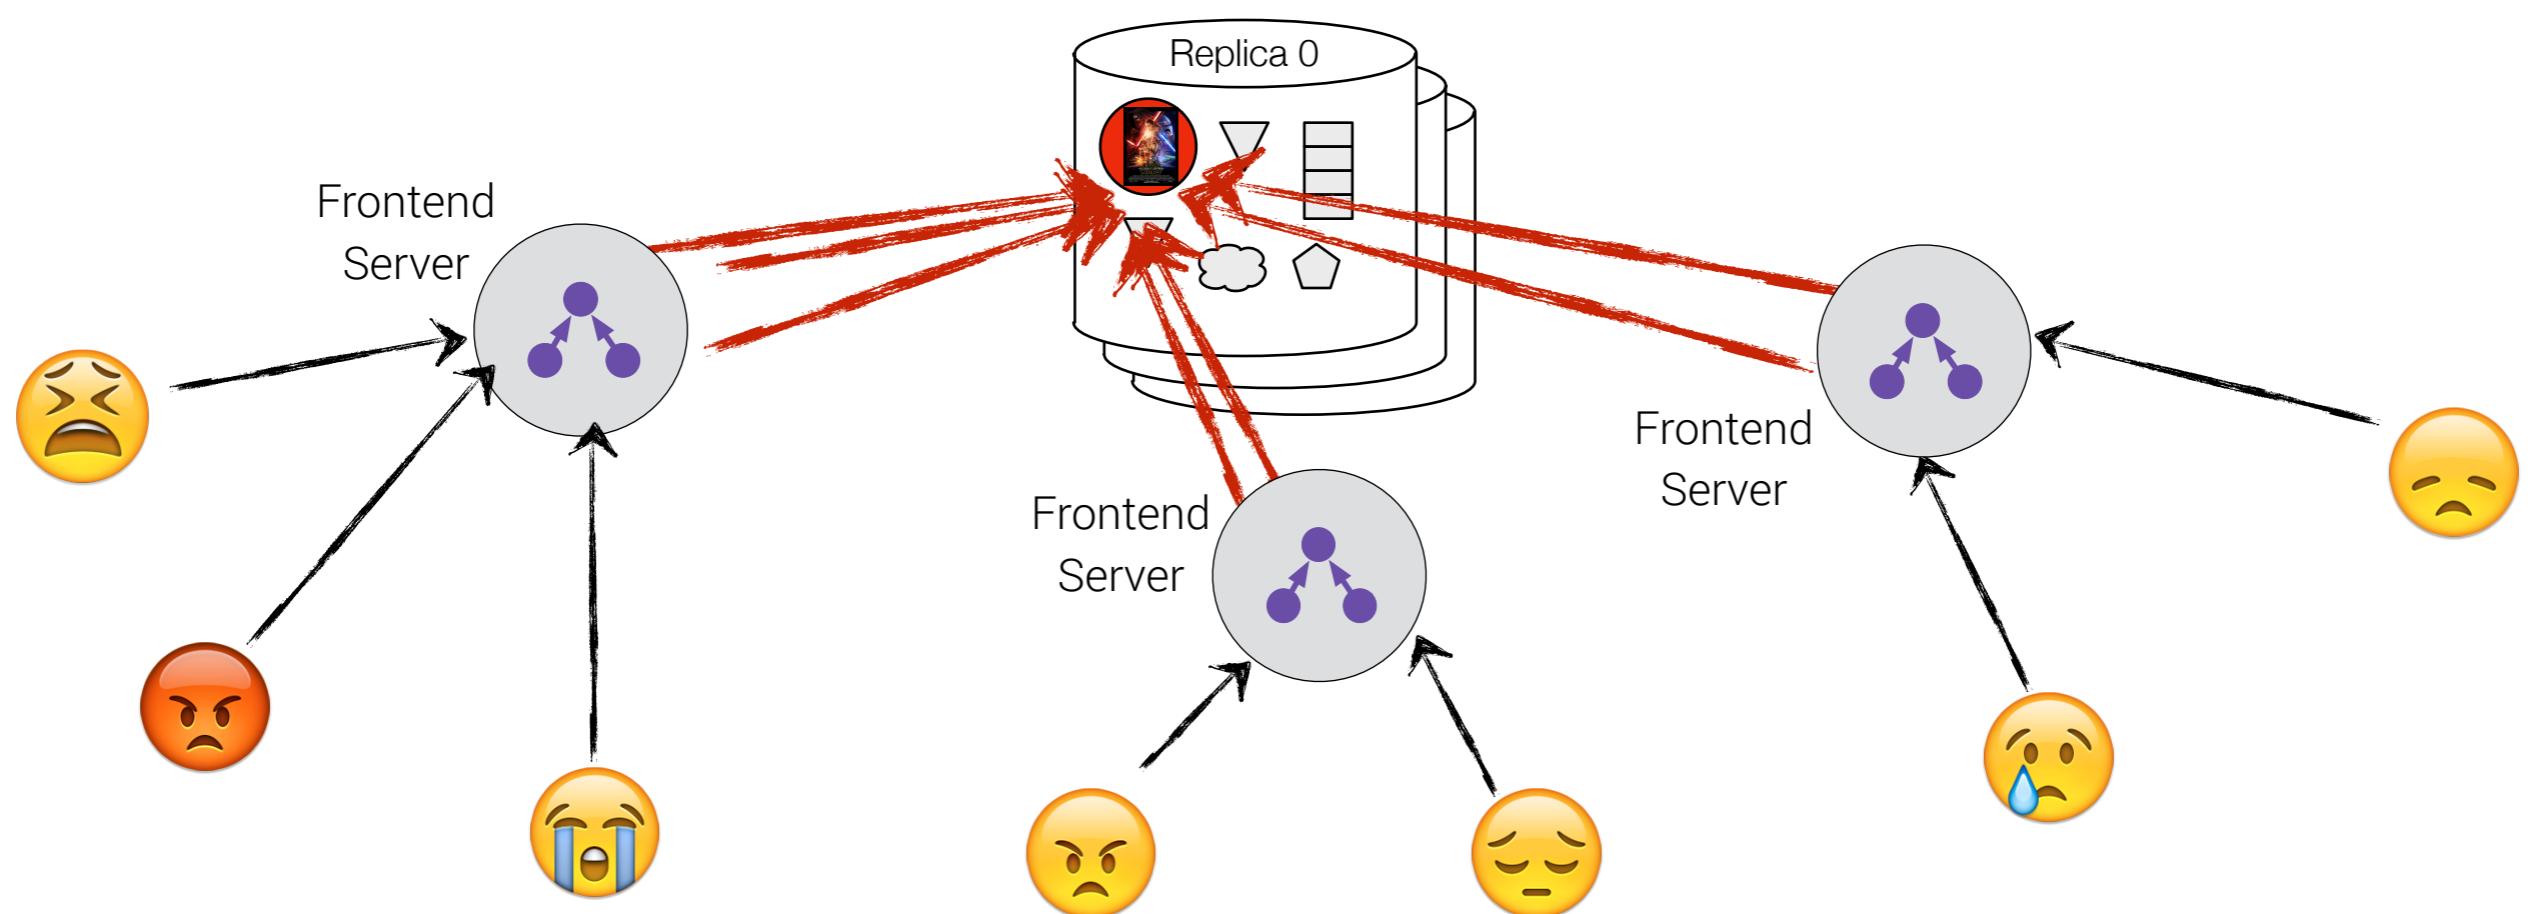
- Make tradeoffs *explicit* and *safe* with types.
- Reason about *values* rather than *ordering*.
- Synthesize synchronization/enforcement logic.



① Offload synchronization to clients.



Flat Combining Global Shared Data Structures
B. Holt, et al. – PGAS'13



Introduction

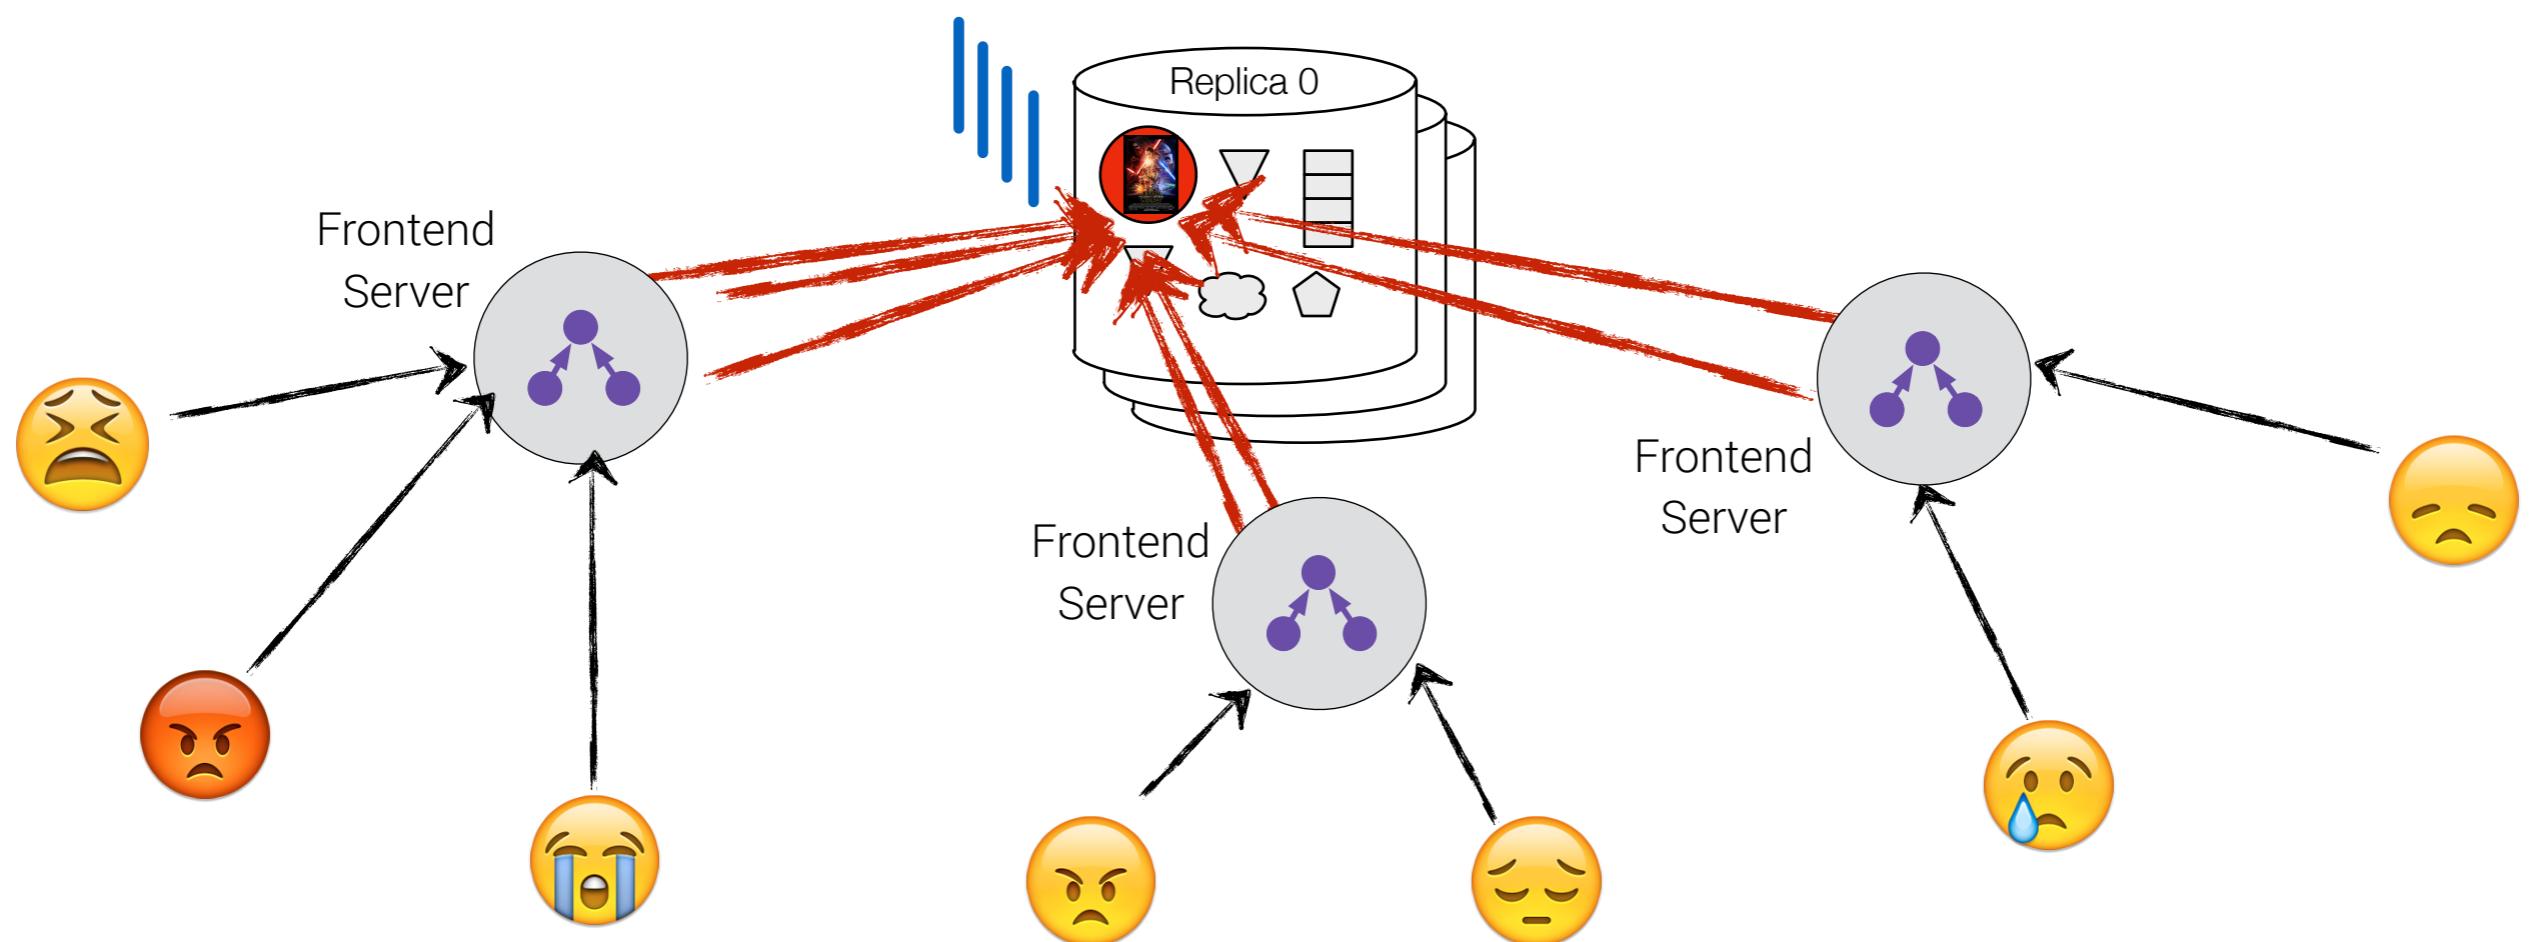
Mitigating contention

① Offload synchronization to clients.

② Unlock safe concurrency within the record.

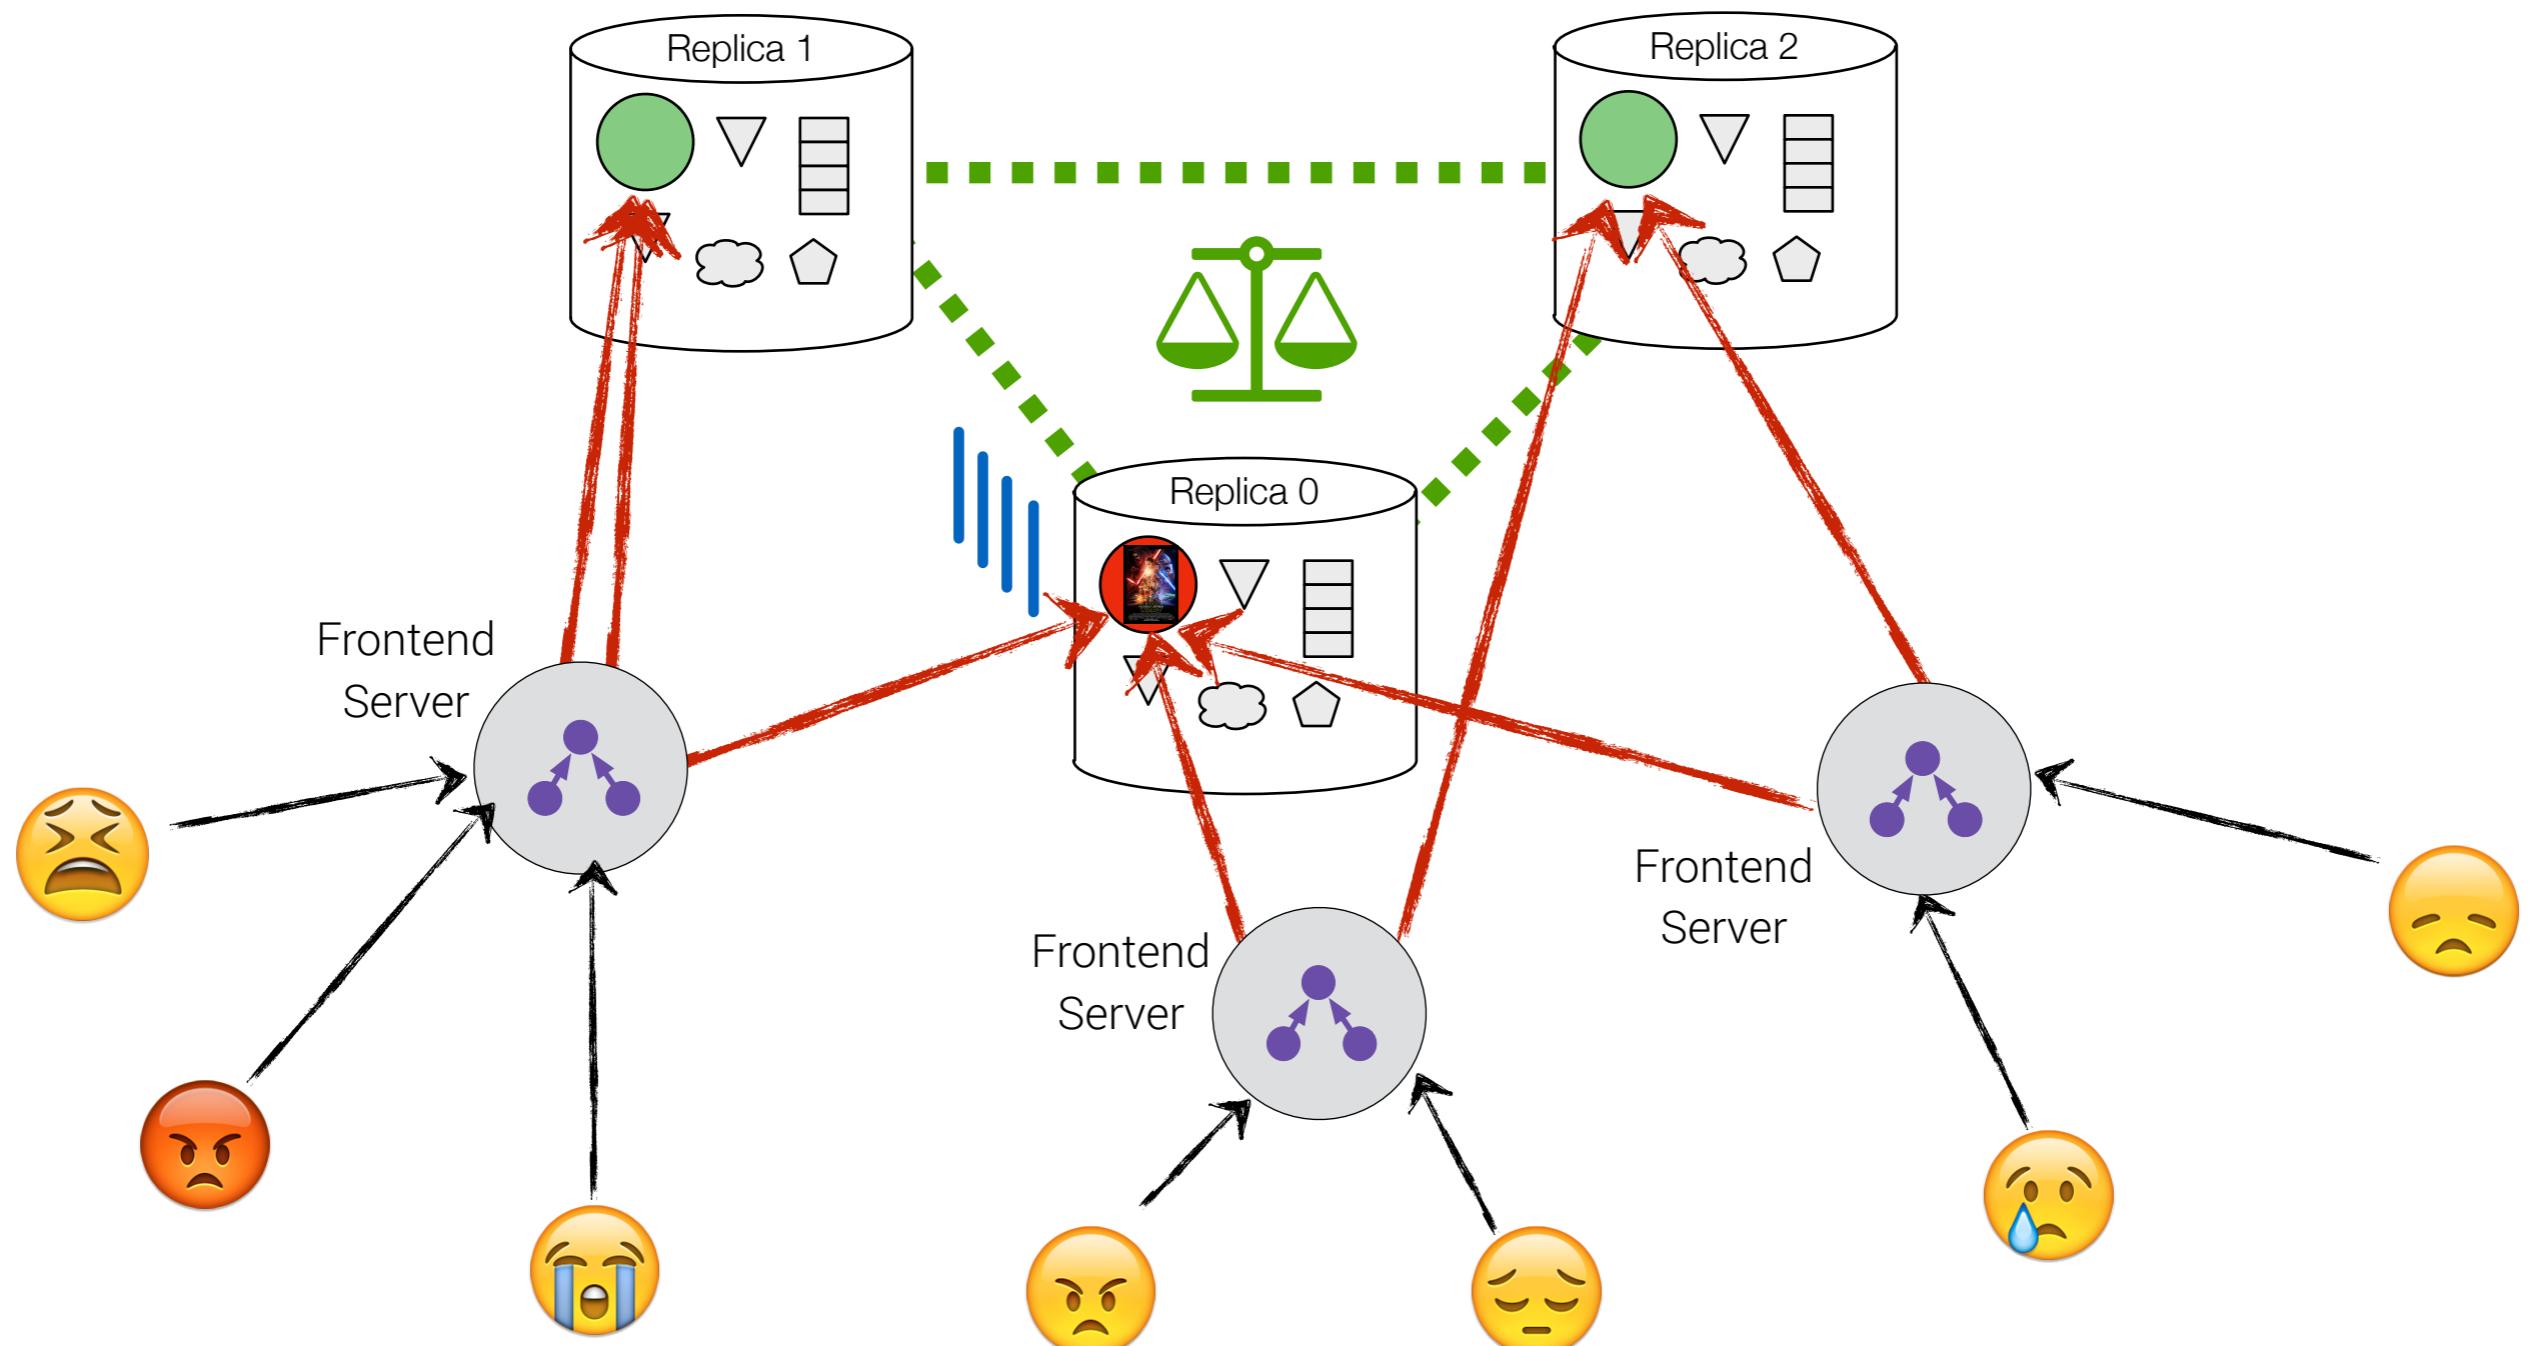


Claret: Avoiding Contention in Distributed Transactions with Abstract Data Types
B. Holt, et al. – *In submission*.



Mitigating contention

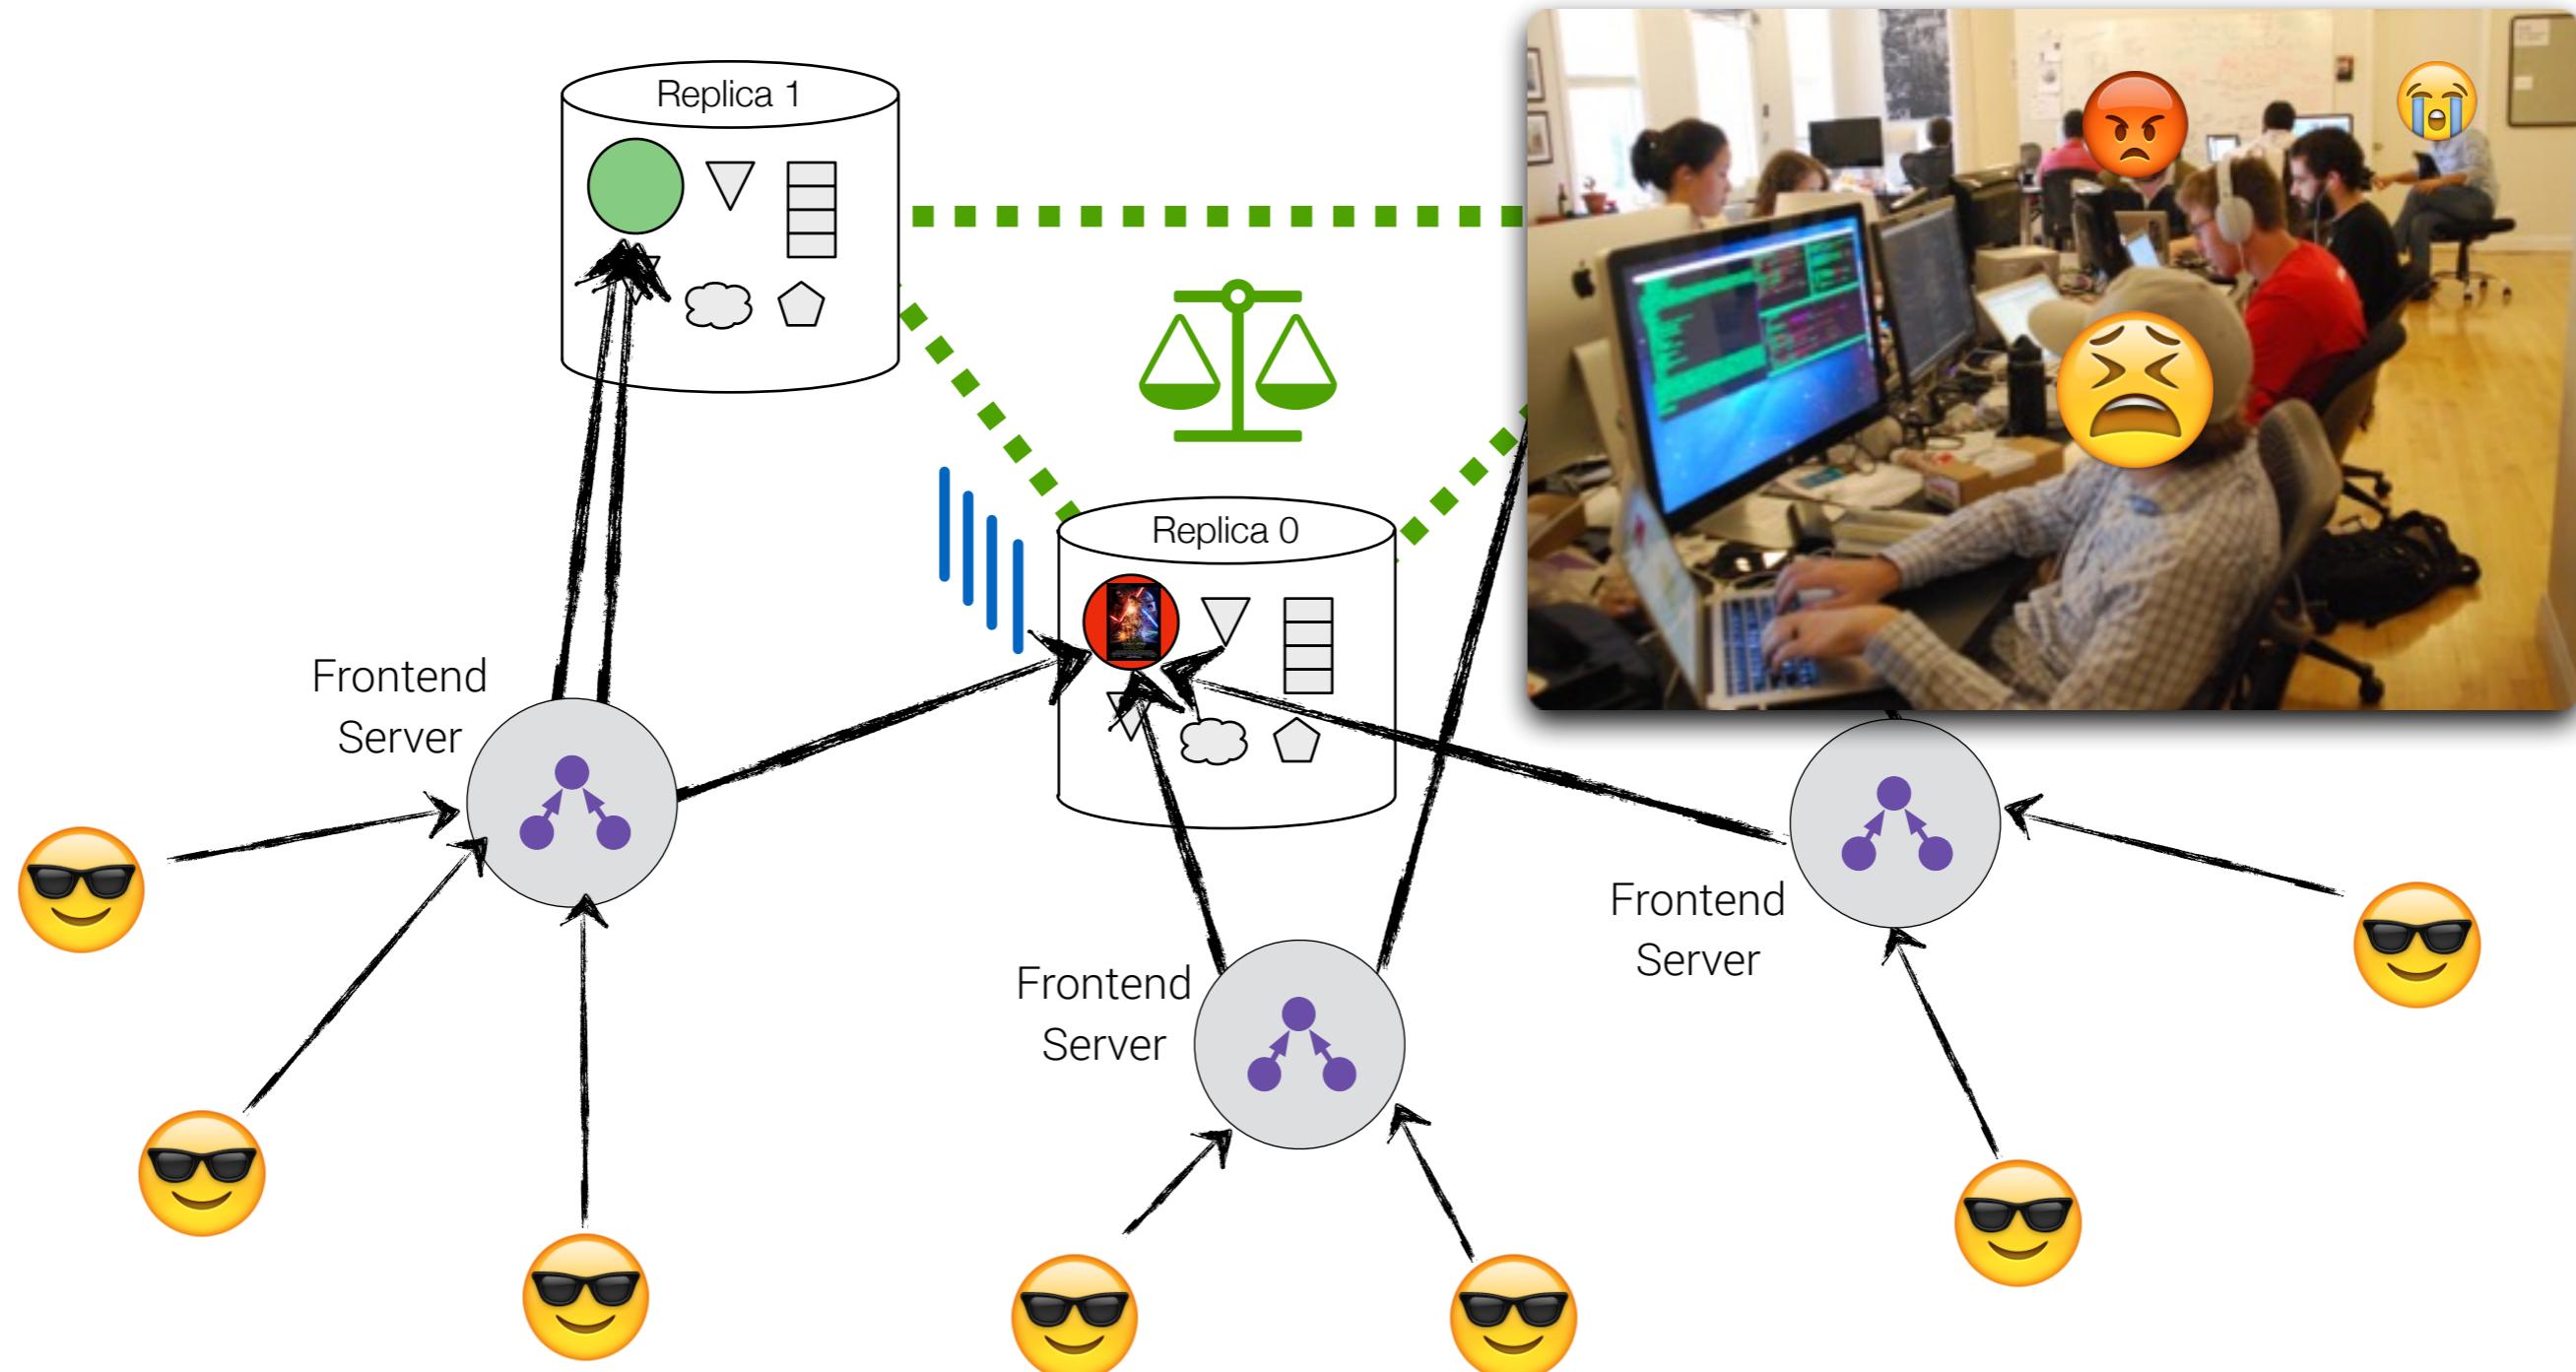
- ① Offload synchronization to clients.
- ② Unlock safe concurrency within the record.
- ③ Trade off consistency with weak replication.



Introduction

Mitigating contention

- ① Offload synchronization to clients.
- ② Unlock safe concurrency within the record.
- ③ Trade off consistency with weak replication.



③ Trade off consistency with weak replication.

How should coordination be specified and enforced?

How can programming models help make these trade-offs?



Introduction: Contention



- ✓ Sources of contention.
- ✓ 3 broad approaches to mitigating contention.

Background: Trading off consistency



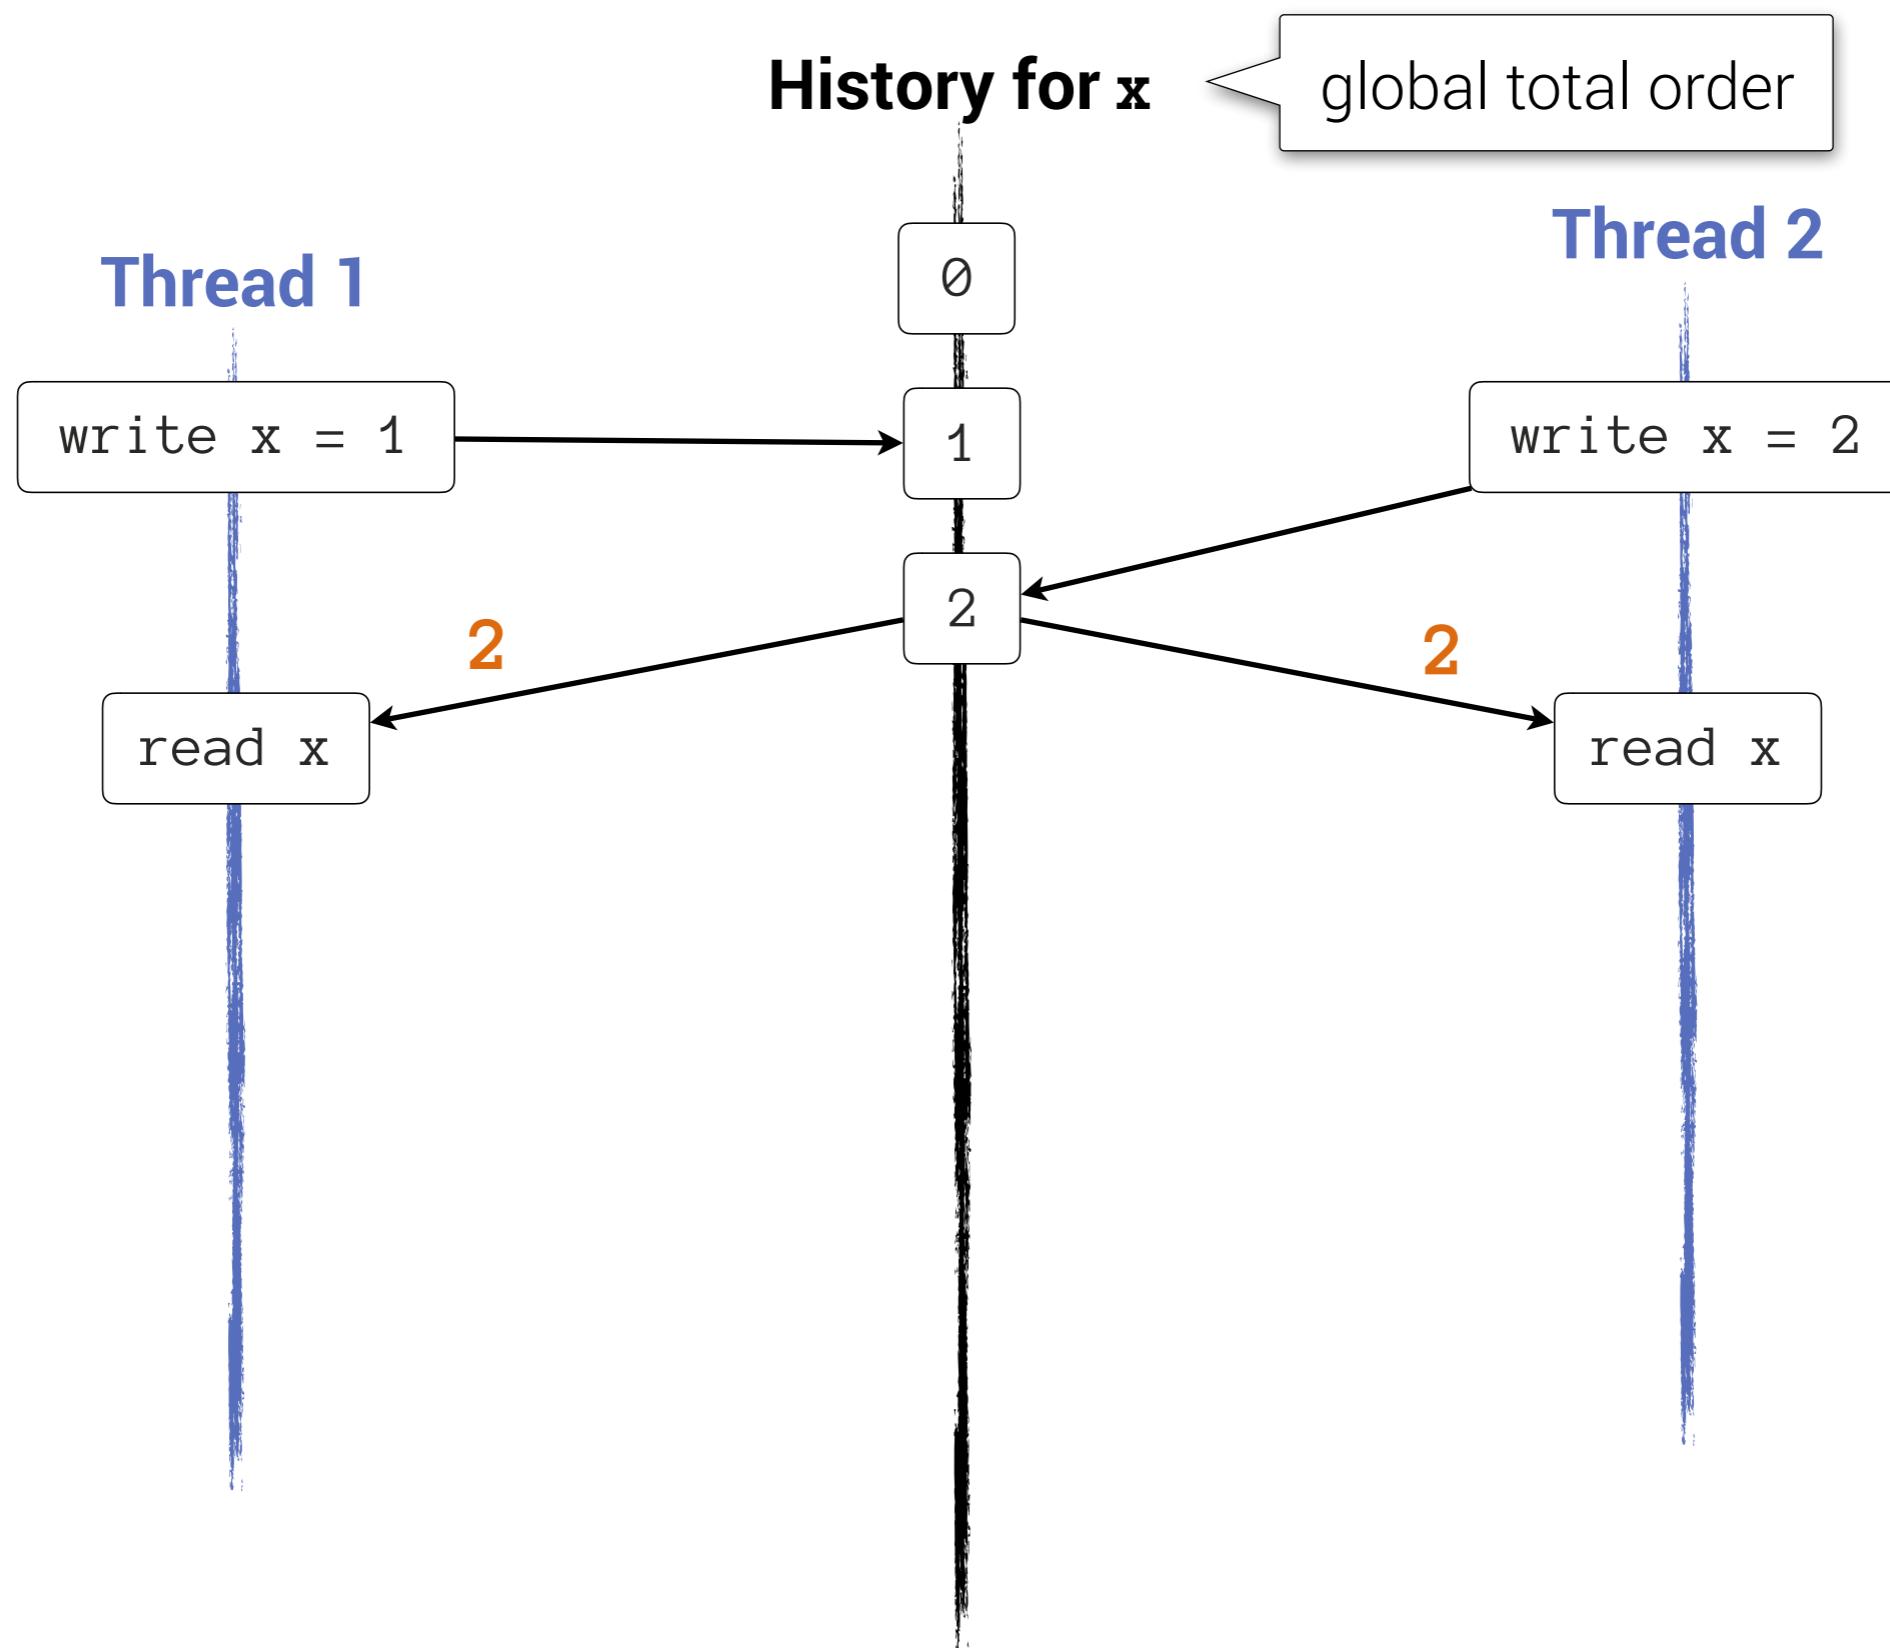
- Consistency for architects.
 - Ordering and visibility constraints.
 - Dealing with uncertainty.
 - Programming with tradeoffs.

Proposal: Disciplined Inconsistency

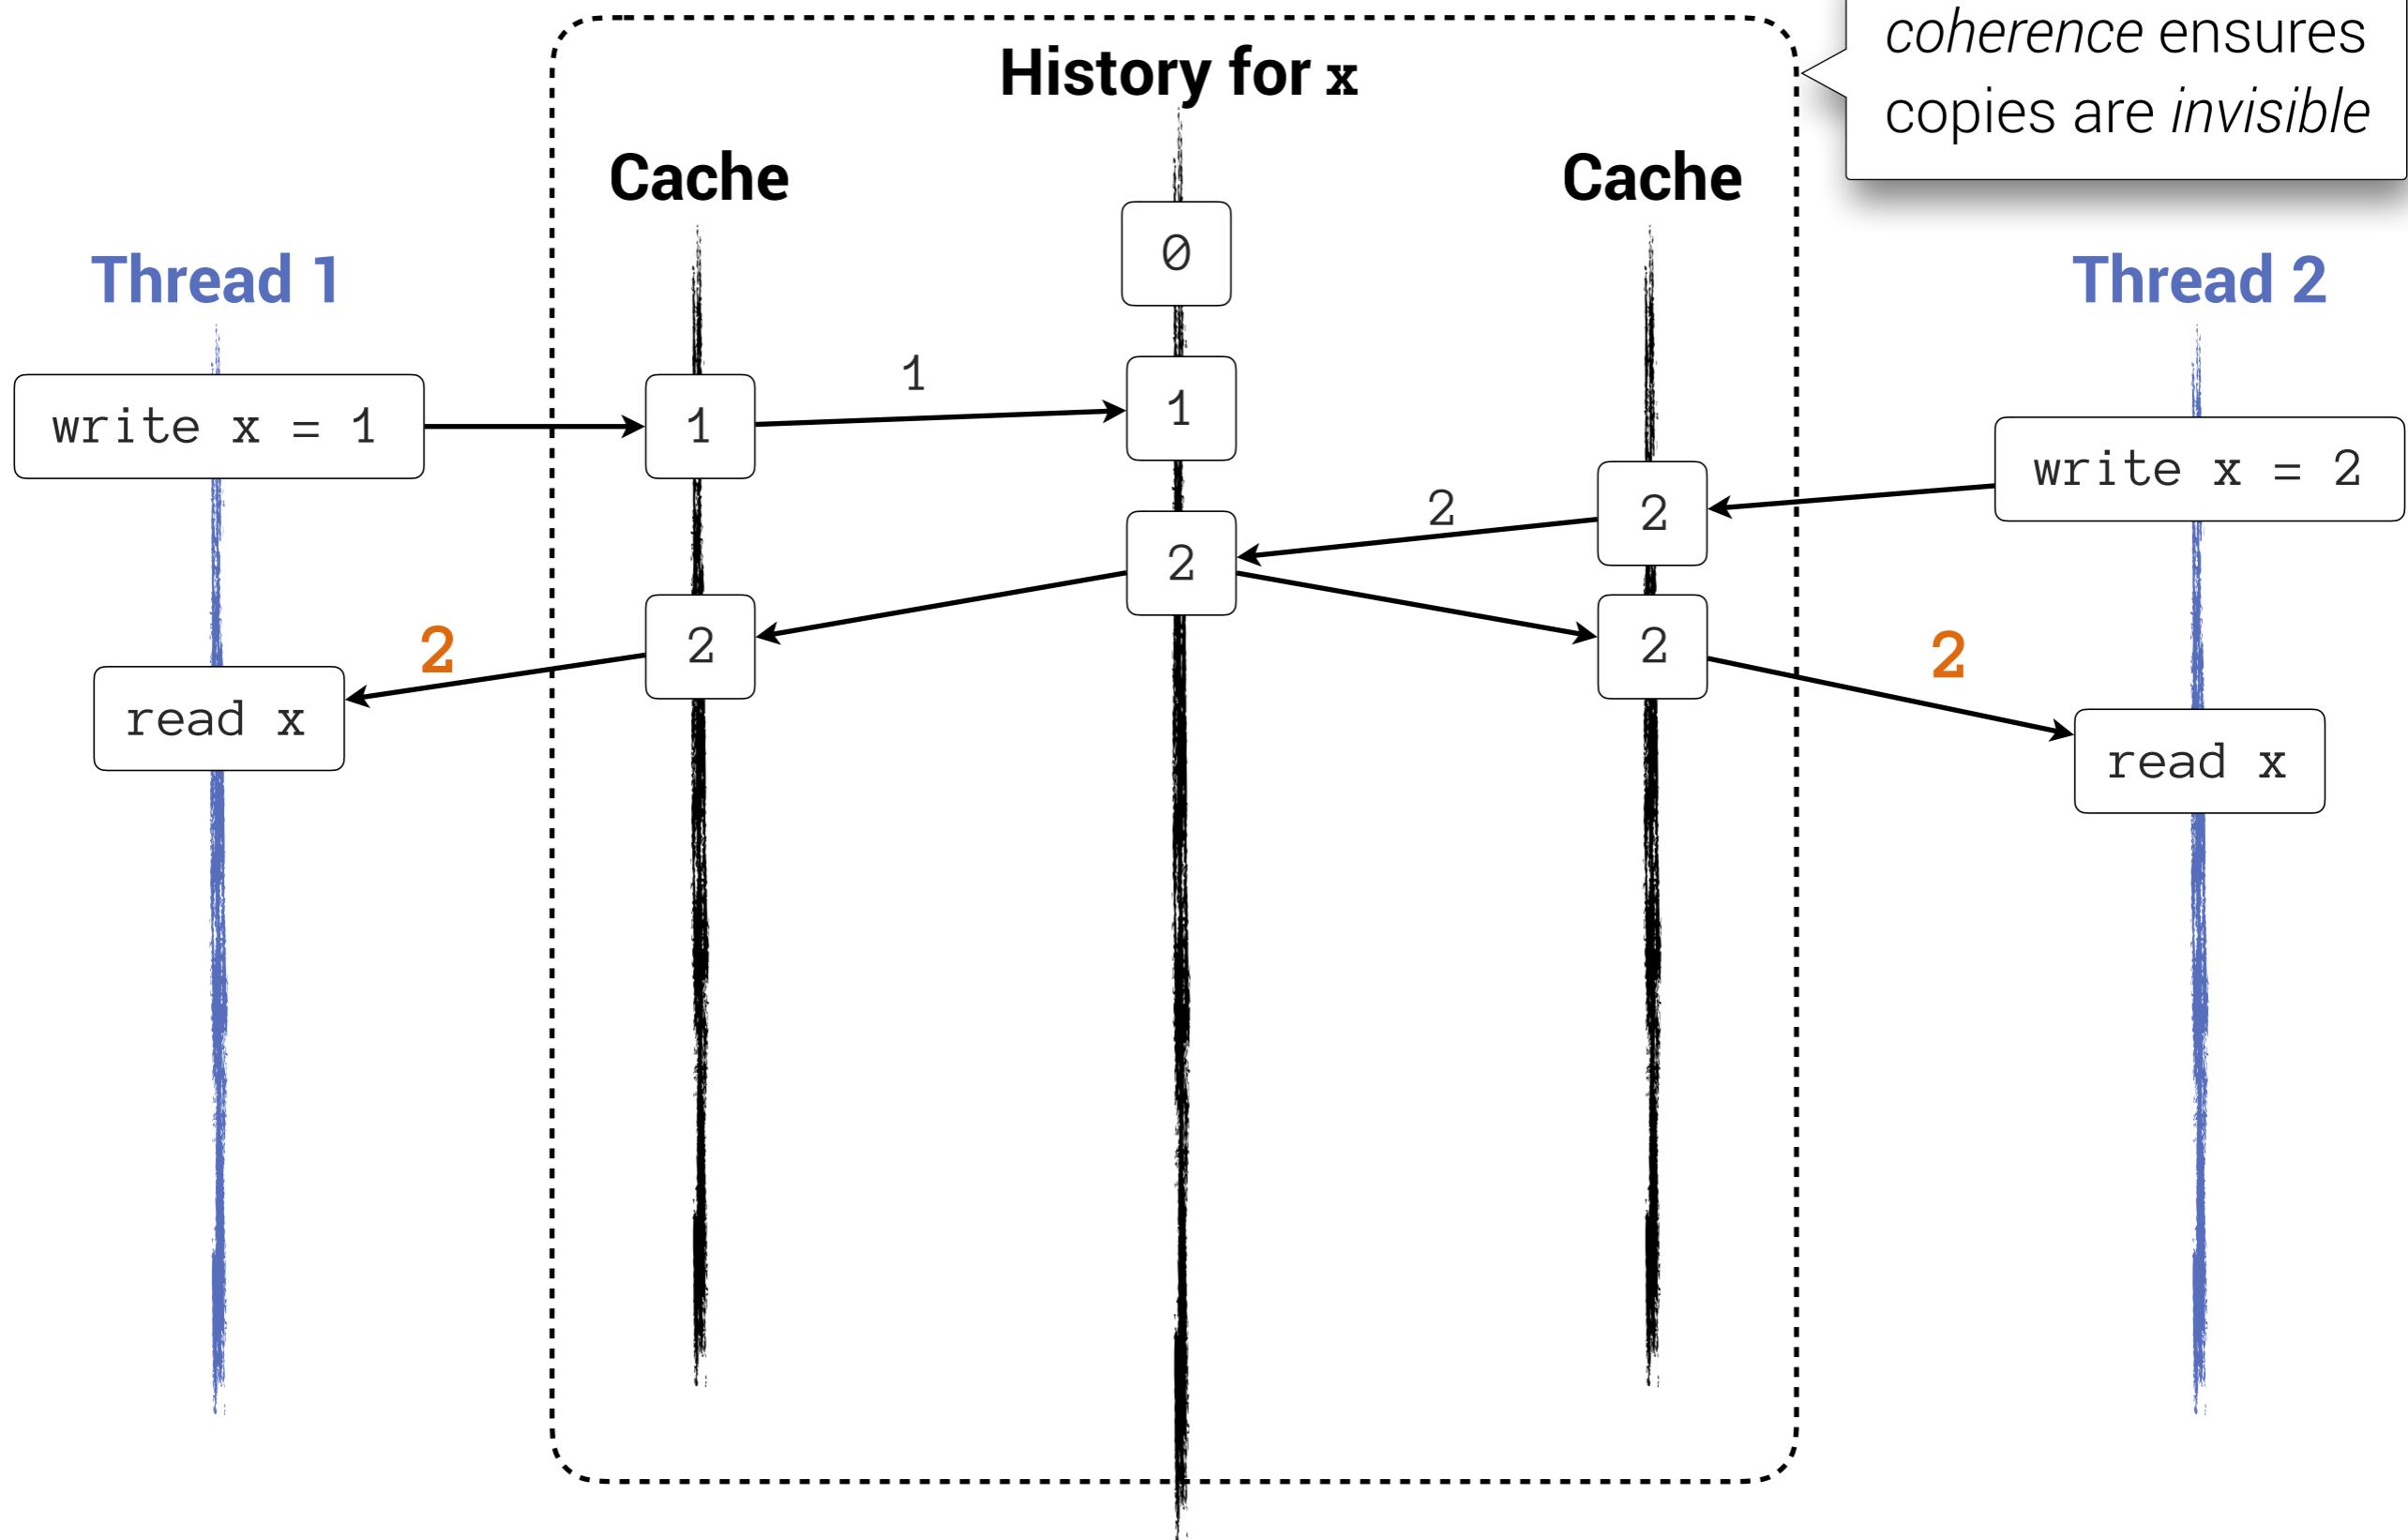


- Make tradeoffs *explicit* and *safe* with types.
- Reason about *values* rather than *ordering*.
- Synthesize synchronization/enforcement logic.

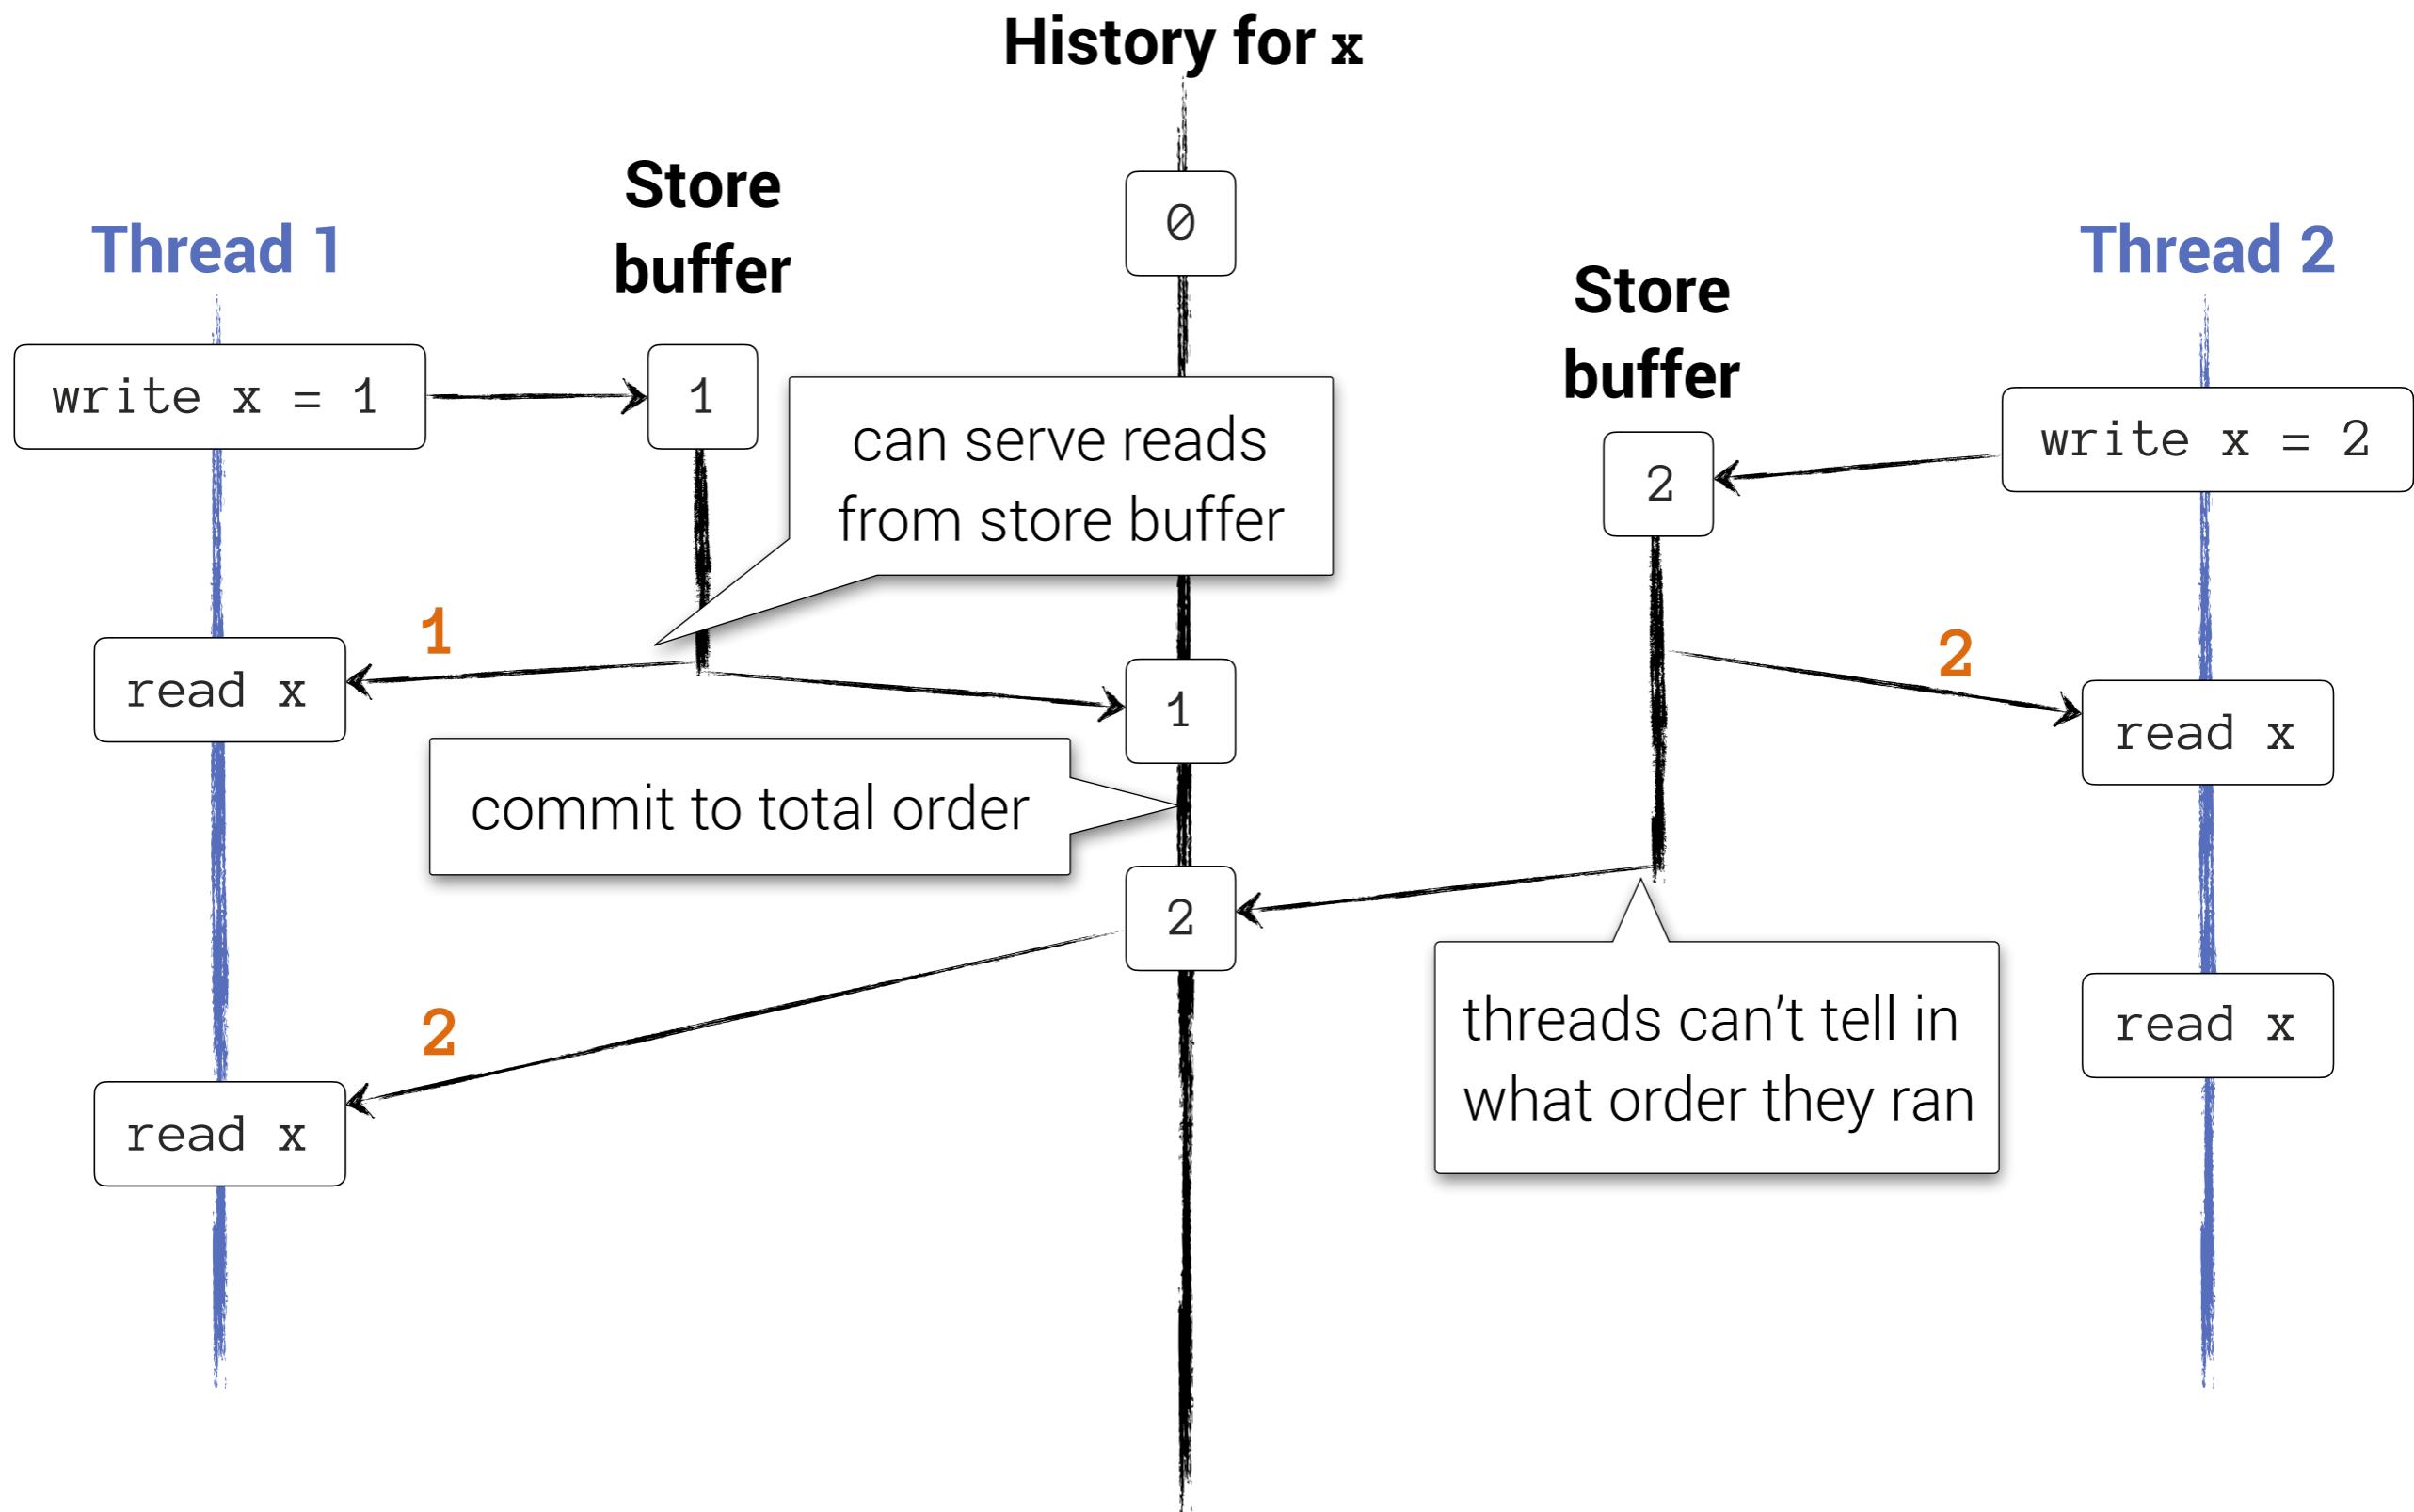
Linearizable



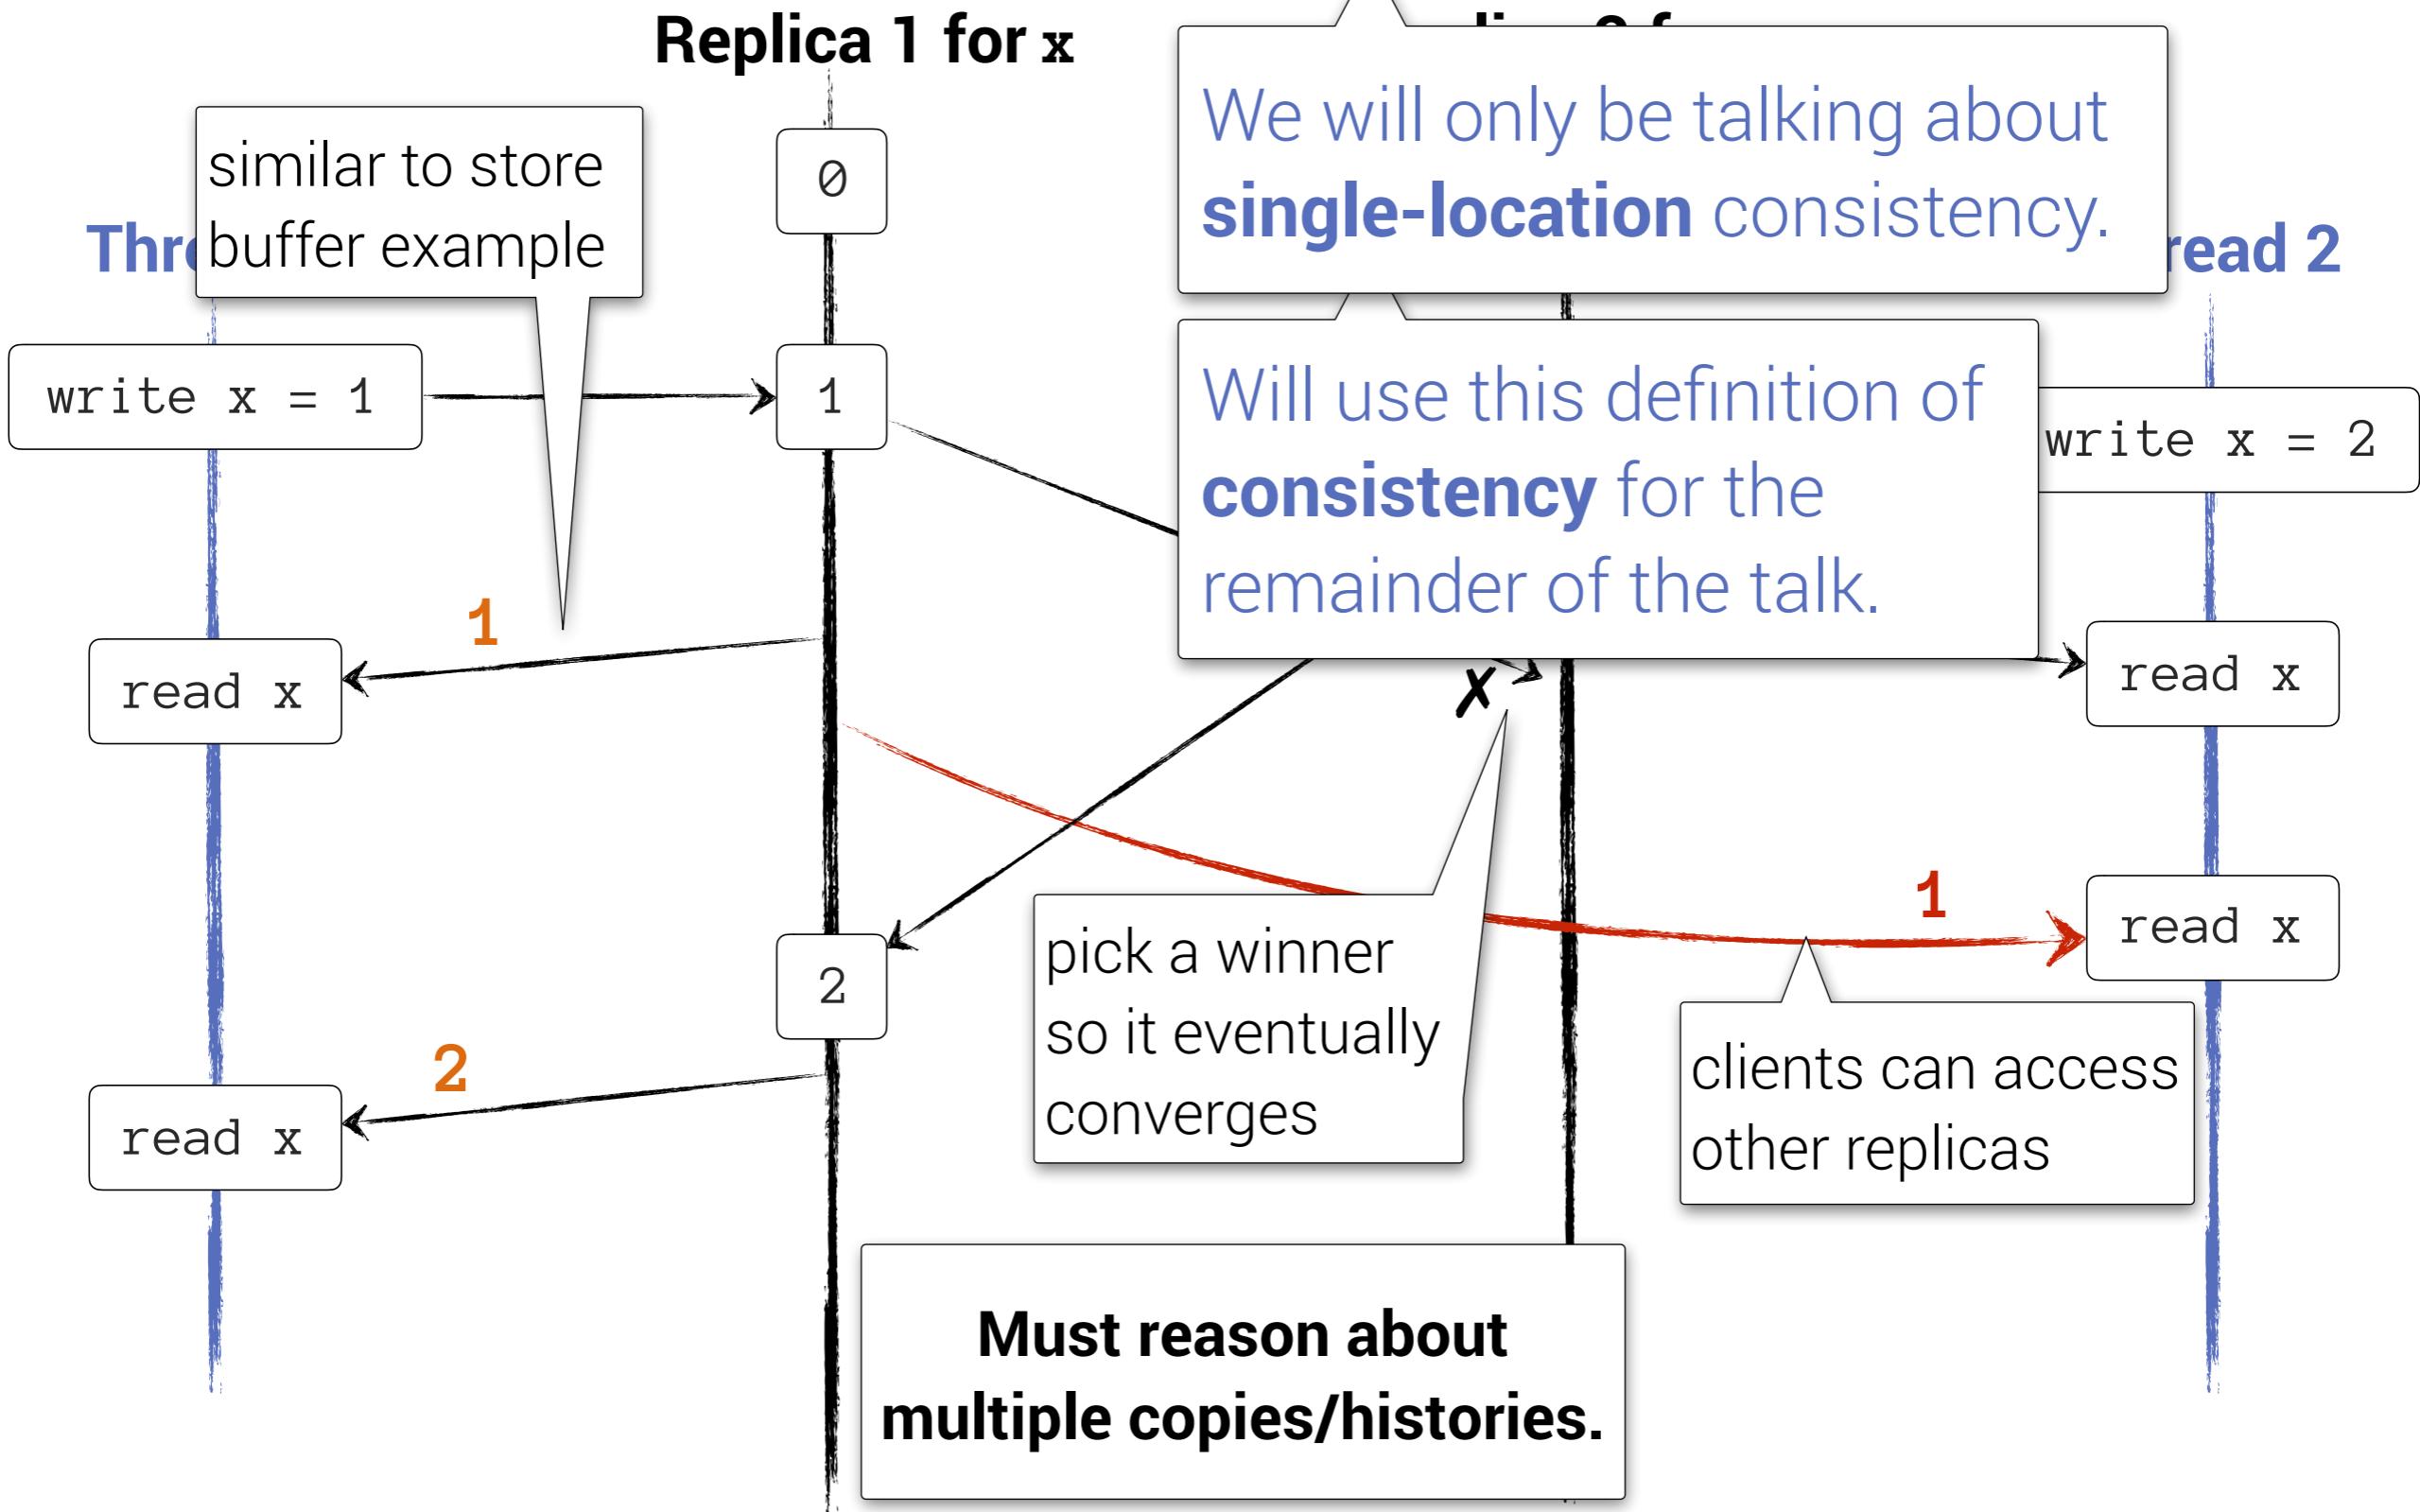
Linearizable



Linearizable



Weak consistency



Introduction: Contention



- ✓ Sources of contention.
- ✓ 3 broad approaches to mitigating contention.

Background: Trading off consistency



- ✓ Consistency for architects.
- Ordering and visibility constraints.
- Dealing with uncertainty.
- Programming with tradeoffs.

Proposal: Disciplined Inconsistency



- Make tradeoffs *explicit* and *safe* with types.
- Reason about *values* rather than *ordering*.
- Synthesize synchronization/enforcement logic.

Eventual Consistency

- Only guarantee: if updates stop, replicas will eventually reflect the same state.
- Fastest response time.
- *High availability, lowest latency*

Strong Consistency

- Herlihy's *linearizability* (only talking about single-location)
- Total global order
- Not highly available

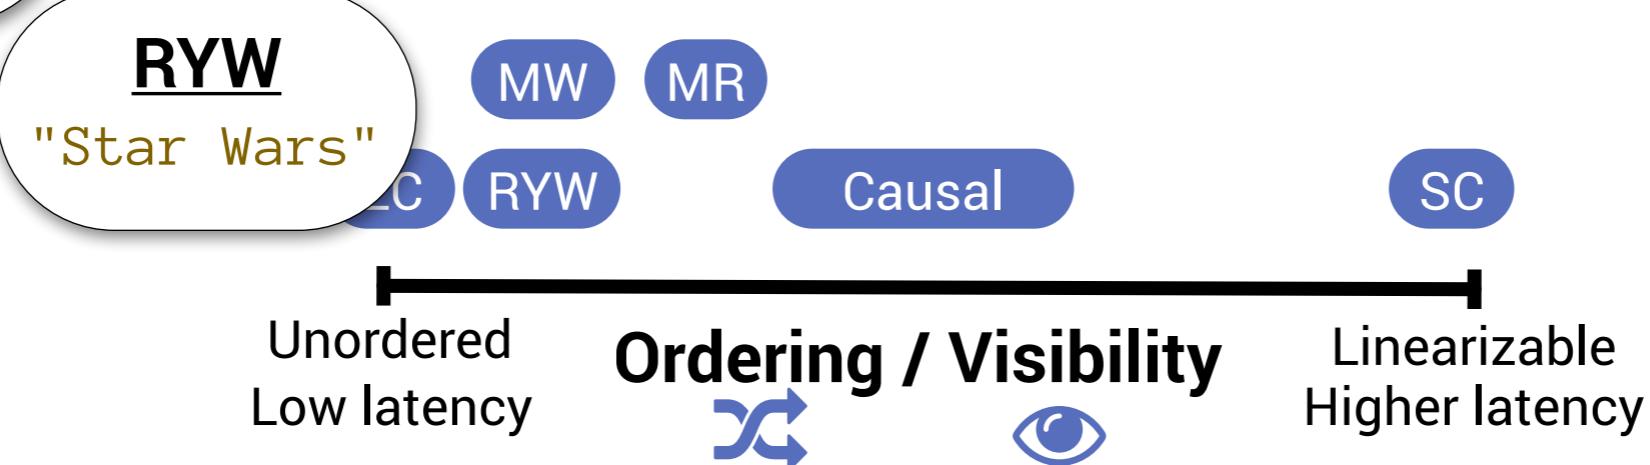


Other weak consistency

- Force coordination in certain cases, or restrict which replicas clients can talk to (e.g. "sticky sessions")
- But, easier to reason about:

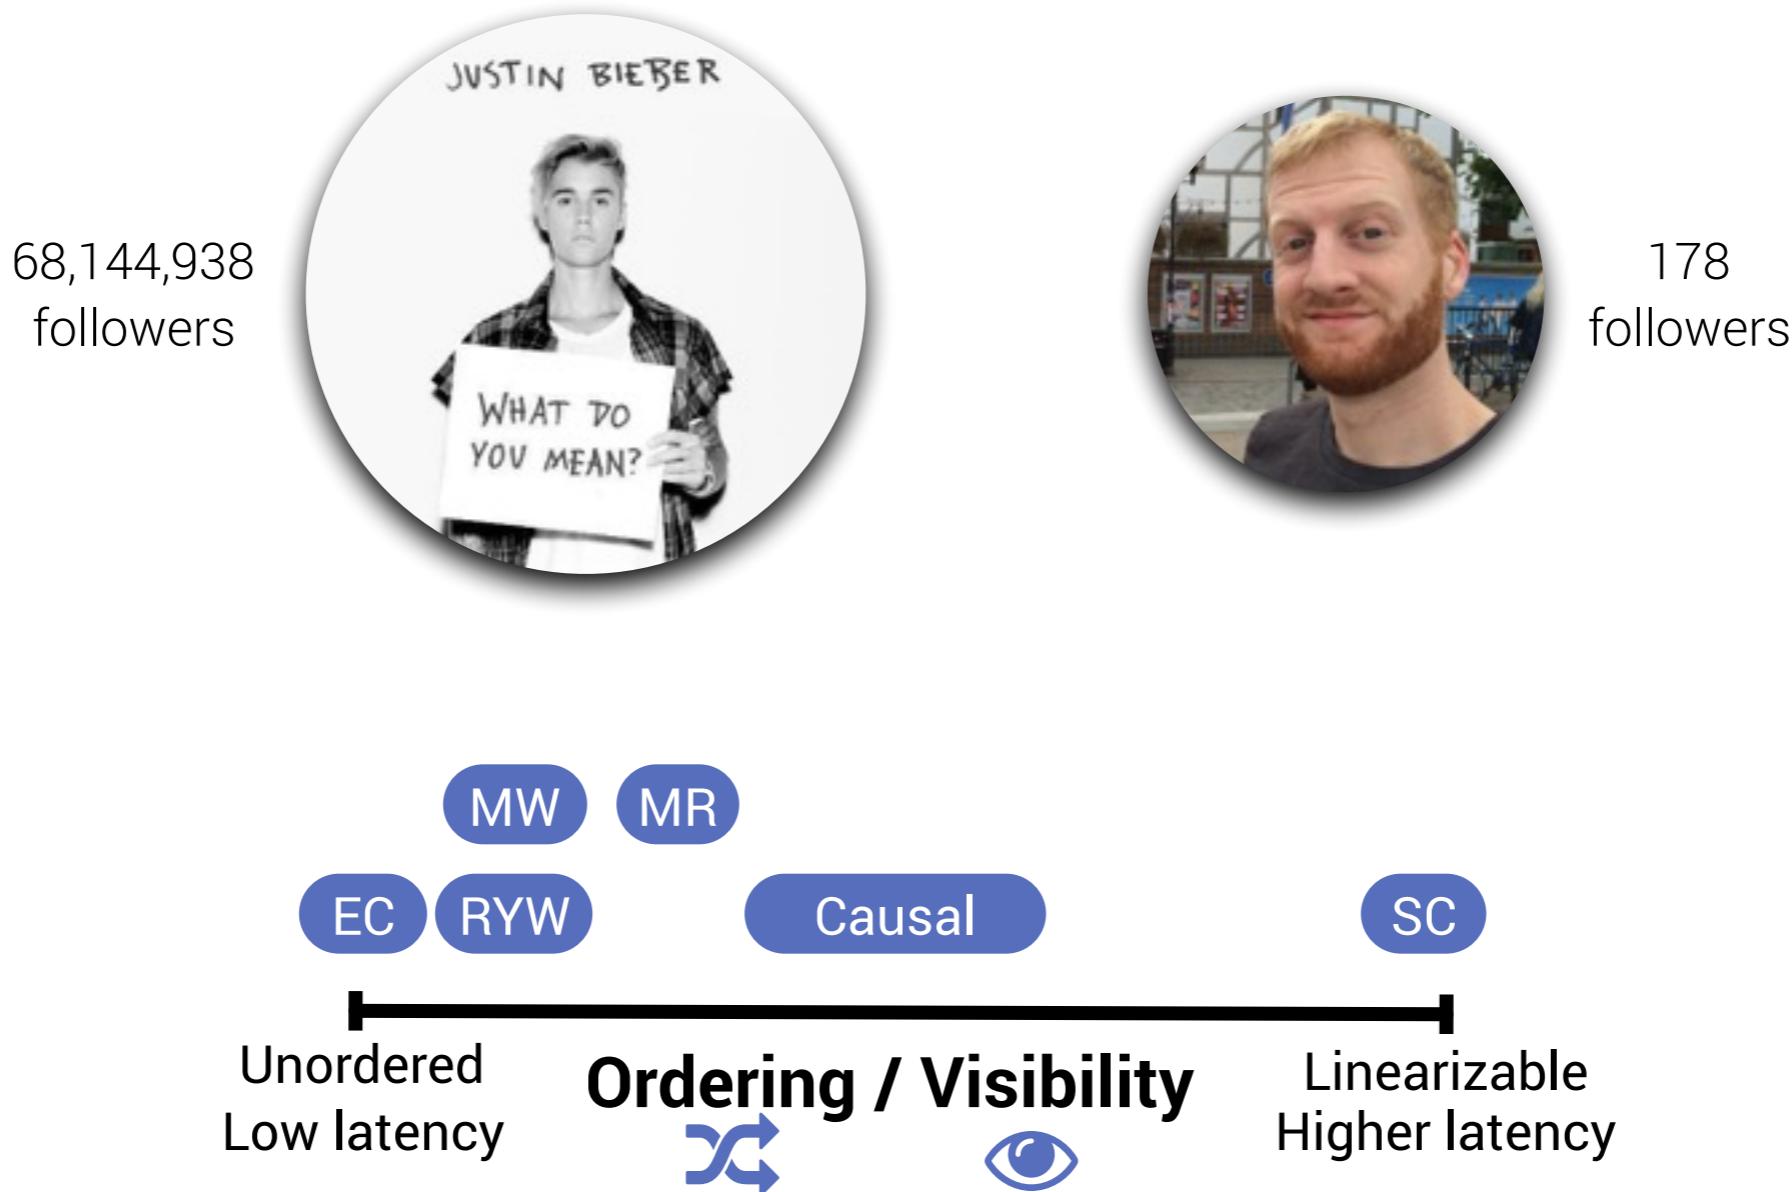
```
# put ticket in my cart  
put("bholt:cart", "Star Wars")  
# view my cart  
cart = get("bholt:cart")
```

EC
(empty)
"Star Wars"



Trading off consistency ➤ Ordering constraints

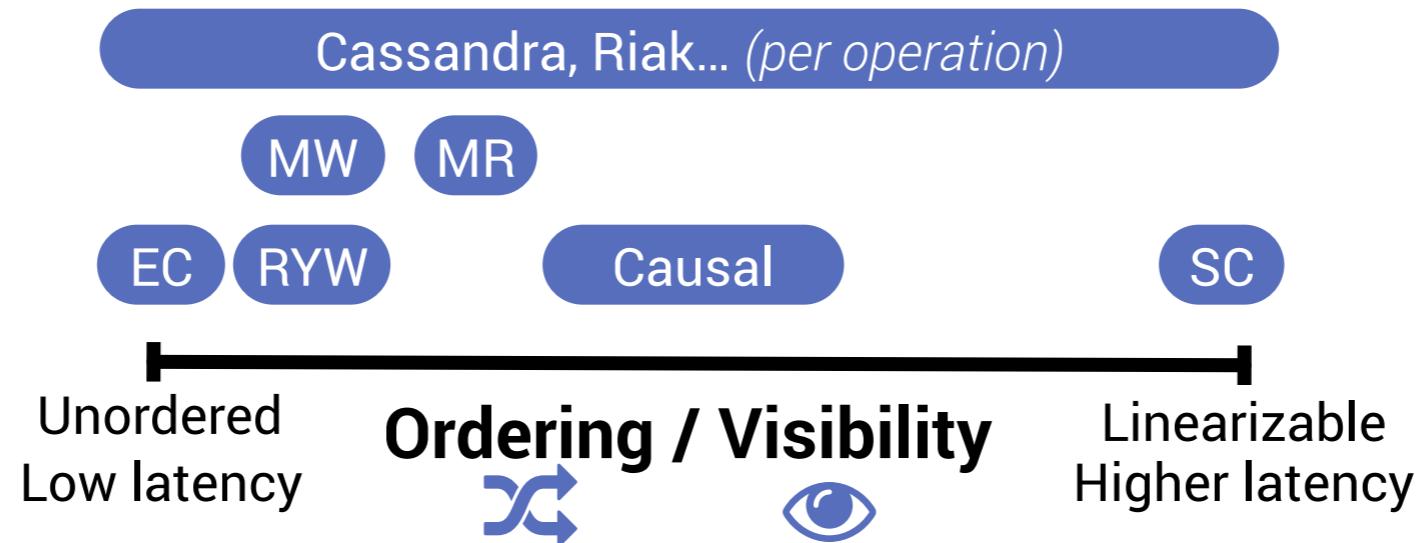
- ⚡ Consistency models are **too coarse-grained**.
- ⚡ Only enforce constraints where **necessary**.



think about a better example? would be better to have something about different parts of app, since that's more typical.

Per-Operation Consistency Level

- Supported by many datastores (e.g. Cassandra, Riak...)
- Only use stronger consistency where necessary.
- Selecting the right one is *error-prone, non-modular*.



add example of difficulty of using multiple consistency levels? get one from Quelea paper.

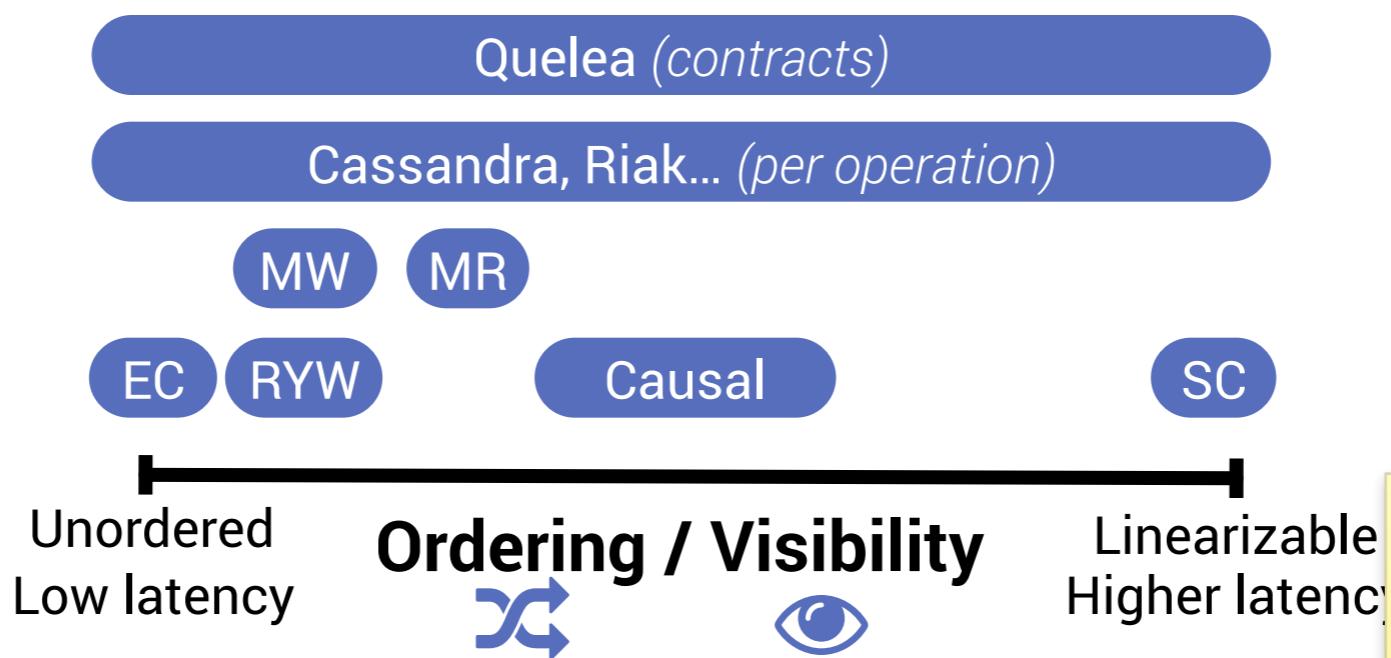
Quelea

- Explicit ordering & visibility contracts:

$$\forall (a : \text{sellTicket}) : \text{sameobj}(a, \eta) \Rightarrow a = \eta \vee \text{vis}(a, \eta) \vee \text{vis}(\eta, a)$$

- Automatically choose correct consistency level.
- But low-level constraints not much better.

*Declarative programming over eventually consistent data stores. (PLDI'15)
KC Sivaramakrishnan, Gowtham Kaki, and Suresh Jagannathan.*



Trading off consistency ➤ Ordering constraints

Indigo

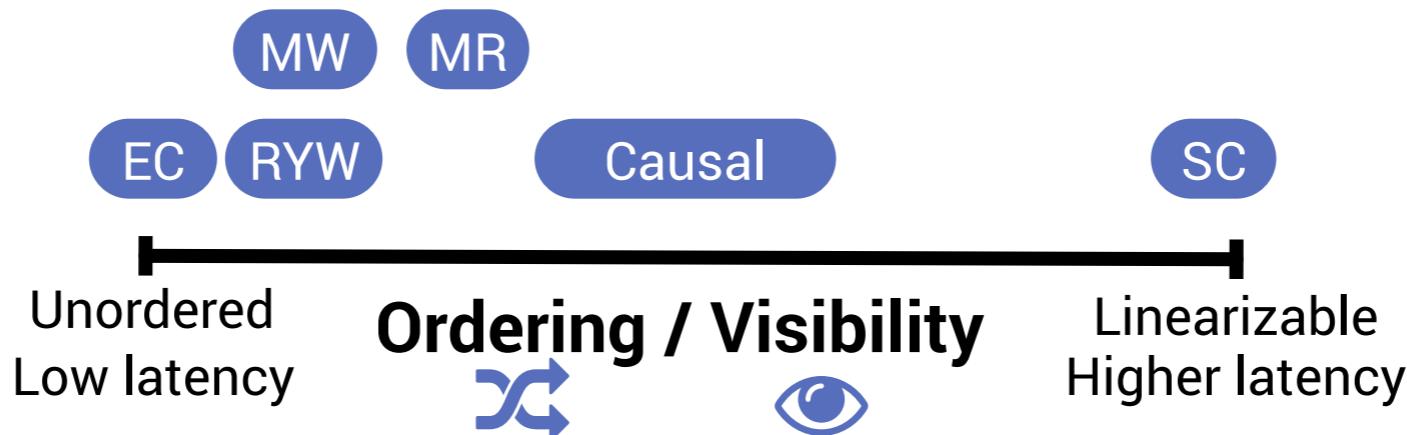
- Annotate *application-level invariants*:

```
@Invariant("forall(Movie : m) :- tickets(m) >= 0")
public interface MovieTickets {
    @PostCondition-Decrements("tickets(m, 1)")
    void sellTicket(Movie m);
}
```

- Synthesize coordination logic (*reservations*).
- But: manually annotate abstract state & behavior.

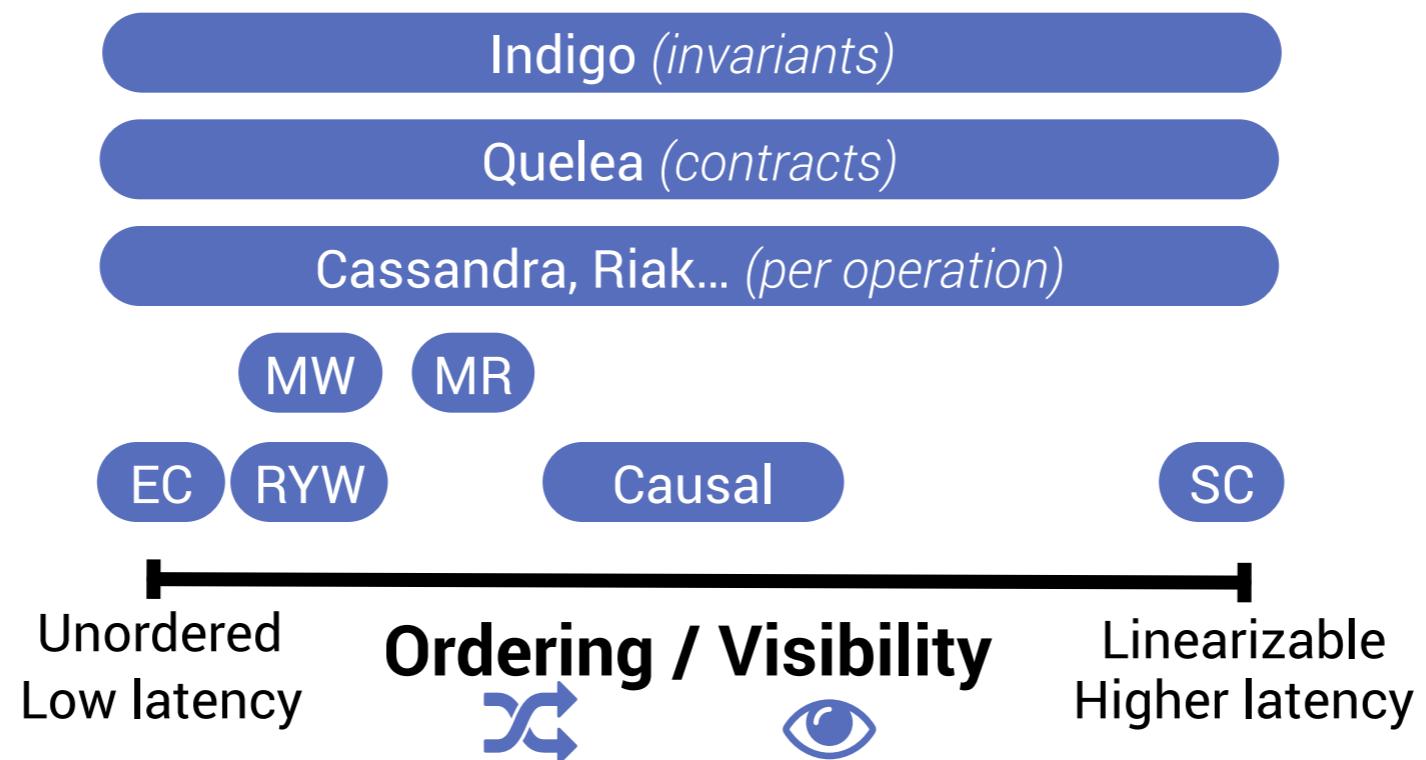
Putting consistency back into eventual consistency. (EuroSys'15)
V. Balegas, ... , M. Shapiro.

Cassandra, Riak... (per operation)



Trading off consistency ➤ Ordering constraints

- ⚡ Consistency models are **too coarse-grained**.
- ⚡ Only enforce constraints where **necessary**.
- ⚡ **Synthesize fine-grained synchronization.**



Introduction: Contention



- ✓ Sources of contention.
- ✓ 3 broad approaches to mitigating contention.

Background: Trading off consistency



- ✓ Consistency for architects.
- ✓ Ordering and visibility constraints.
- Dealing with uncertainty: **convergence vs. staleness**
- Programming with tradeoffs.



Proposal: Disciplined Inconsistency

- Make tradeoffs *explicit* and *safe* with types.
- Reason about *values* rather than *ordering*.
- Synthesize synchronization/enforcement logic.

Trading off consistency → Uncertainty

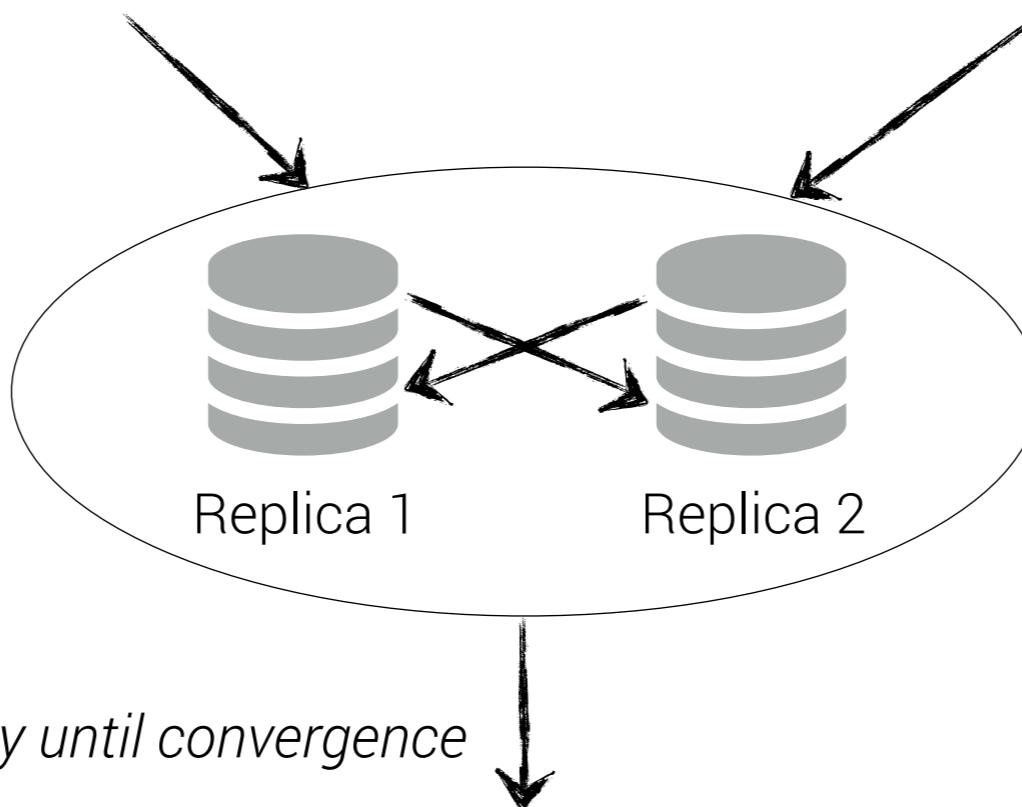
Converged states

Client 1

```
x = read("people")  
write("people", x+"Alice")
```

Client 2

```
x = read("people")  
write("people", x+"Bob")
```



wait indefinitely until convergence

Client 3

```
read("people")
```

uncertainty
in converged
states

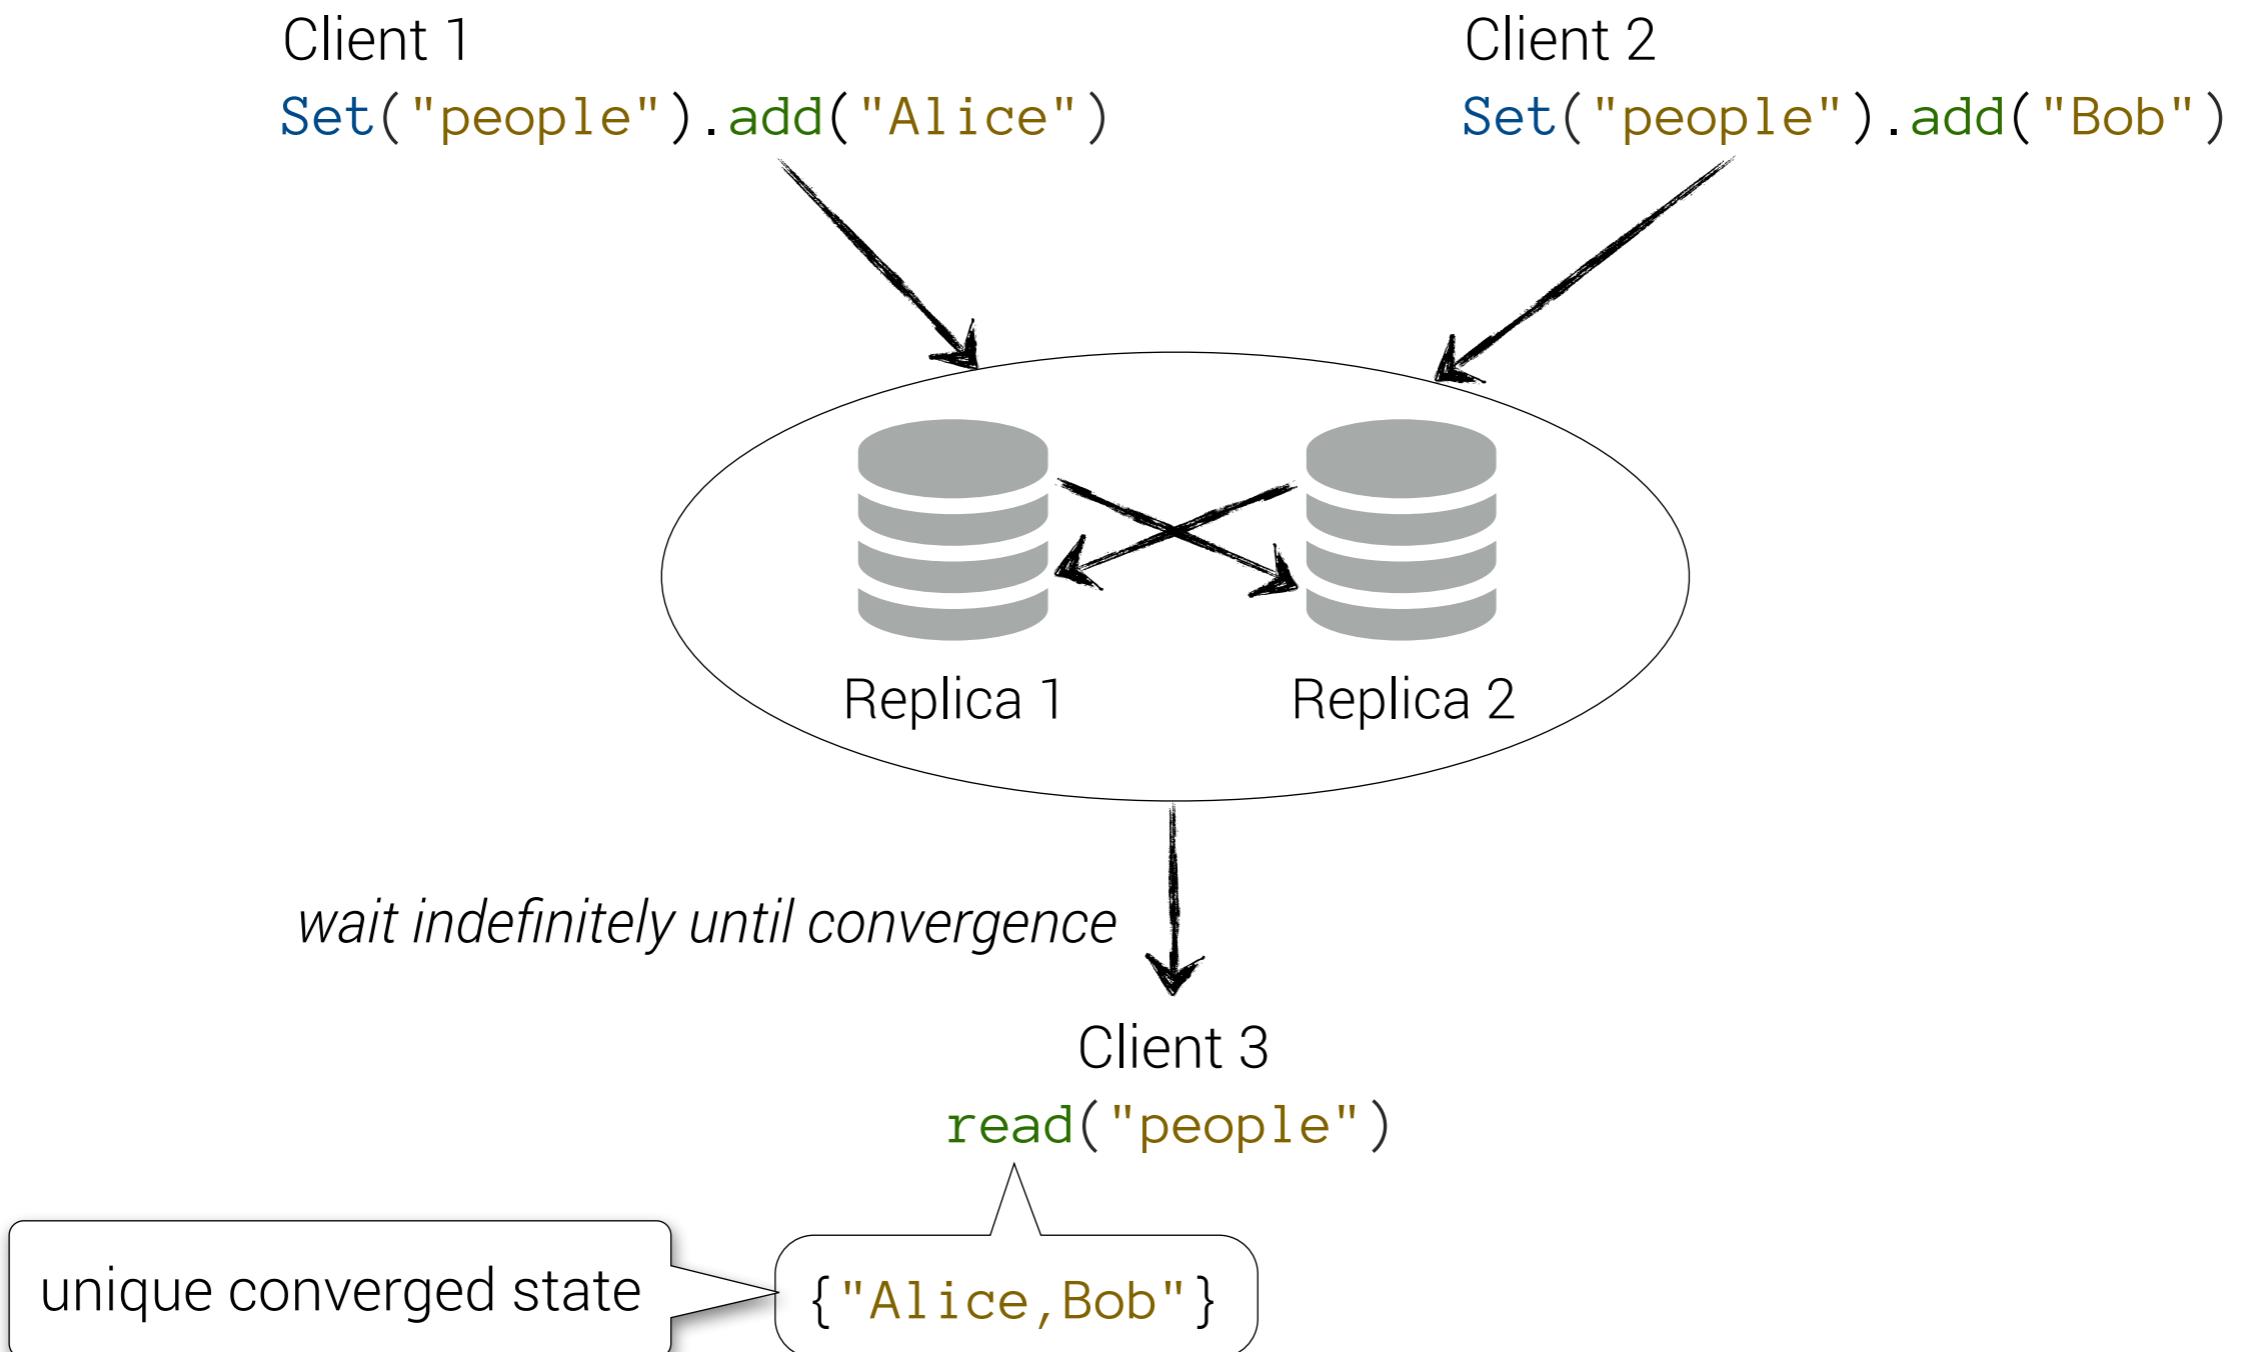
"Alice"
"Bob"
"Alice, Bob"
"Bob, Alice"

missed update

commutative

Trading off consistency → Uncertainty

Converged states



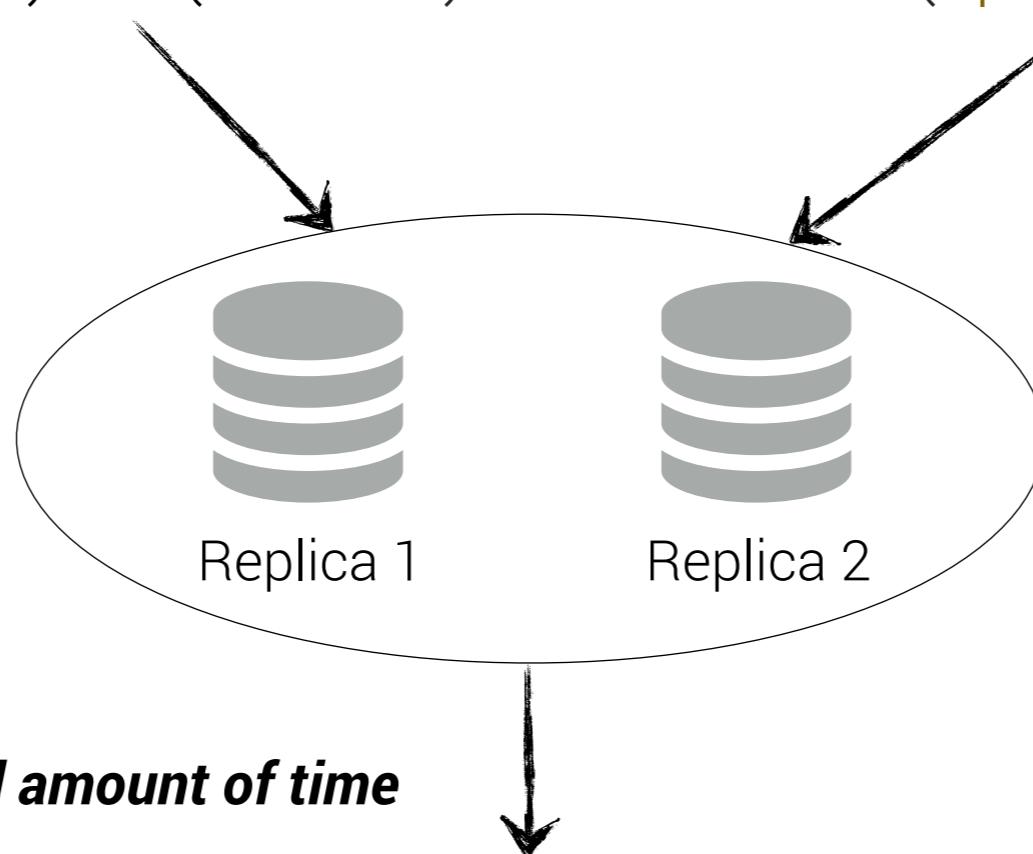
Trading off consistency

Uncertainty

Staleness

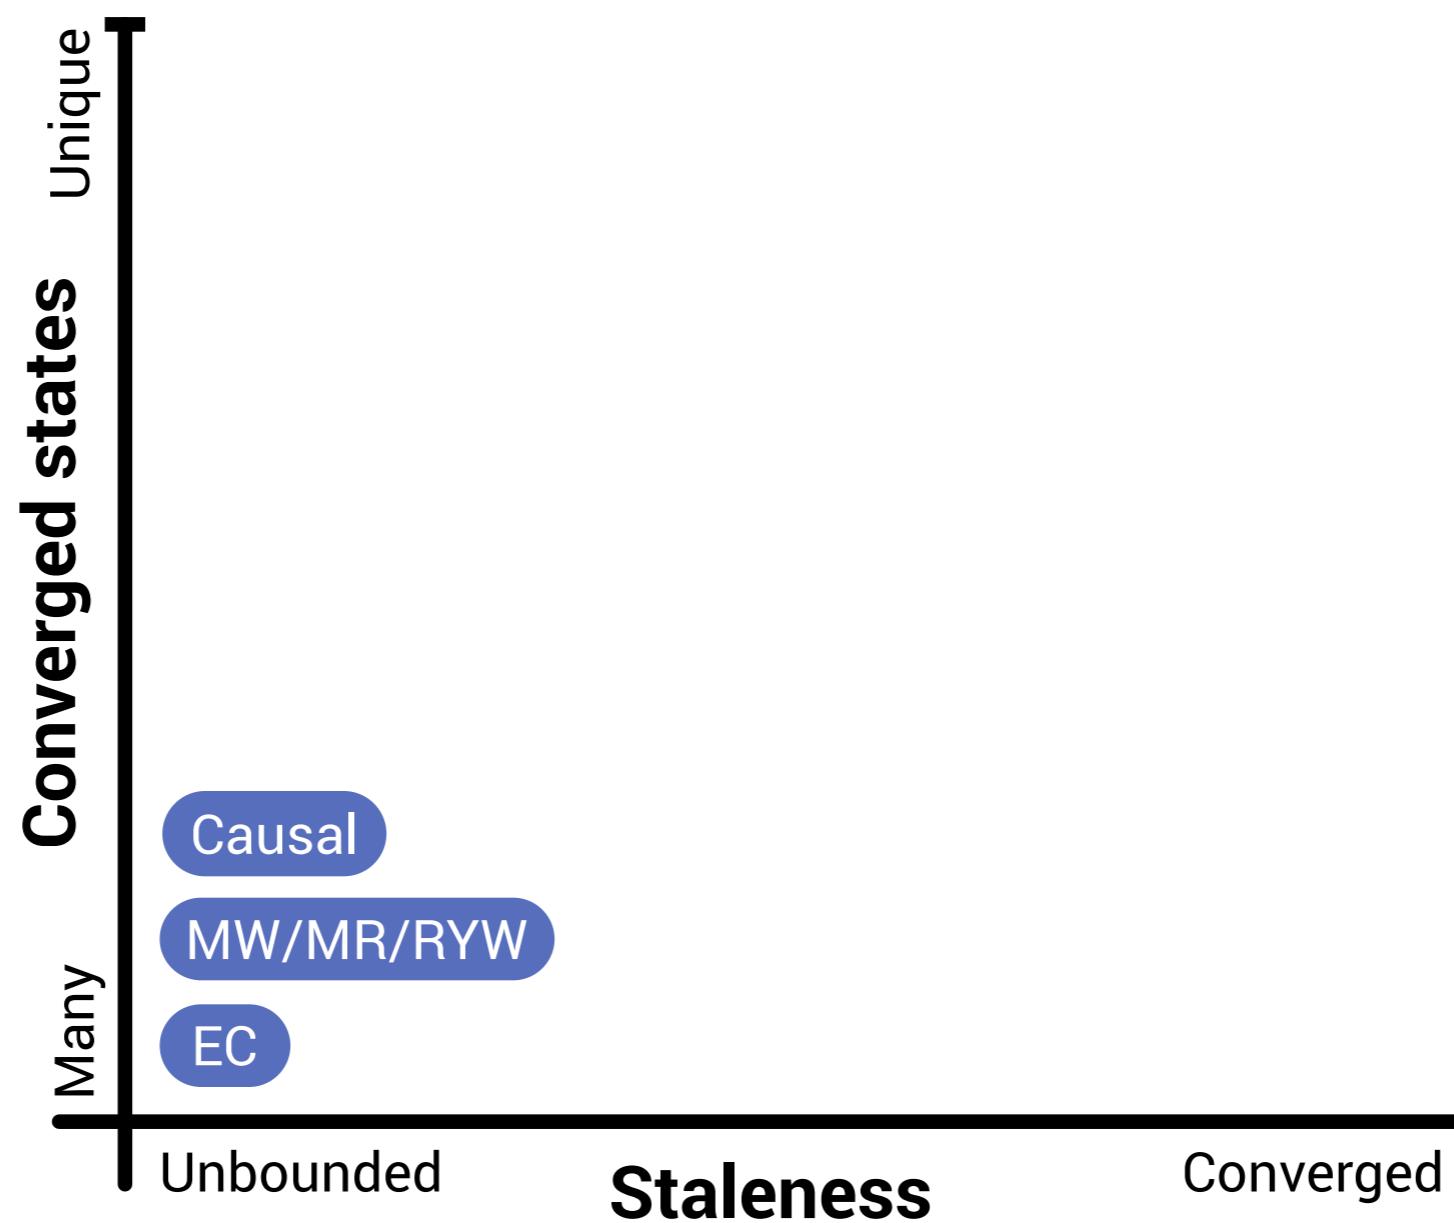
Client 1
`Set("people").add("Alice")`

Client 2
`Set("people").add("Bob")`

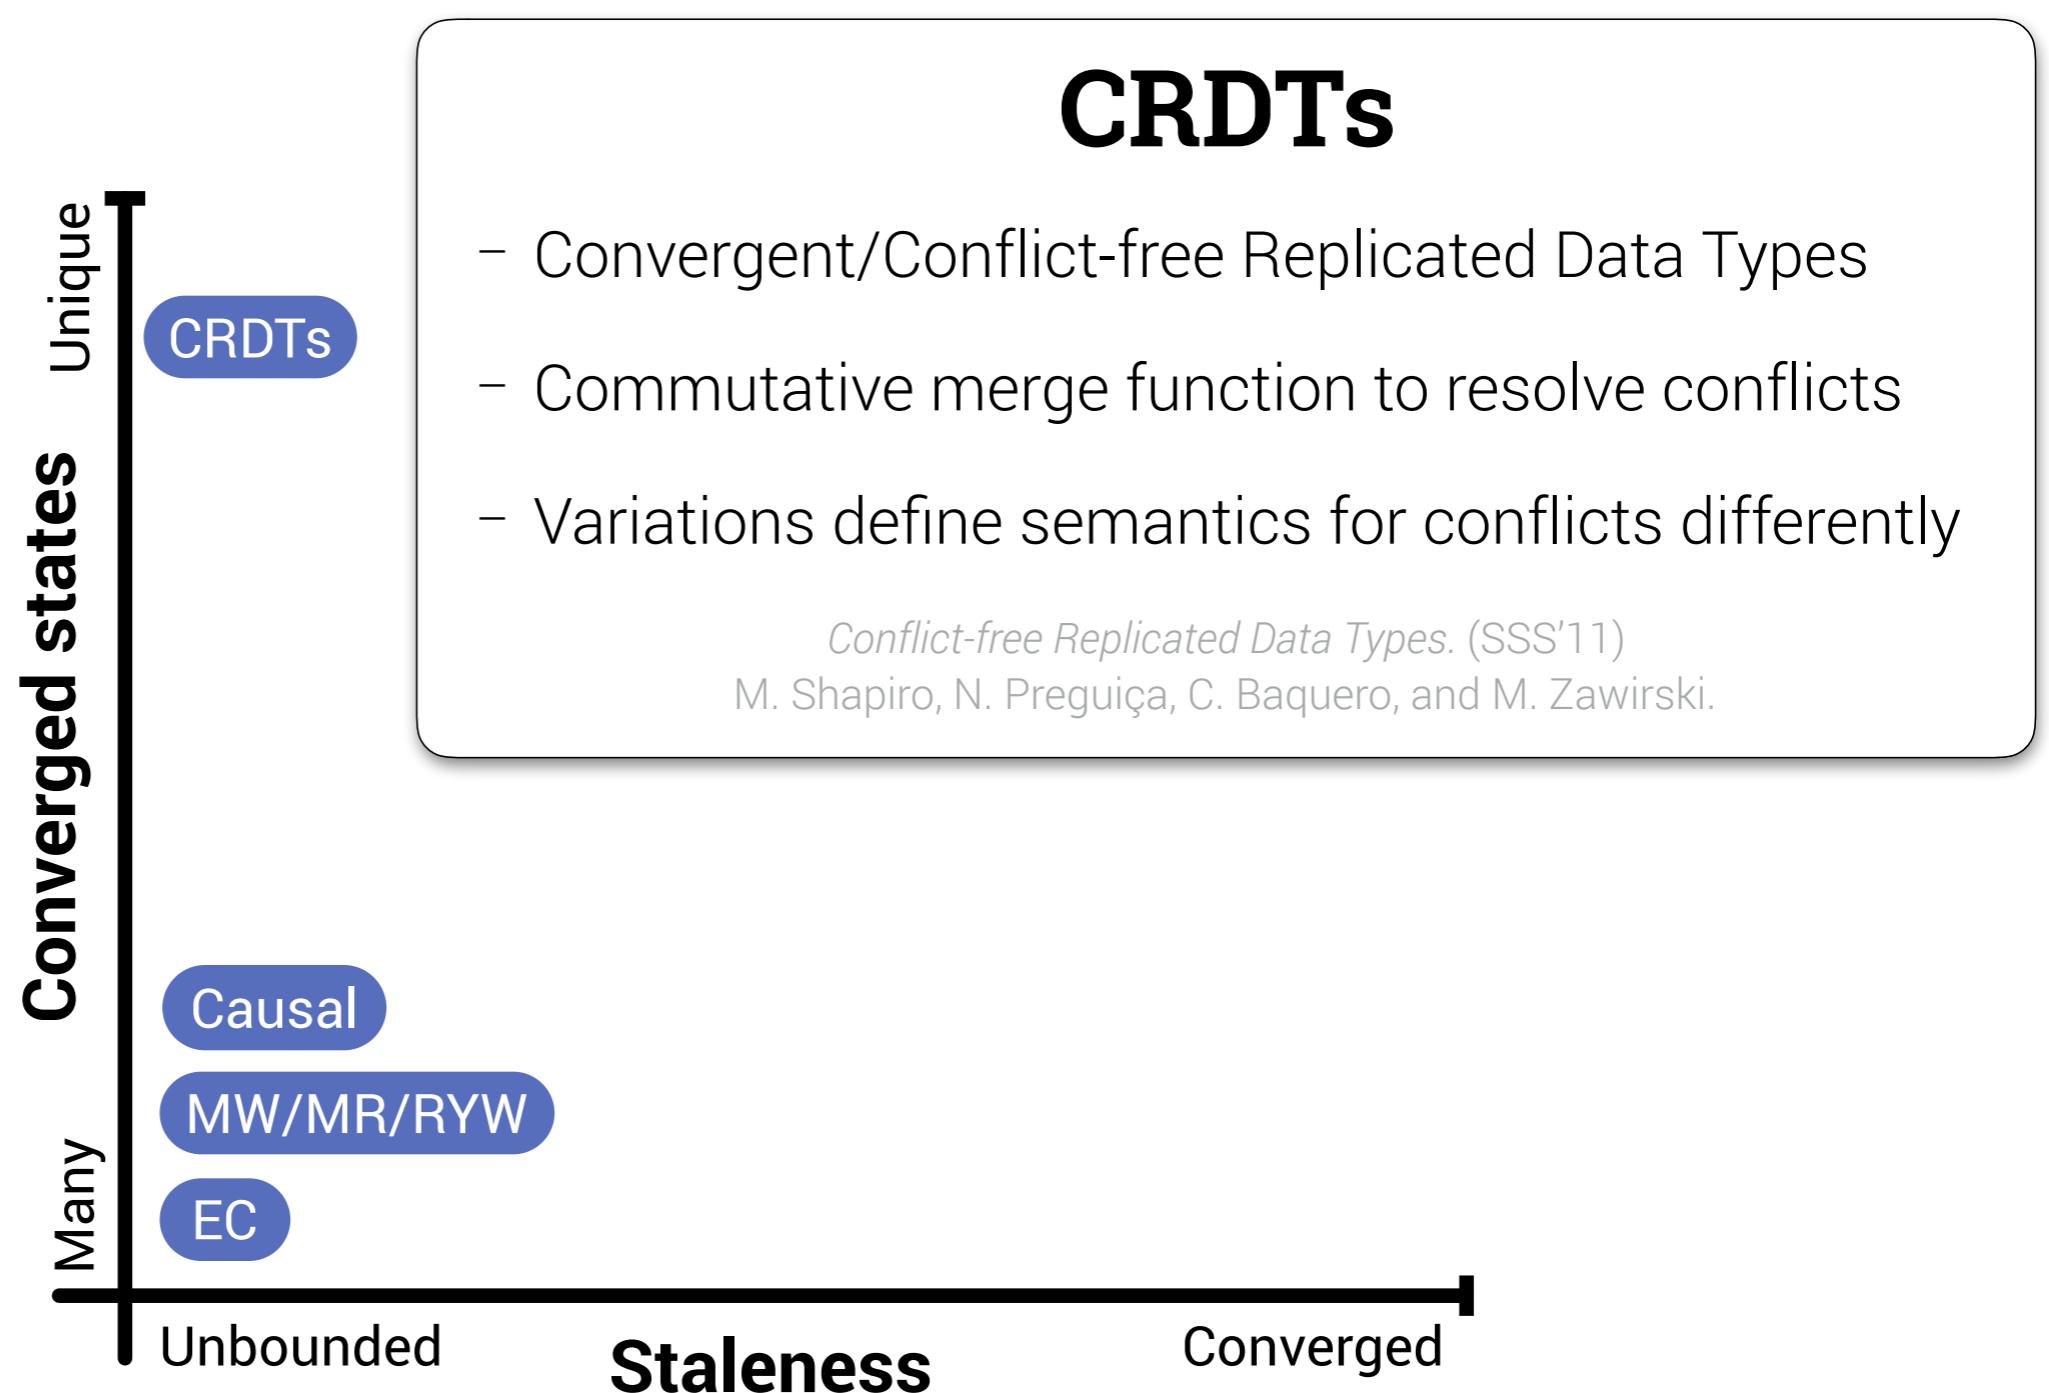


staleness uncertainty

Trading off consistency → Uncertainty



Trading off consistency → Uncertainty



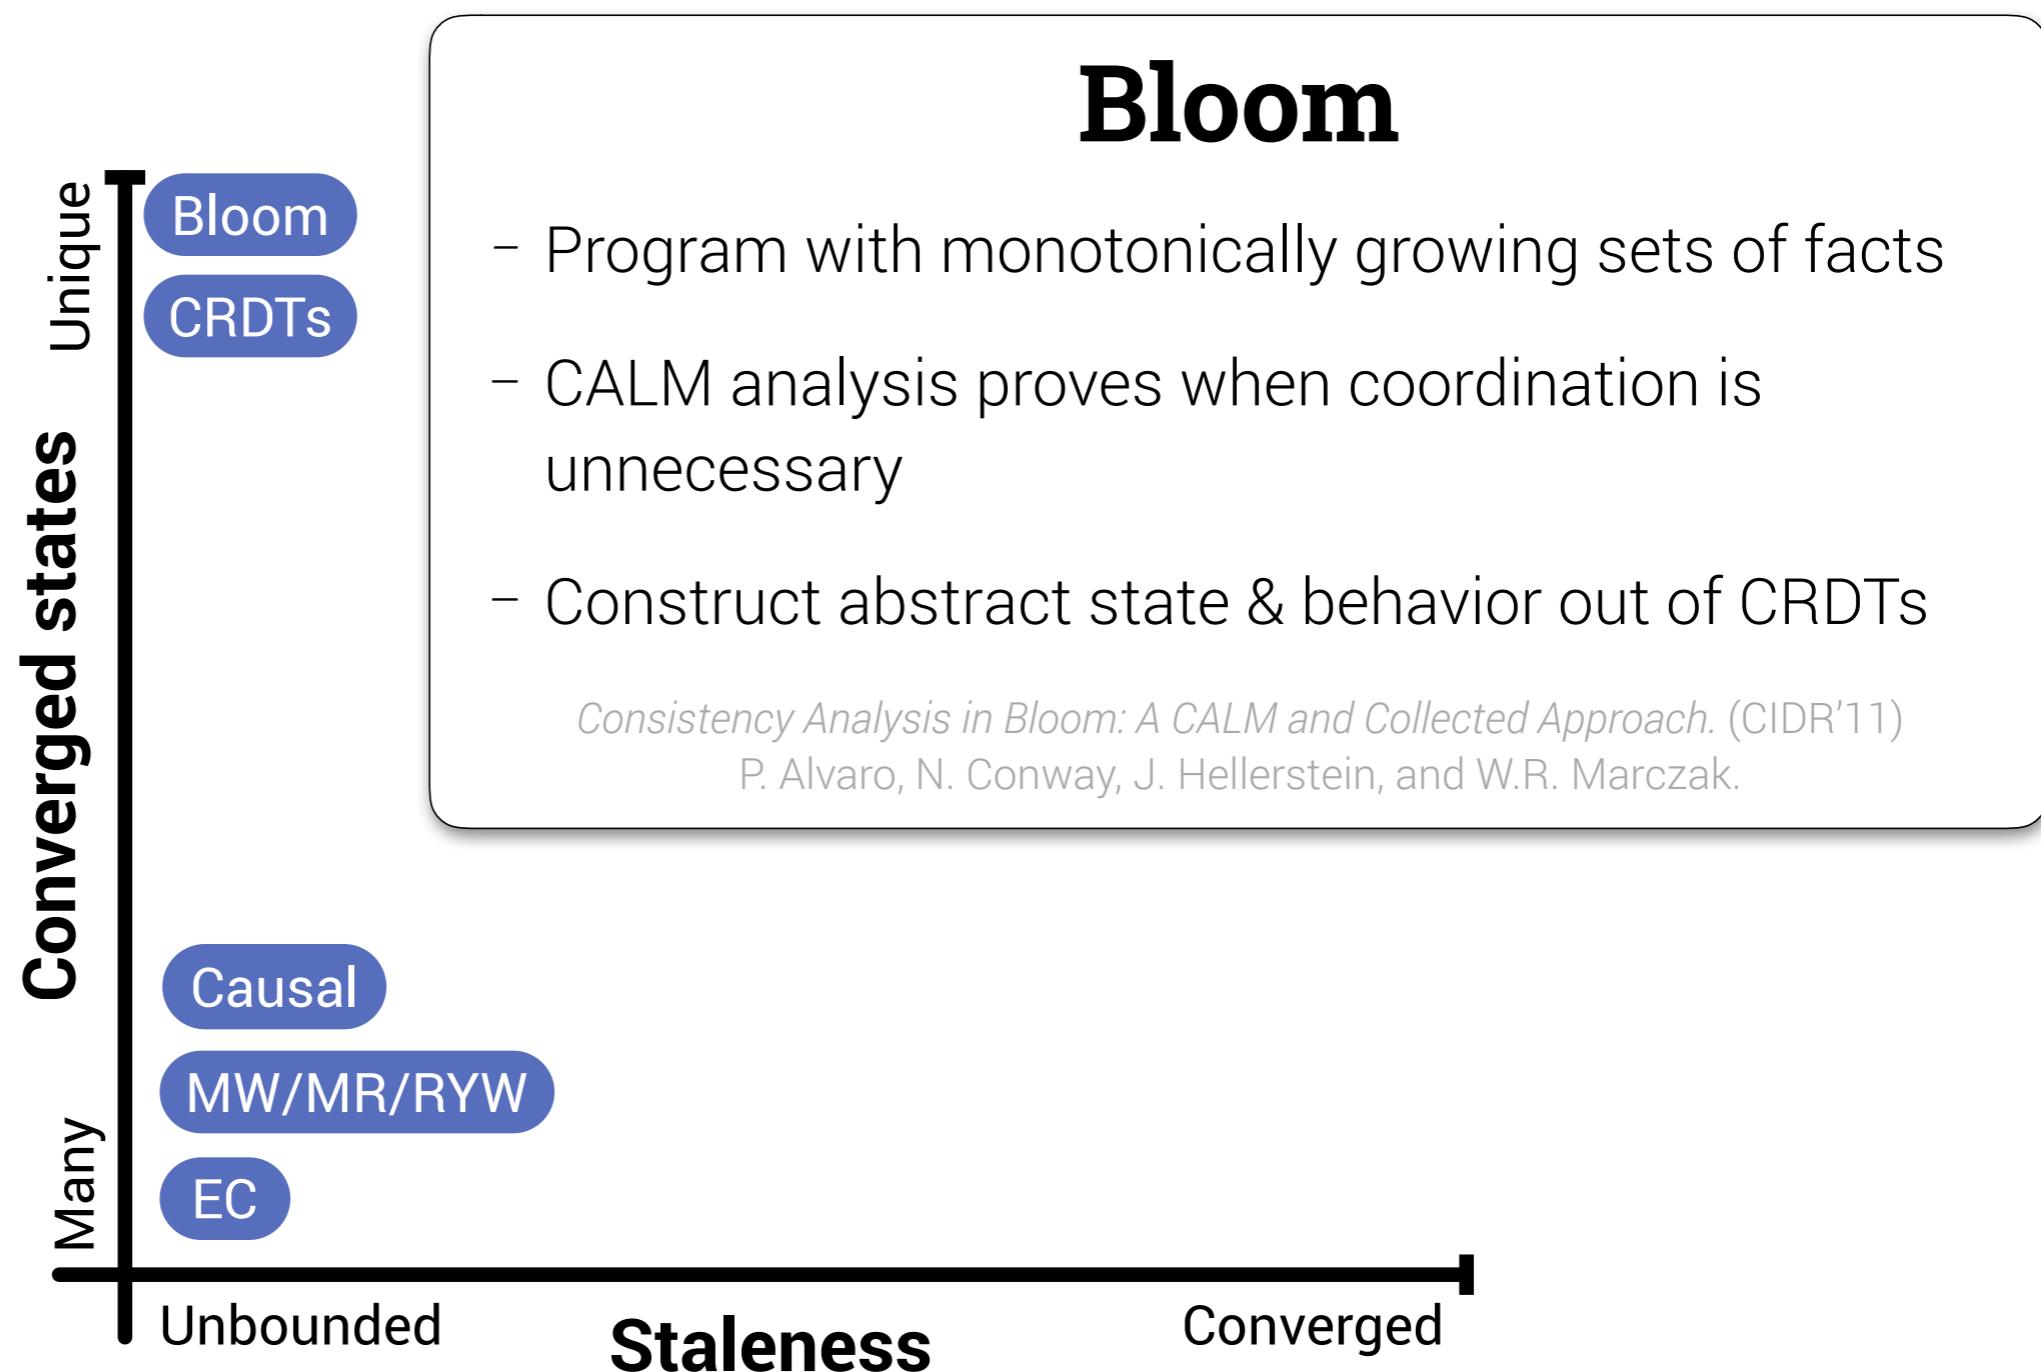
CRDTs

- Convergent/Conflict-free Replicated Data Types
- Commutative merge function to resolve conflicts
- Variations define semantics for conflicts differently

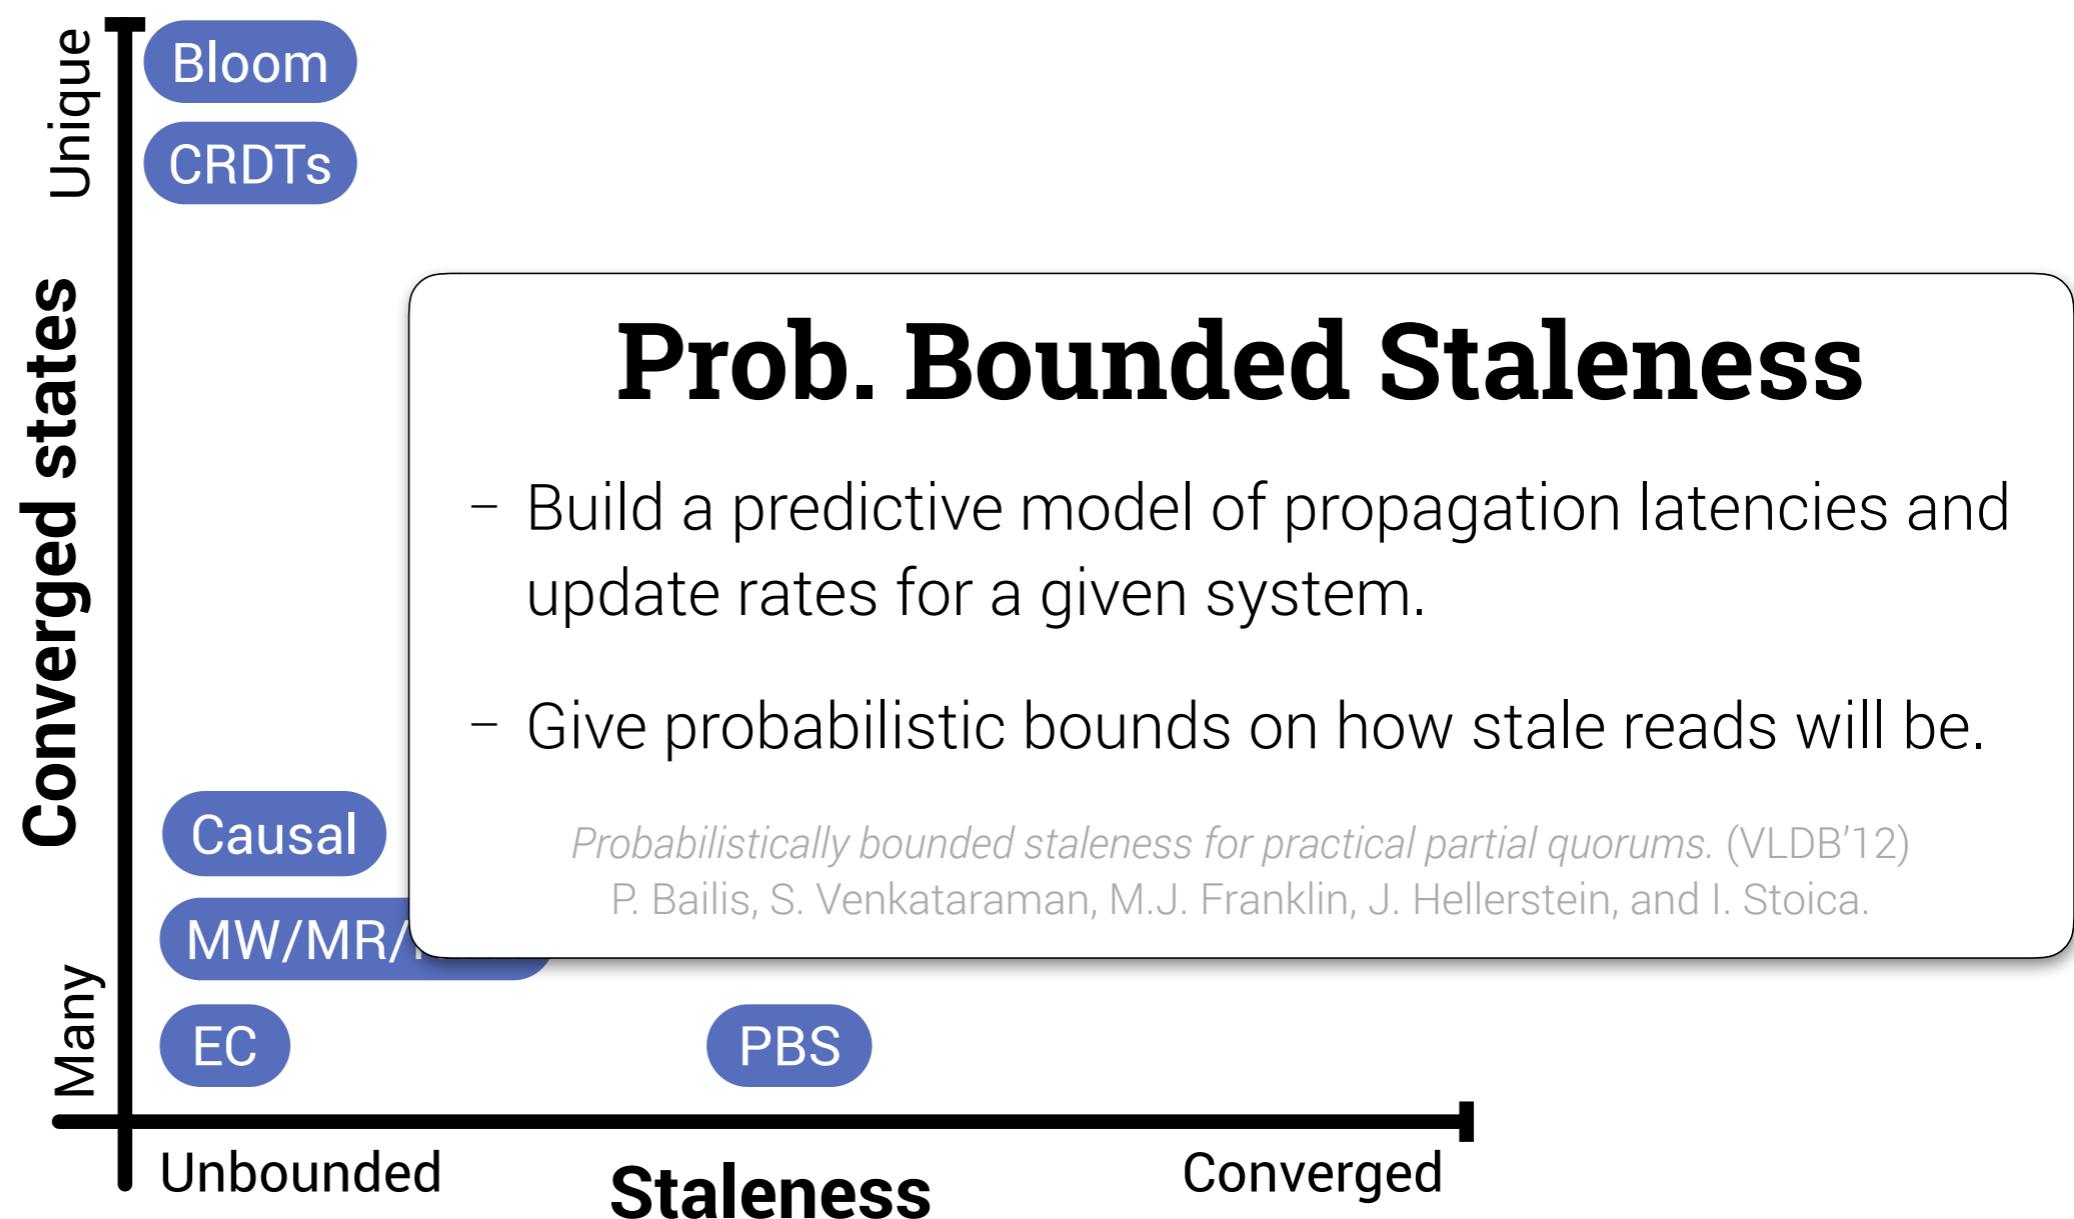
Conflict-free Replicated Data Types. (SSS'11)

M. Shapiro, N. Preguiça, C. Baquero, and M. Zawirski.

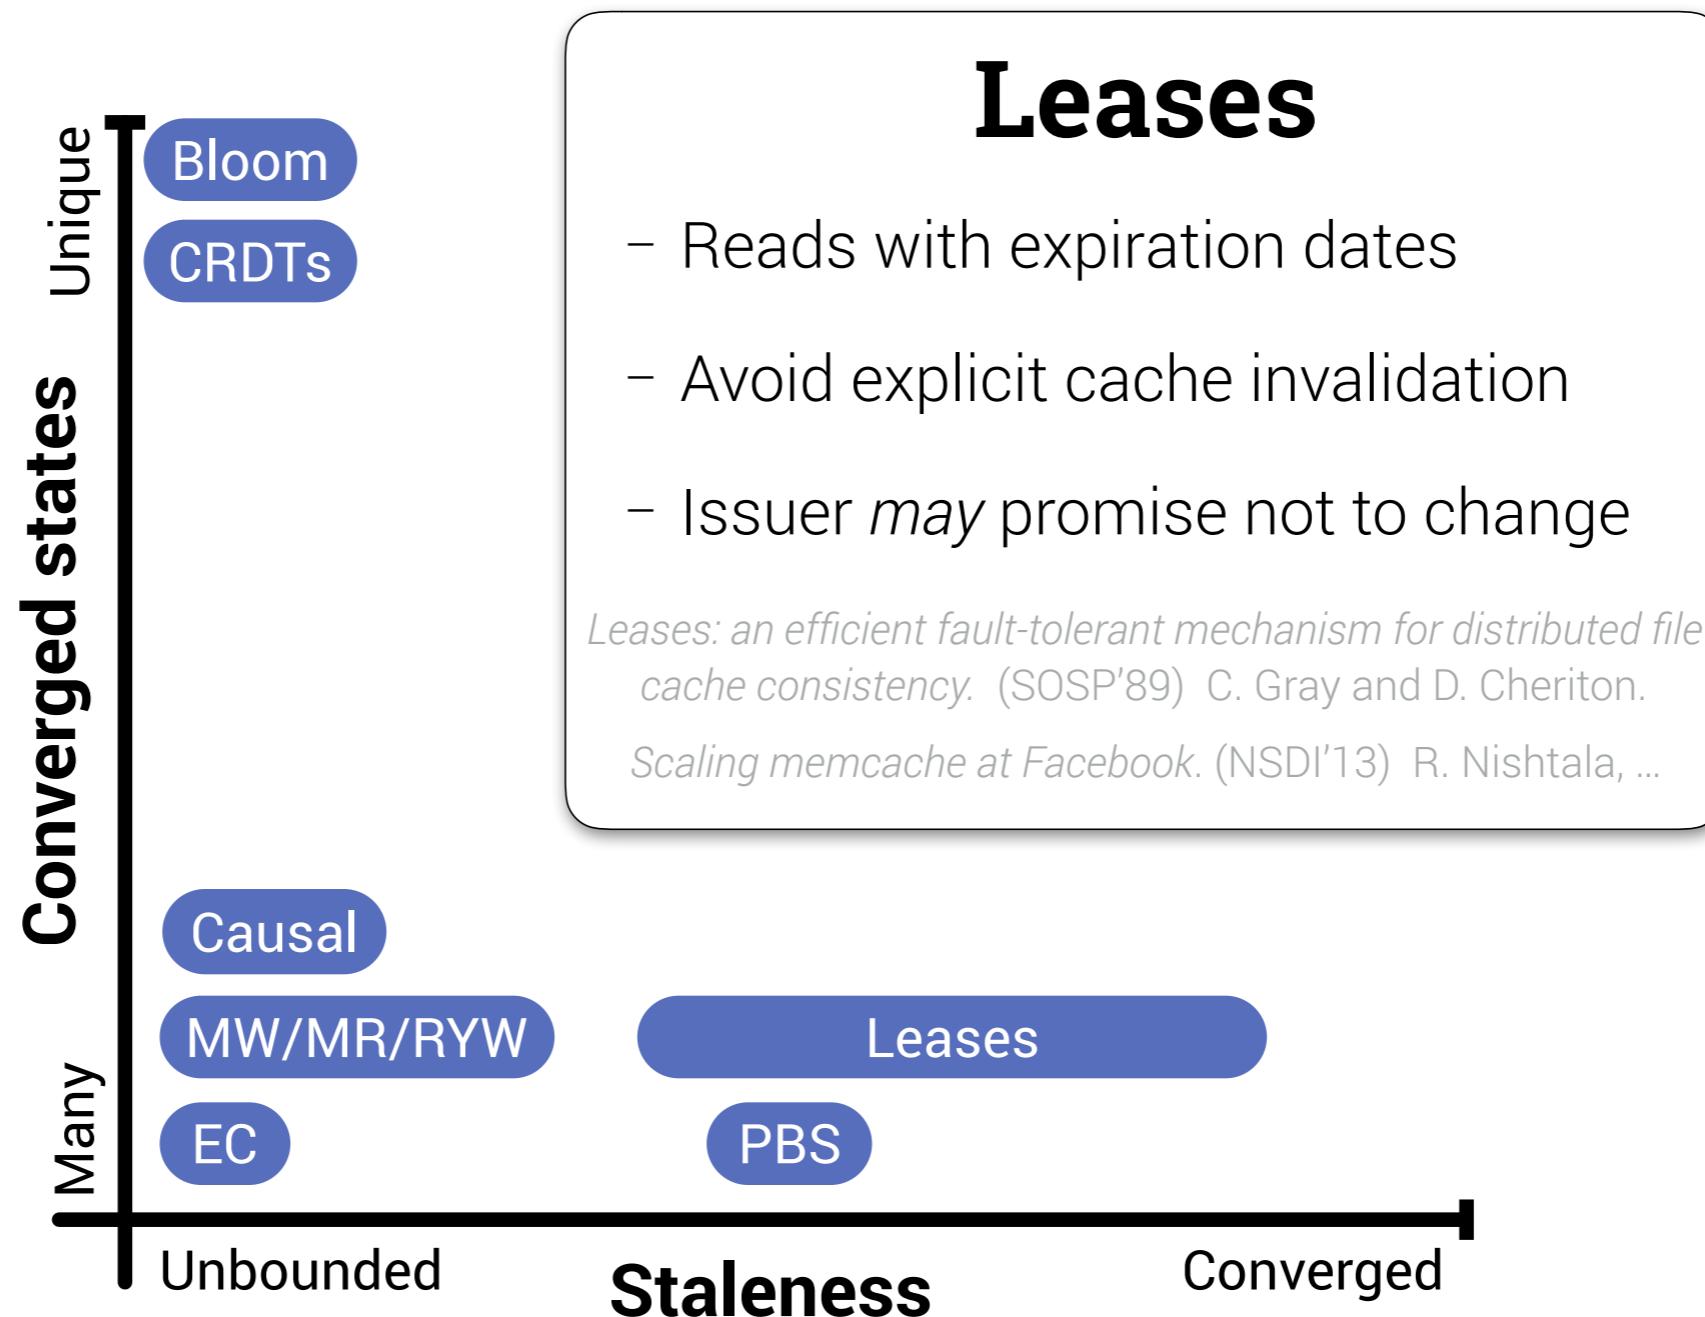
Trading off consistency → Uncertainty



⚡ Specify abstract state and behavior **by-construction**.



Trading off consistency → Uncertainty



Consistency-based SLAs

- System-enforced SLAs in a system called *Pileus*
- SLAs contain *target latency* and *consistency level*
- System monitors replicas, predicts which SLA it can meet.
- Specify SLA on reads, return achieved consistency level.

```
# add ticket to cart
put("bob:cart", "Star Wars Ticket")  
  
# get current cart
cart, consistency = get("bob:cart", MovieTicketSLA)
```

Consistency	Latency	Utility
Strong	100 ms	1.0
Eventual	100 ms	0.5

Movie Ticket SLA

Consistency-based service level agreements for cloud storage. (SOSP'13)
D. Terry, V. Prabhakaran, R. Kotla, M. Balakrishnan, M. Aguilera, and H. Abu-Libdeh.

STATELESS

⚡ Explicit performance **targets, feedback about results.**

Introduction: Contention



- ✓ Sources of contention.
- ✓ 3 broad approaches to mitigating contention.

Background: Trading off consistency



- ✓ Consistency for architects.
 - ✓ Ordering and visibility constraints.
 - ✓ Dealing with uncertainty.
- **Review:** Programming tradeoffs.



Proposal: Disciplined Inconsistency

Make tradeoffs *explicit* and *safe* with types.

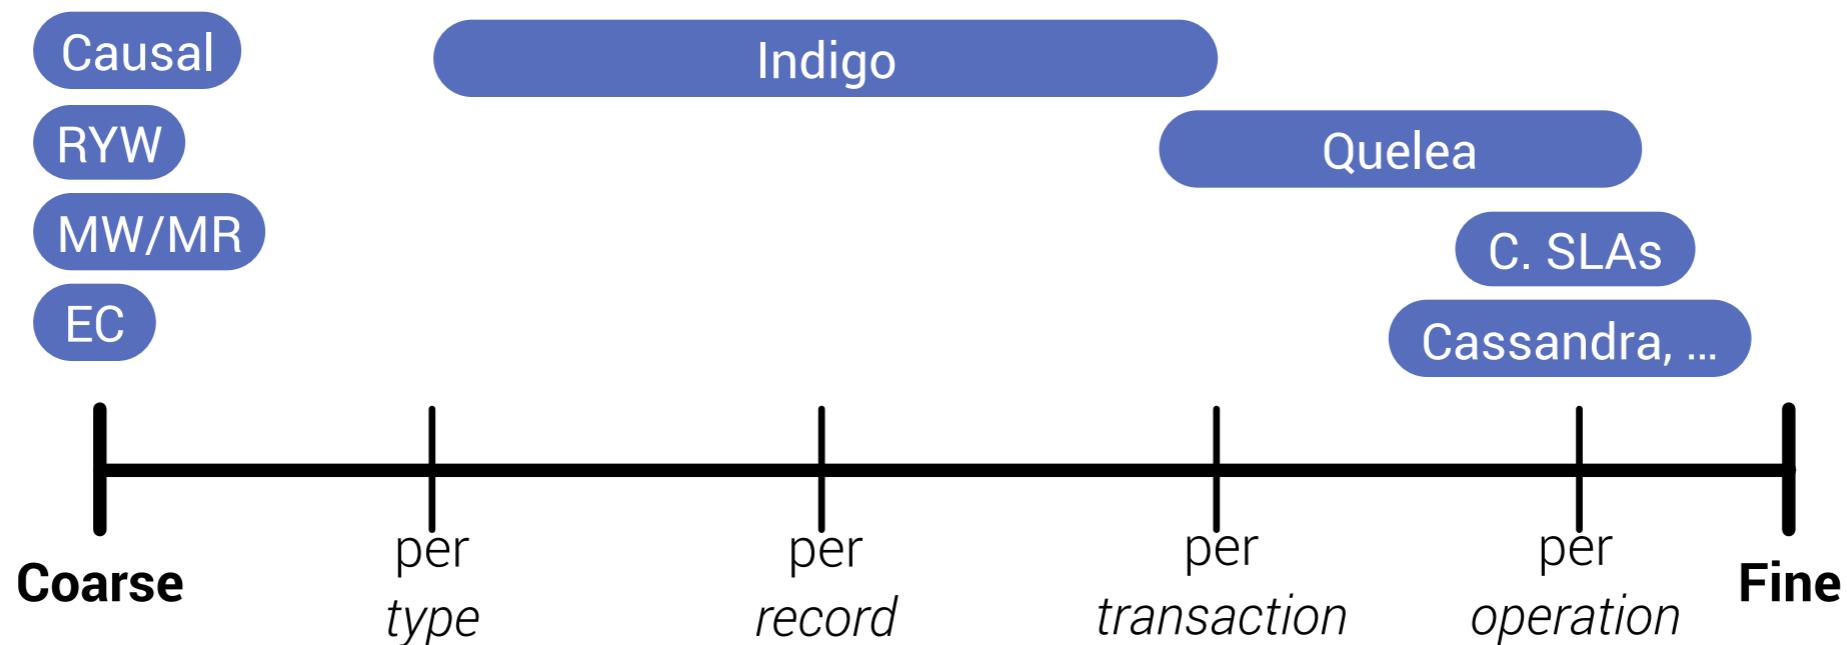
Reason about *values* rather than *ordering*.

Synthesize synchronization/enforcement logic.

Trading off consistency → Programming tradeoffs

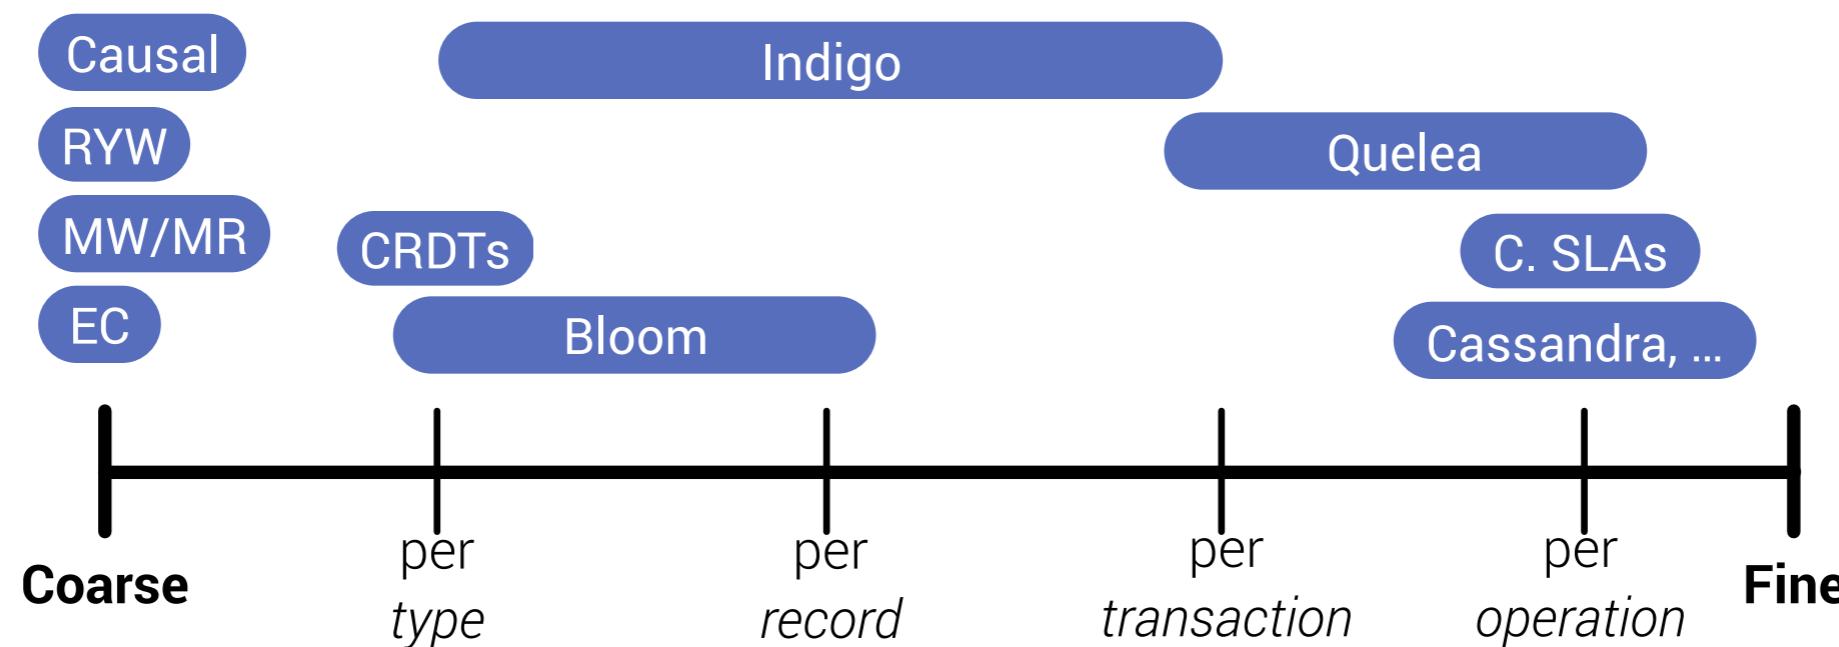
How should coordination be specified and enforced?

- ⚡ Only enforce constraints where **necessary**.
- ⚡ Consistency models are **too coarse-grained**.
- ⚡ **Synthesize fine-grained synchronization.**



How can programming models help make these trade-offs?

- ⚡ Specify abstract state and behavior **by-construction**.
- ⚡ Explicit performance **targets, feedback** about results.



Introduction: Contention



- ✓ Sources of contention.
- ✓ 3 broad approaches to mitigating contention.

Background: Trading off consistency



- ✓ Consistency for architects.
- ✓ Ordering and visibility constraints.
- ✓ Dealing with uncertainty.
- ✓ **Review:** Programming tradeoffs.

➤ Proposal: Disciplined Inconsistency



Make tradeoffs *explicit* and *safe* with types.
Synthesize synchronization/enforcement logic.

Bounds \Leftrightarrow Uncertainty

- *Bounds* on reads determine the *uncertainty* of results

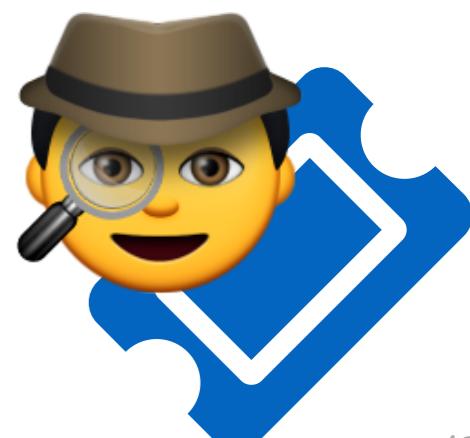
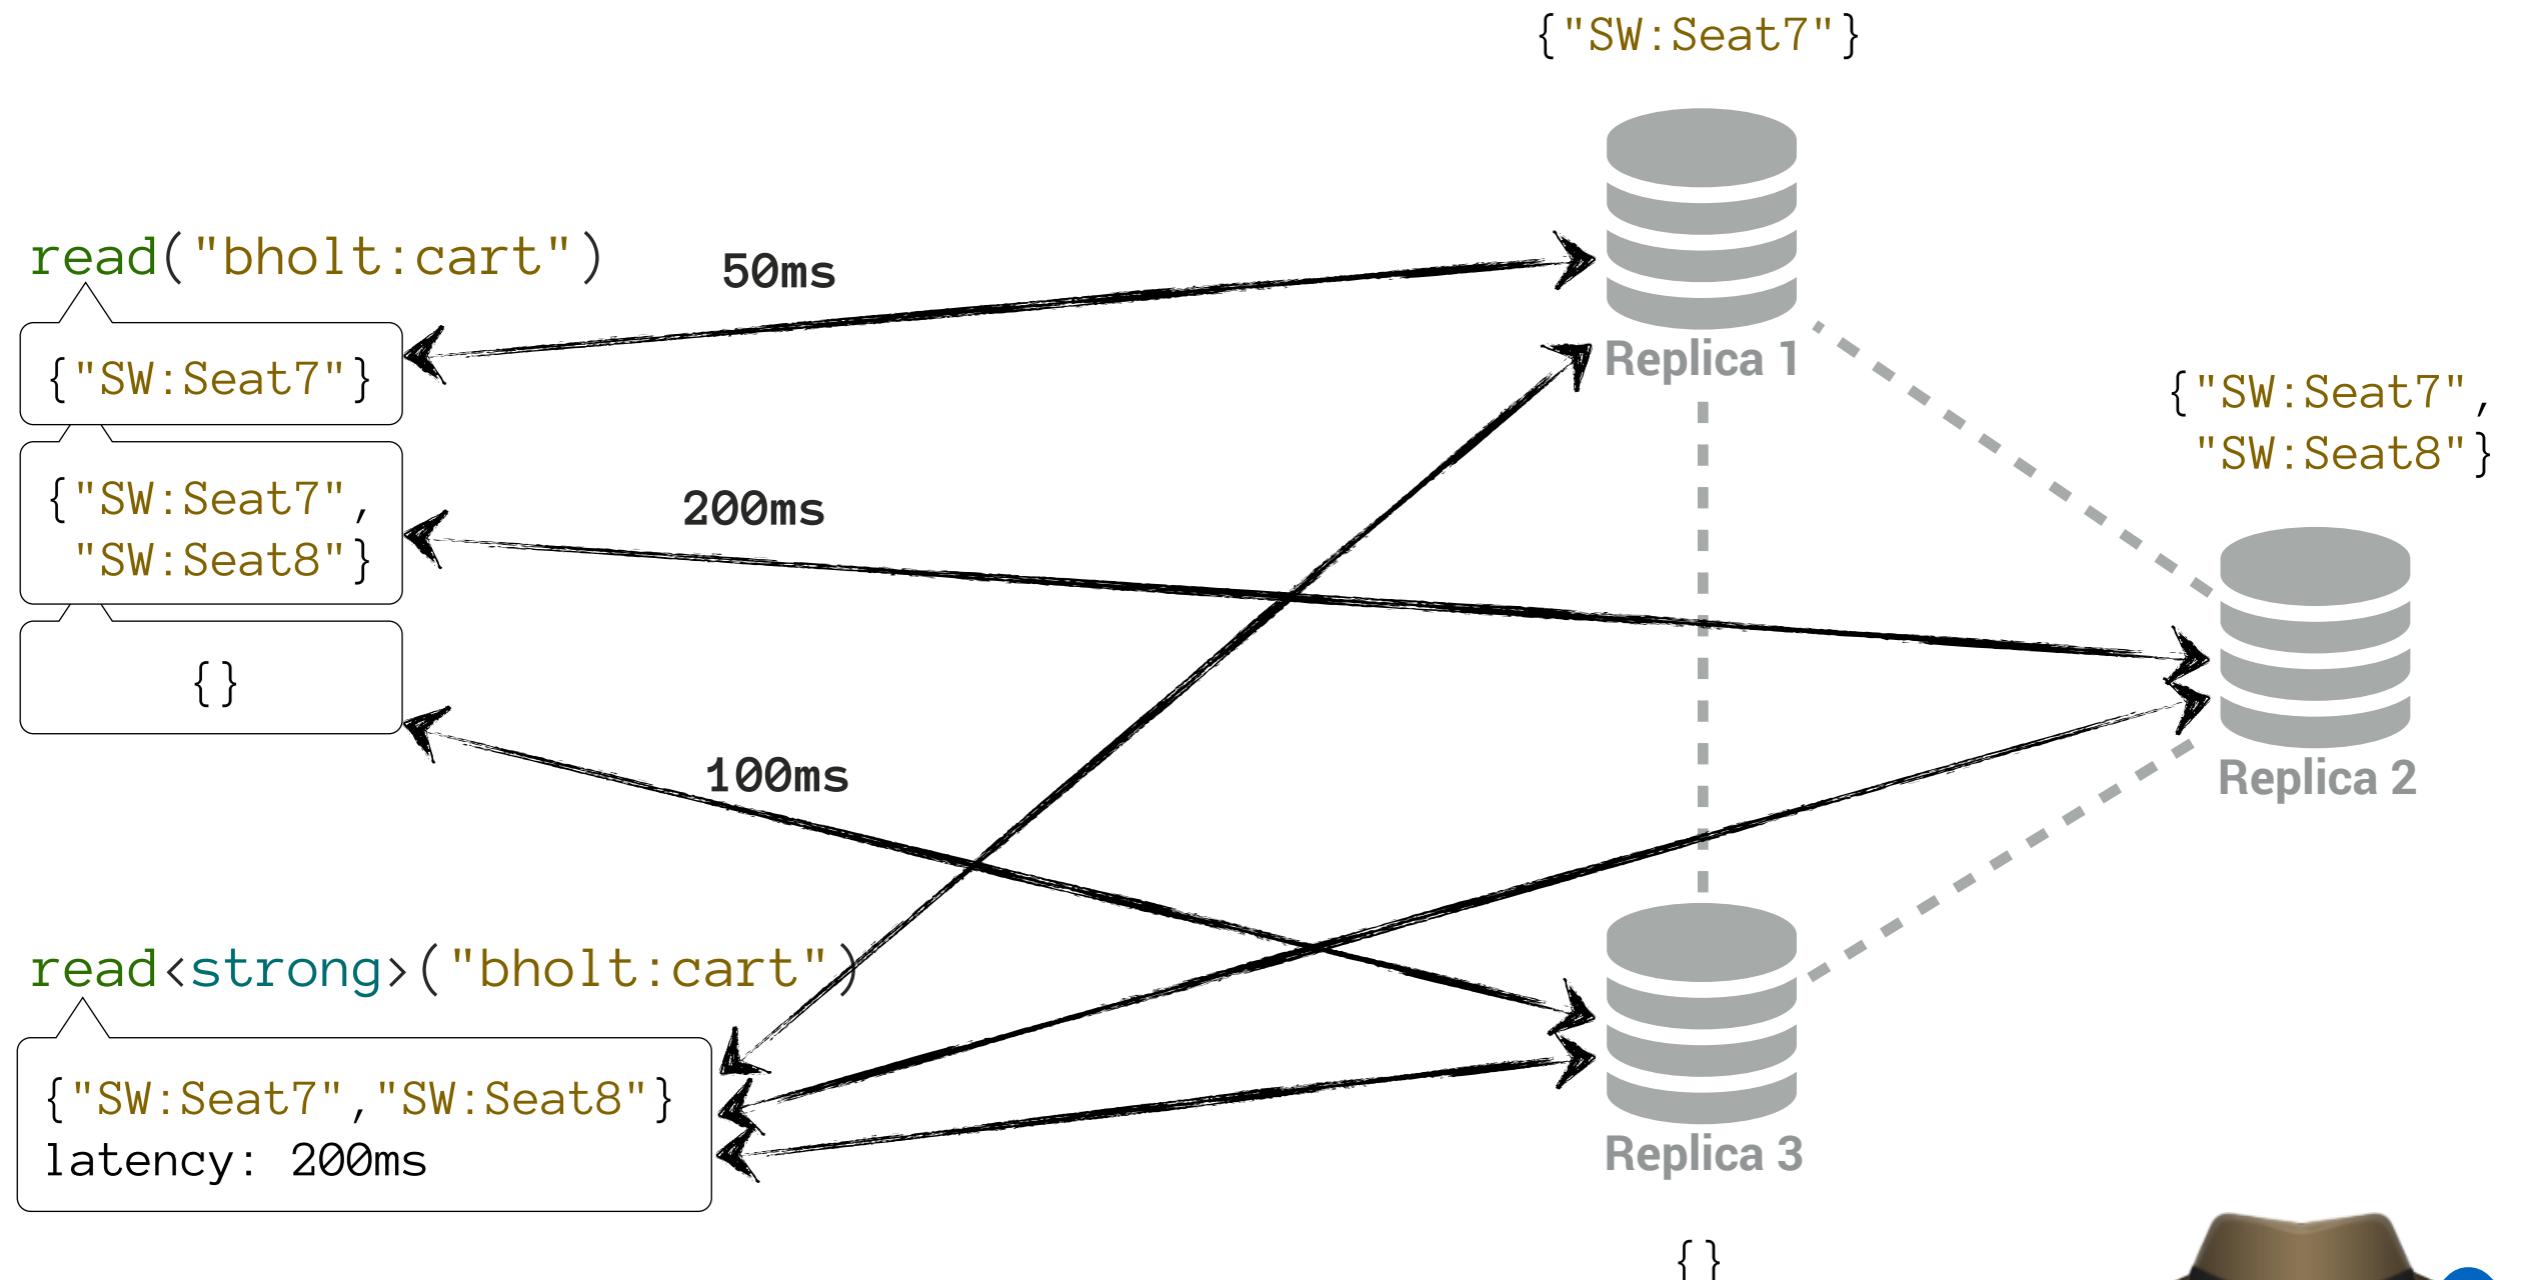
System Design

- Statically determine the type of uncertainty implied by bounds
- IPA Types
- Enforcement System

Case Studies

- TicketSleuth
- Twitter

Disciplined Inconsistency → Bounds ⇔ Uncertainty



Disciplined Inconsistency → Bounds ⇔ Uncertainty

⚡ Bounds on reads determine **uncertainty** of result.

`read("bholt:cart")`

{ "SW:Seat7" }

{ "SW:Seat7",
 "SW:Seat8" }

{ }

`read("bholt:cart")`

{ "SW:Seat7", "SW:Seat8" }

latency: 200ms

`read<latency:100ms>("bholt:cart")`

{ "SW:Seat7" }

latency: 51ms



R1: { latency: ~50ms, version: ~v-1 }

R2: { latency: ~200ms, version: ~v }

Monitor

R3: { latency: ~100ms, version: ~v-5 }

{ "SW:Seat7" }



Replica 1

{ "SW:Seat7",
 "SW:Seat8" }

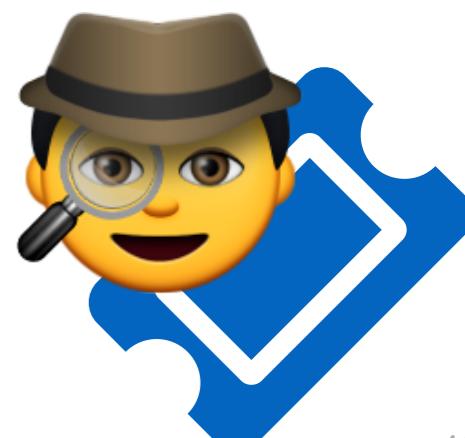


Replica 2



Replica 3

{ }



Disciplined Inconsistency → Bounds ⇔ Uncertainty

⚡ Bounds on reads determine **uncertainty** of result.

Bounds

`read` ⇒

`read<ec>` ⇒
`read<available>`

`latency < 100ms` ⇒

`value ≤ 0`
`value ± 5%`
`unique(_)` ⇒

Uncertainty

`T`

`Inconsistent<T>`

`Stale<T>`

`Probabilistic<T>`

`Interval<T>`

`Interval<T>`

`Leased<T>`

Specified with
abstract data
types

IPA Types

Disciplined Inconsistency → System Design

Abstract Data Types

encode **semantics** and
value & performance **bounds**

IPA Types

represent **uncertainty**

Type System

- **statically check** or **infer** that IPA types for return values match the specified bounds
- determine which techniques to use to enforce the bounds

Enforcement System

library of coordination/synchronization strategies

IPA Types

Inconsistent

`Inconsistent<T>`

`Stale<T>`

`Leased<T>`

Probabilistic

`Probabilistic<T>`

Approximate

`Interval<T>`

- if you don't know anything more, at least prevent bad data flow (a la EnerJ)

```
// initial load: as fast as possible
Inconsistent<List> x = Cart.read()
display_cart(x)
```

```
// on checkout, be precise
List y = Cart.read<strong>()
display_cart(y)
```

```
checkout(x) // catch bug on compile
```



Disciplined Inconsistency ➔ System Design

IPA Types

Inconsistent

`Inconsistent<T>`

`Stale<T>`

`Leased<T>`

- when did the replica last synchronize?
- may want to handle older data differently

```
// initial load: as fast as possible  
Stale<List> x = Cart.get()  
  
if (x.older_than(Seconds(5))) {  
    display_warning("May be out of date!")  
}  
display_cart(x)
```

Probabilistic

`Probabilistic<T>`

Approximate

`Interval<T>`



IPA Types

Inconsistent

`Inconsistent<T>`

`Stale<T>`

`Leased<T>`

- value associated with a *lease* promising no updates for the duration

```
// initial load: as fast as possible
```

```
Leased<List> x = Cart.get()
```

```
if (x.expired()) {  
    display_warning("Out of date!")  
} else {  
    display_cart("Current: " + x)  
}
```

Probabilistic

`Probabilistic<T>`

Approximate

`Interval<T>`



IPA Types

Inconsistent

`Inconsistent<T>`

`Stale<T>`

`Leased<T>`

Probabilistic

`Probabilistic<T>`

Approximate

`Interval<T>`

- probabilistic model (like PBS) that predicts staleness & update rate
- either a *single value* with a *confidence* (staleness domain)
- or a distribution in the *value domain*

```
// bounded latency (e.g. < 200ms)
Prob<Int> x = Tickets("StarWars").remaining()

if ( (x < 10).confidence(0.5) ) {
    display_warning("Purchase soon!")
}
display_tickets(x)
```

Uncertain<T>: A First Order Type for Uncertain Data (ASPLOS'14)

J. Bornholt, T. Mytkowicz, K.S. McKinley



IPA Types

Inconsistent

`Inconsistent<T>`

`Stale<T>`

`Leased<T>`

Probabilistic

`Probabilistic<T>`

Approximate

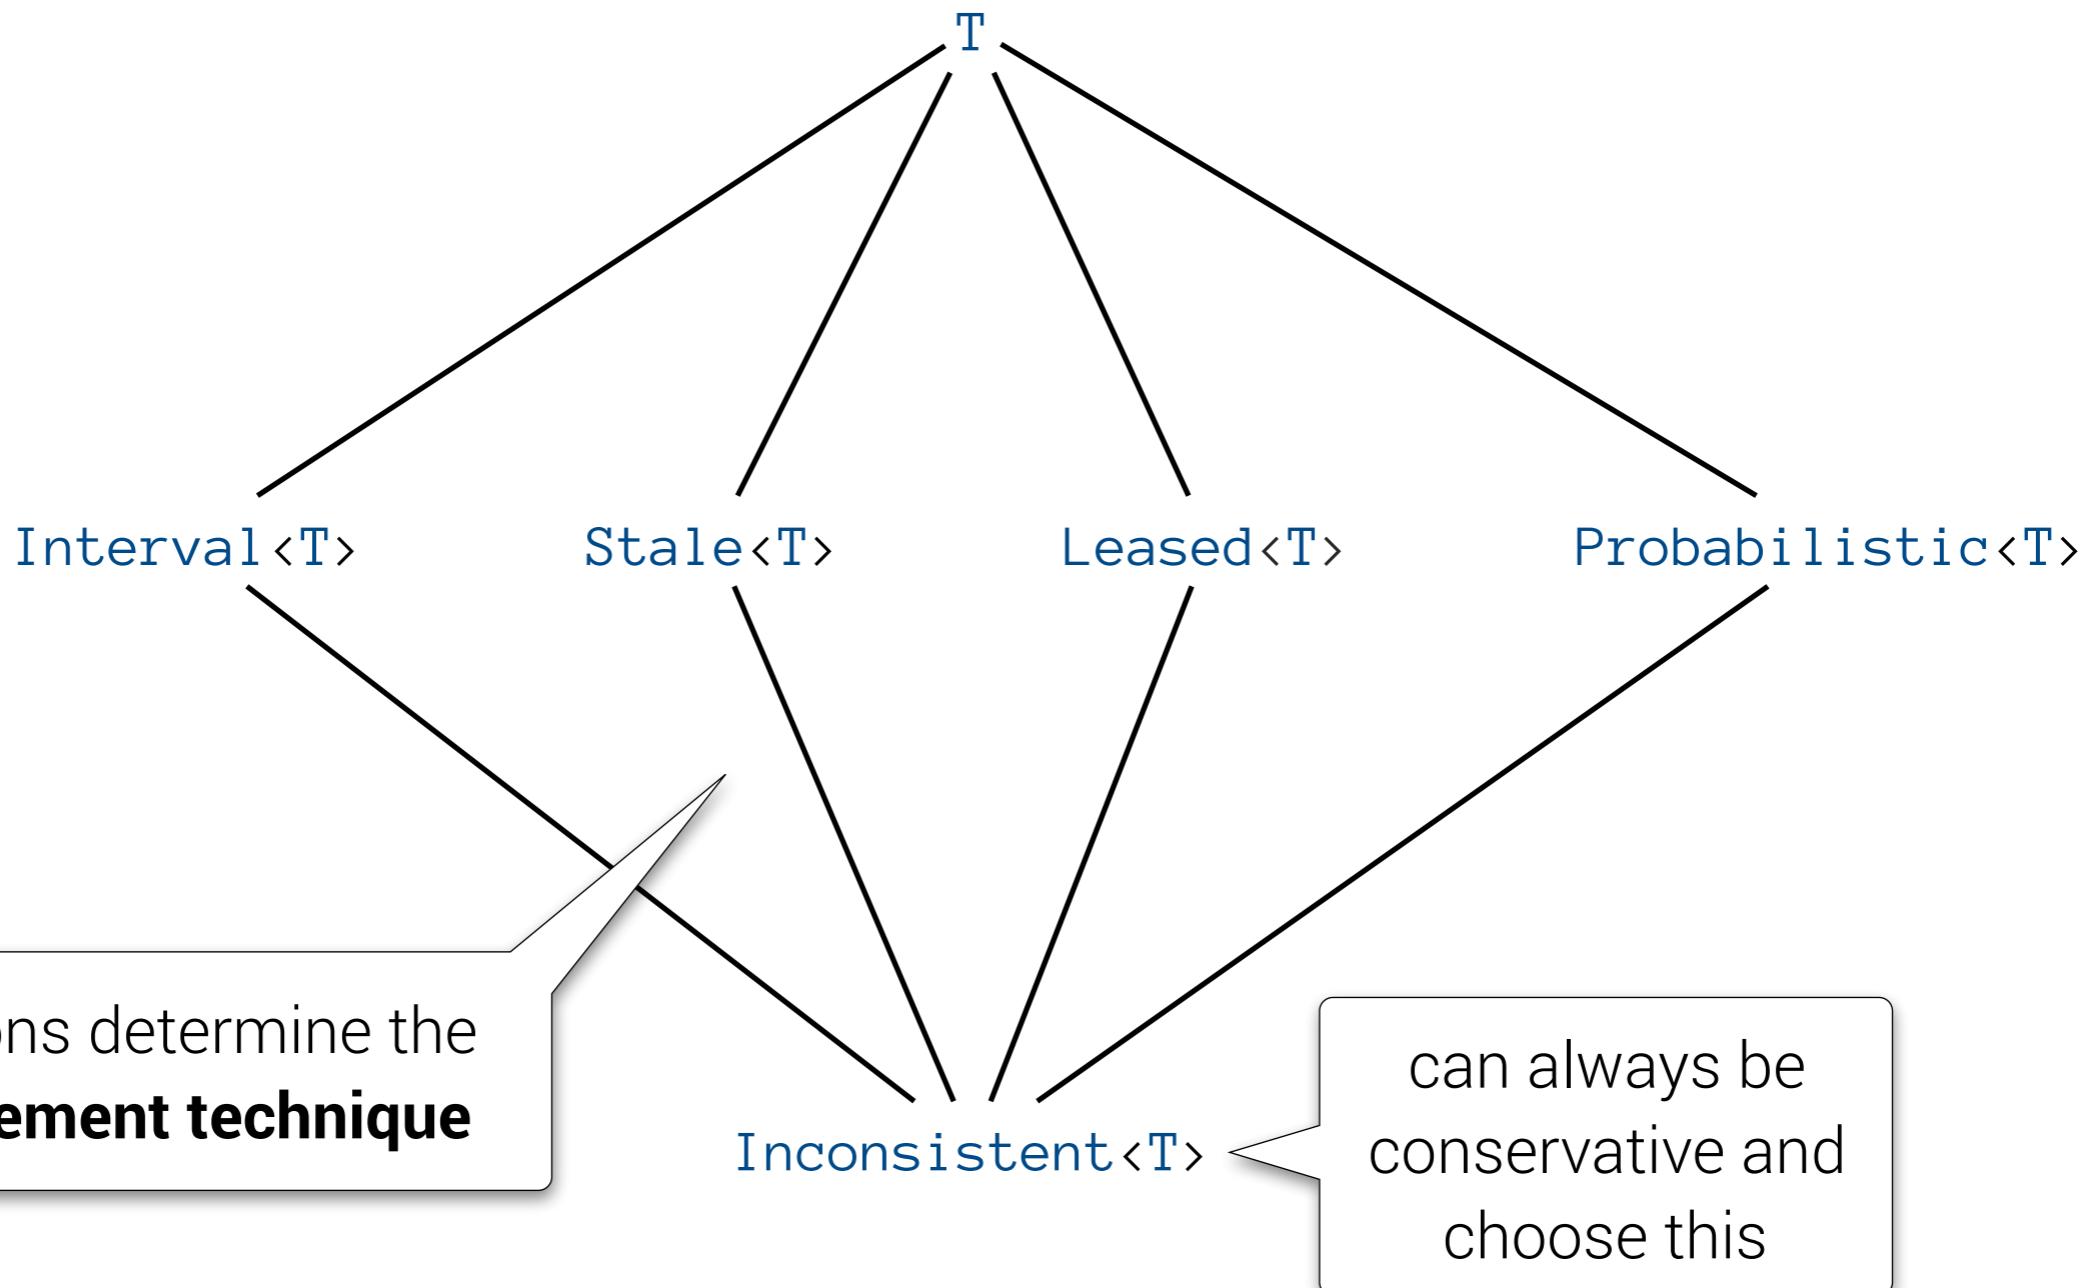
`Interval<T>`

- range of possible values
- any type with a *partial order* over values

```
// bounded latency (e.g. < 200ms)
Interval<Int> x = Tickets("StarWars").remaining()

if (x.min > 1M) {
    // e.g. '1.4M'
    display_tickets((x/100k) + "M")
} else if (x.min > 1000) {
    // e.g. '2.7k'
    display_tickets((x.min/100) + "k")
} else {
    // e.g. '5-20'
    display_tickets(x.min + "-" + x.max)
}
```

IPA Type Hierarchy

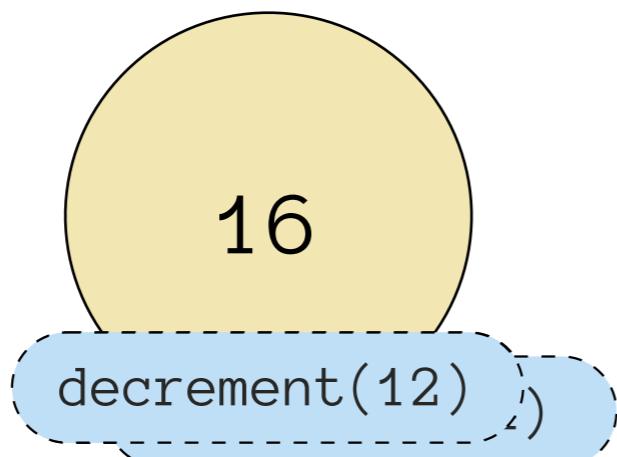


Disciplined Inconsistency ➔ Enforcement System

*How can you enforce **hard value bounds** on weak replication?*

Escrow

- Fragment value, hold parts *in escrow* until ready to commit.
- Take from *pool* of available updates / fragments.



Txn 1
decrement(4)
...

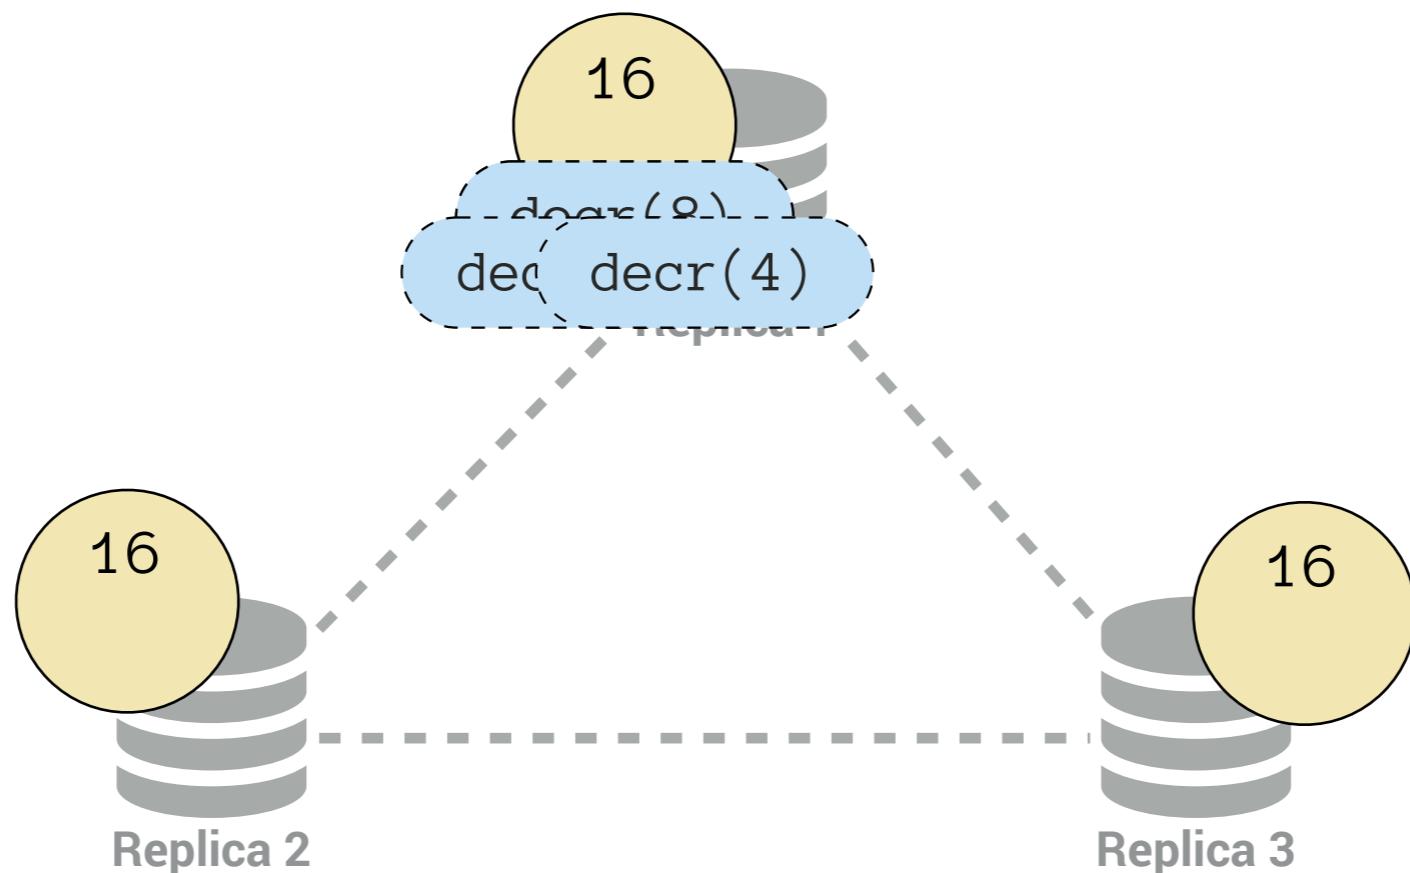
The Escrow Transactional Method.
(Transactions on Database Systems, 1986)
Patrick E. O'Neil

Disciplined Inconsistency ➔ Enforcement System

*How can you enforce **hard value bounds** on weak replication?*

Reservations

- Pre-allocate permission to do updates and distribute to replicas.
- Moves coordination off the critical path.

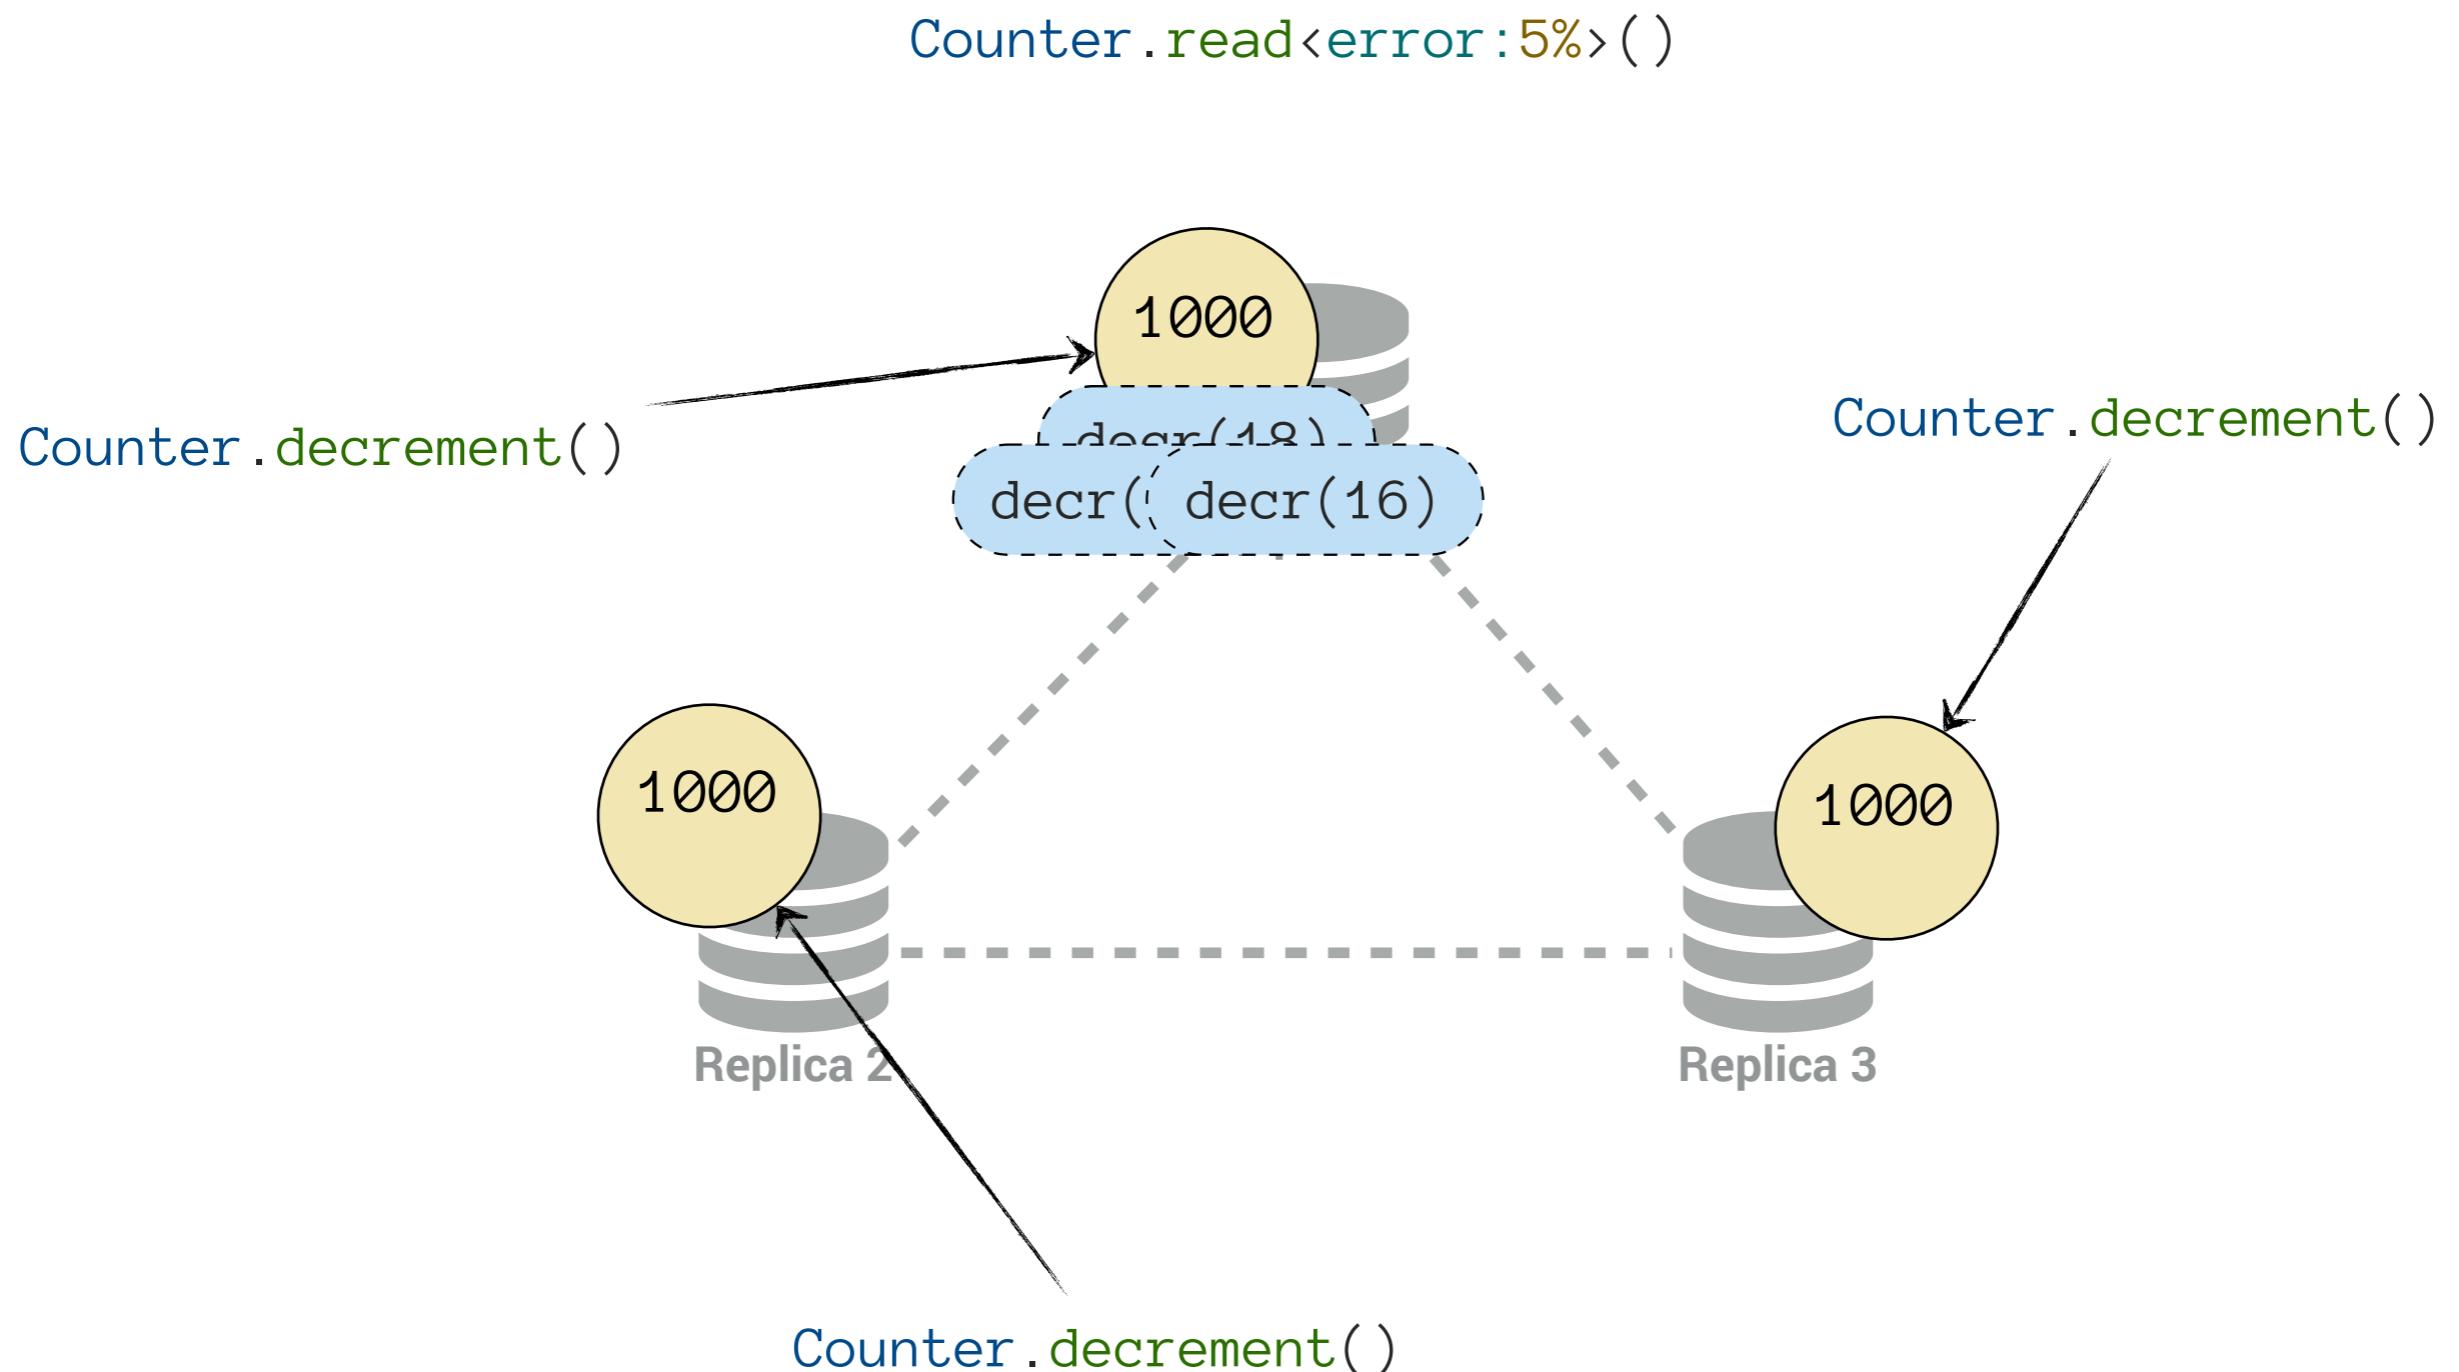


Reservations for conflict avoidance in a mobile database system.

N. Preguiça, J.L. Martins, M. Cunha, and H. Domingos.

Disciplined Inconsistency ➔ Enforcement System

How can you enforce **hard value bounds** on weak replication?



Escrow / Reservations

- provide Interval<T>, Leased<T>

Monitoring & prediction system

- like those in PBS/Consistency-based SLAs
- provide Probabilistic<T>, Stale<T>

More...

Bounds \Leftrightarrow Uncertainty

- *Bounds* on reads determine the *uncertainty* of results

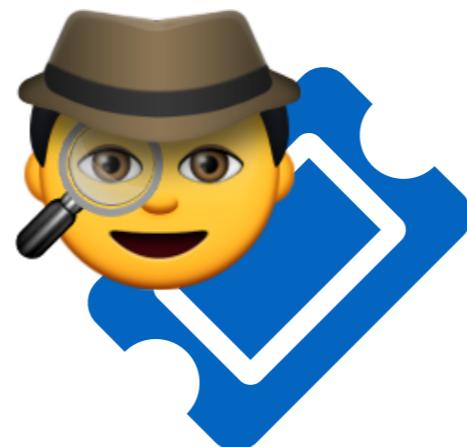
System Design

- Statically determine the type of uncertainty implied by bounds
- IPA Types
- Enforcement System

Case Studies

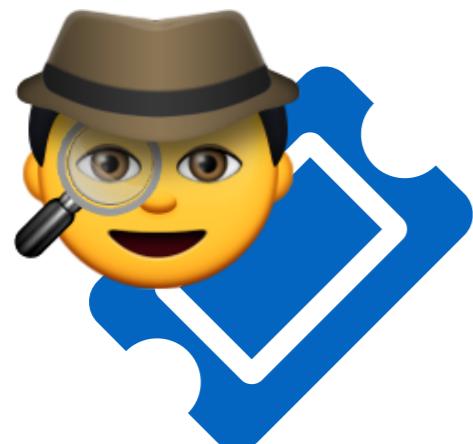
- TicketSleuth
- Twitter

Disciplined Inconsistency ➤ Case Studies



TICKETSLEUTH

Disciplined Inconsistency ➔ Case Studies

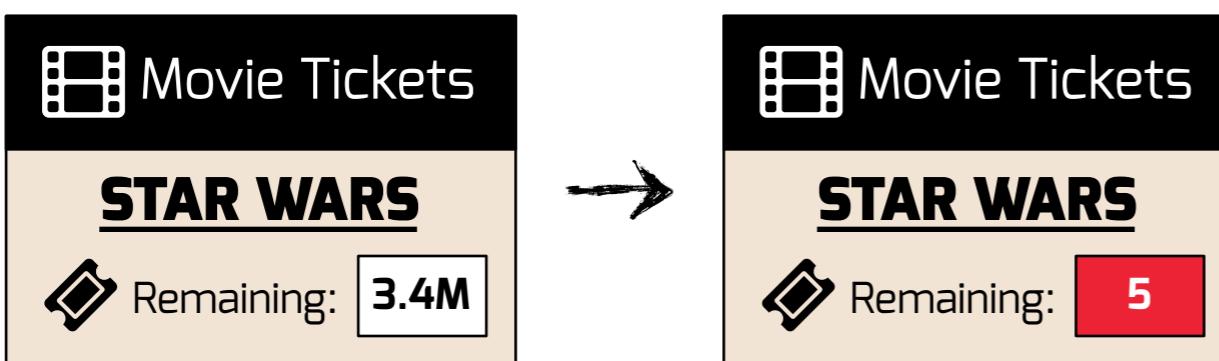


hard value constraint

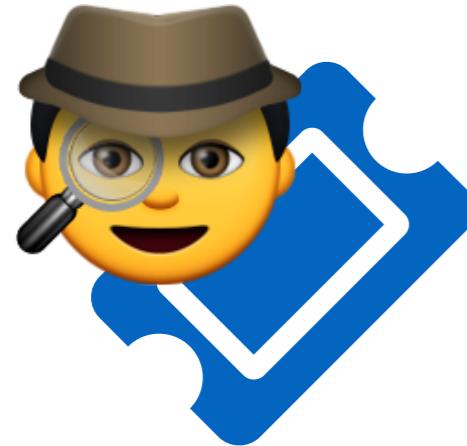
1. Never over-sell tickets.
2. Show estimate of remaining tickets.
3. Initial page must load in 100ms (99th percentile)

IPA type

performance constraint



Disciplined Inconsistency ➔ Case Studies



- 1. Never over-sell tickets.
 - 2. Show estimate of remaining tickets.
 - 3. Initial page must load in *100ms* (99th percentile)
- hard value constraint*
- IPA type*
- performance constraint*

```
// Pool supporting unique `take`s
// bounded by initial size, and
// reading an approximate size.
class Pool< type T, Latency L > {

    fn take() → T

    // return the range of possible
    // current sizes
    fn size<latency: L>() → Interval<int>
}
```



```
// derive app-specific type from generic IPA type,
// select constraints: 100 ms latency
type MovieTickets = Pool<TicketID, Latency(100ms)>;
```

Disciplined Inconsistency ➤ Case Study: Twitter

Timeline <: List<TweetID>

```
fn slice<latency:100ms>(int start, int end) → Stale<List<TweetID>>
```

Retweets <: Set<UserID>

```
fn size<tolerance:1%>() → Interval<int>
```

```
def timeline(user):
```

```
    # Timeline.range: 100 ms latency bound
```

```
    tweets: Stale<List<TweetID>> = Timeline(user).slice(0, 10)
```

```
    # check if timeline may be missing tweets older than 5s
```

```
    if posts.older_than(Seconds(5)):
```

```
        display_warning("May be out of date.")
```

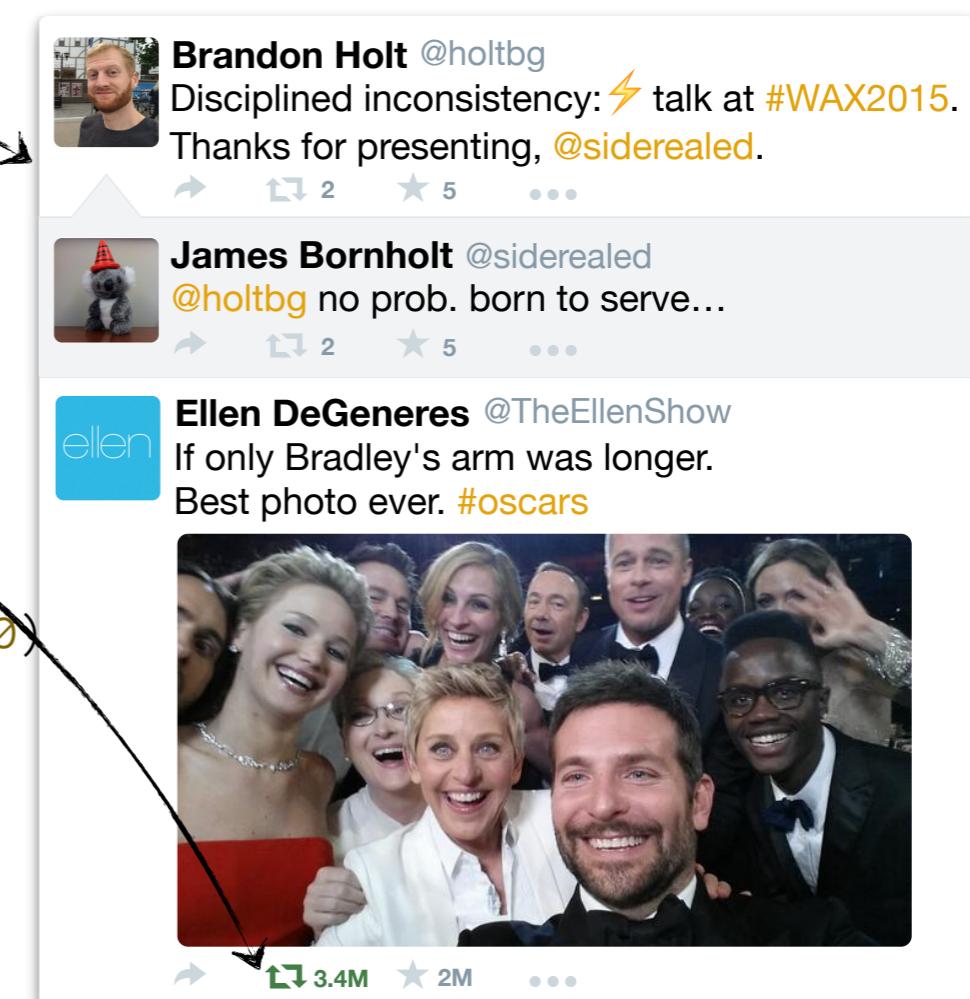
```
for t in tweets:
```

```
    tweet: Map = Tweet(t).get()
```

```
    # Retweeters.size: 1% tolerance
```

```
    retweets: Interval<Int> = Retweets(t).size()
```

```
    display_tweet(tweet, retweets)
```



Introduction: Contention



- ✓ Sources of contention.
- ✓ 3 broad approaches to mitigating contention.

Background: Trading off consistency



- ✓ Consistency for architects.
- ✓ Ordering and visibility constraints.
- ✓ Dealing with uncertainty.
- ✓ Programming tradeoffs.



Proposal: Disciplined Inconsistency

- ✓ Make tradeoffs *explicit* and *safe* with types.
- ✓ Reason about *values* rather than *ordering*.
- ✓ Synthesize synchronization/enforcement logic.

Mitigating Contention in Distributed Systems

Generals Exam

Brandon Holt

Committee

Luis Ceze, Mark Oskin, Dan Ports, Jevin West